

Questionnaire & Join MS teams today

- Please help me understand you more



Questionnaire (before 09/13)



How to apply MS account

- And, Join MS teams right after the class
- Code: ig0604f



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

Operating System

Lecture 01: Introduction

Shuo-Han Chen (陳碩漢),

shch@nycu.edu.tw

Wed. 10:10 - 12:00 EC115 +

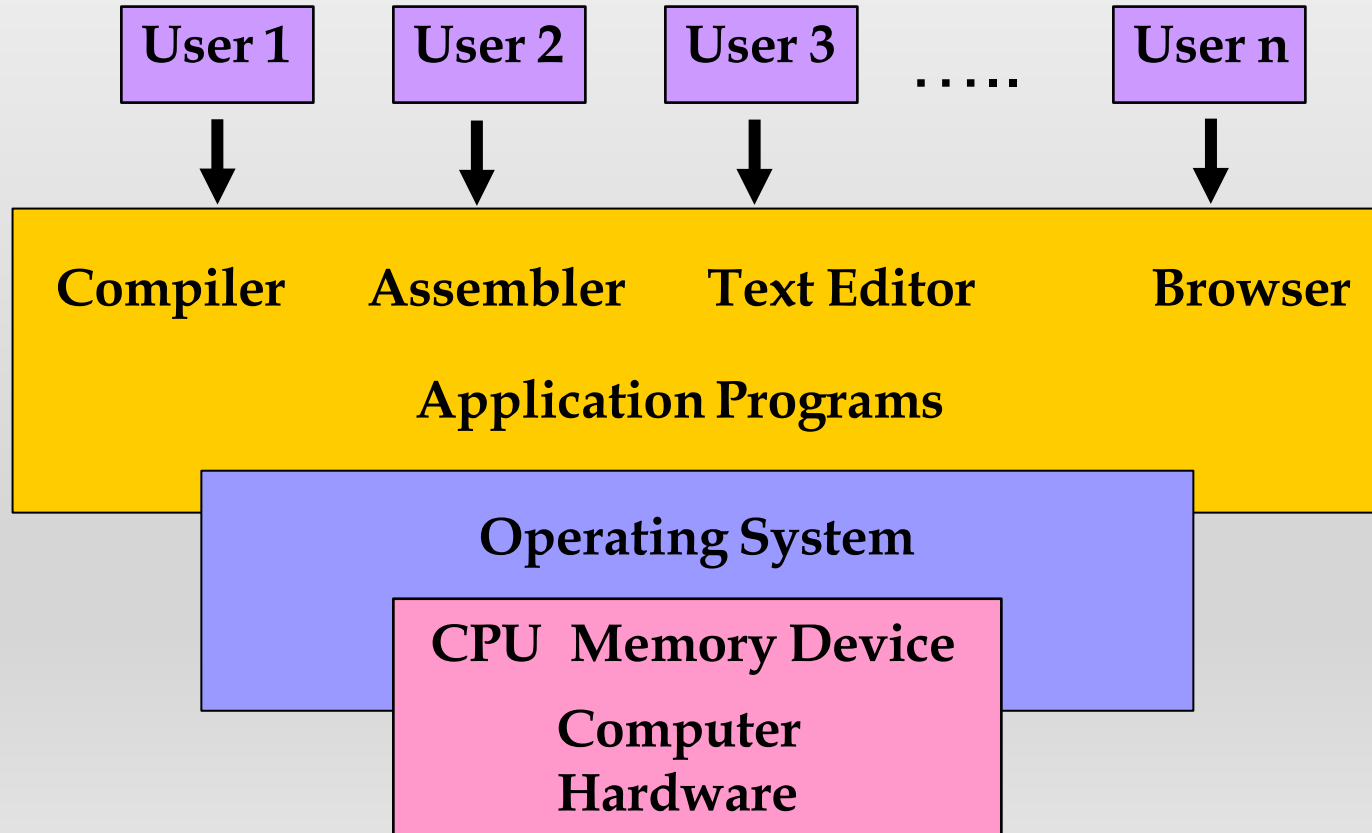
Fri. 11:10 – 12:00 Online

Course Schedule

| W | Date | Lecture | Online | Homework |
|----|----------|---|--------|----------|
| 1 | Sept. 4 | Lec01: Course Overview & Historical Prospective | | |
| 2 | Sept. 11 | Lec02: Introduction | V | |
| 3 | Sept. 18 | Lec03: OS Structure | V | HW01 |
| 4 | Sept. 25 | Lec04: Processes Concept | V | |
| 5 | Oct. 2 | Lec08: Memory Management | V | |
| 6 | Oct. 9 | Lec09: Virtual Memory Management | V | HW02 |
| 7 | Oct. 16 | Lec05: Process Scheduling | V | |
| 8 | Oct. 23 | School Midterm Exam | | |
| 9 | Oct. 30 | Lec06: Process Synchronization | V | |
| 10 | Nov. 6 | Lec07: Deadlocks | V | HW03 |
| 11 | Nov. 13 | Lec10: File System Interface | V | |
| 12 | Nov. 20 | School Event – No class | | |
| 13 | Nov. 27 | Lec11: File System Implementation | V | HW04 |
| 14 | Dec. 4 | Lec12: Mass Storage System | V | |
| 15 | Dec. 11 | Lec13: IO Systems | V | |
| 16 | Dec. 18 | School Final Exam | | |

Computer System

- Four components:
 - Hardware, OS, Application, User




Computer System

- **User** — people, machines, other computers
- **Application** — define the ways in which the system resources are used to **solve the computing problems**
- **Operating System** — **controls** and **coordinates** the use of the **hardware/resources**
- **Hardware** — provides basic **computing resources** (CPU, memory, I/O devices)


What is an Operating System

- An operating system is the “*permanent*” software that **controls/abstracts hardware resources** for user applications

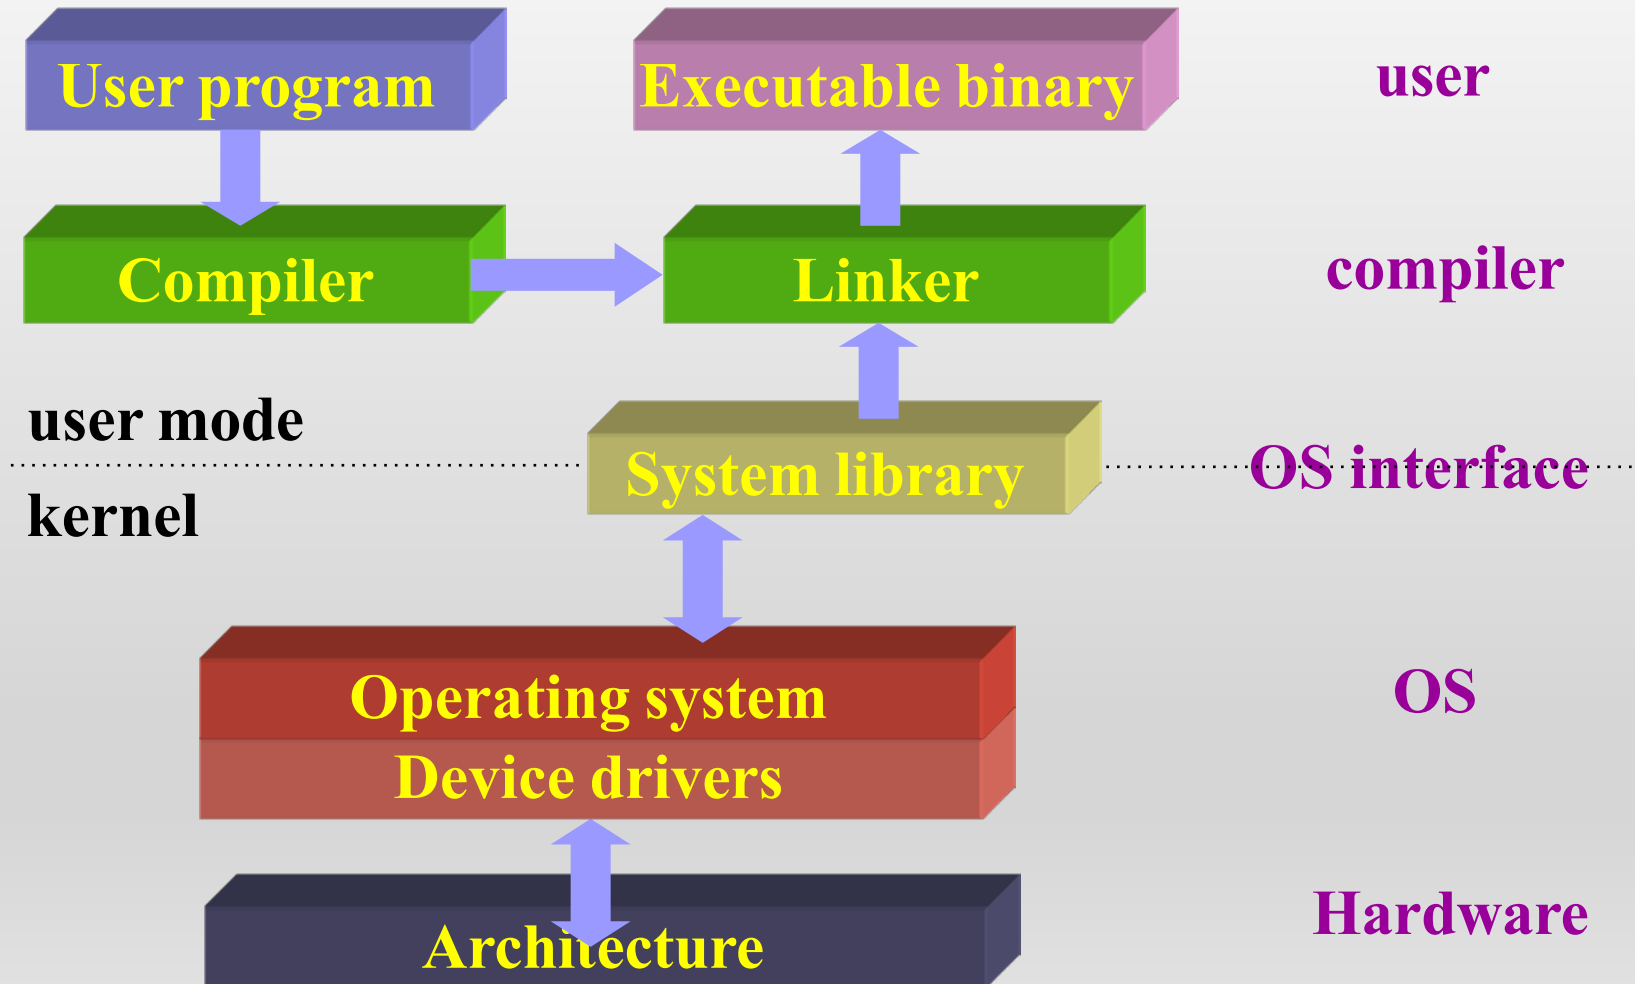


Multi-tasking Operating Systems

- Manages resources and processes to support different user applications
- Provides Applications Programming Interface (API) for user applications



General-Purpose Operating Systems



Definition of an Operating System

- **Resource allocator** – manages and allocates resources to insure efficiency and fairness
- **Control program** – controls the execution of user programs and operations of I/O devices to prevent errors and improper use of computer
- **Kernel** – the one program running at all times (all else being system/application programs)
 - No universally accepted definition

Goals of an Operating System

- **Convenience**

- make computer system easy to use and compute
- In particular for small PC

- **Efficiency**

- use computer hardware in an efficient manner
- Especially for large, shared, multiuser systems

- Two goals are sometimes **contradictory**
- In the past, efficiency is more important

Importance of an Operating System

- System API are the *only* interface between user applications and hardware
 - API are designed for general-purpose, not performance driven
- OS code cannot allow any bug
 - Any break (e.g. invalid access) causes reboot
- The owner of OS technology *controls* the software & hardware industry
- Operating systems and computer architecture influence each other

Modern Operating Systems

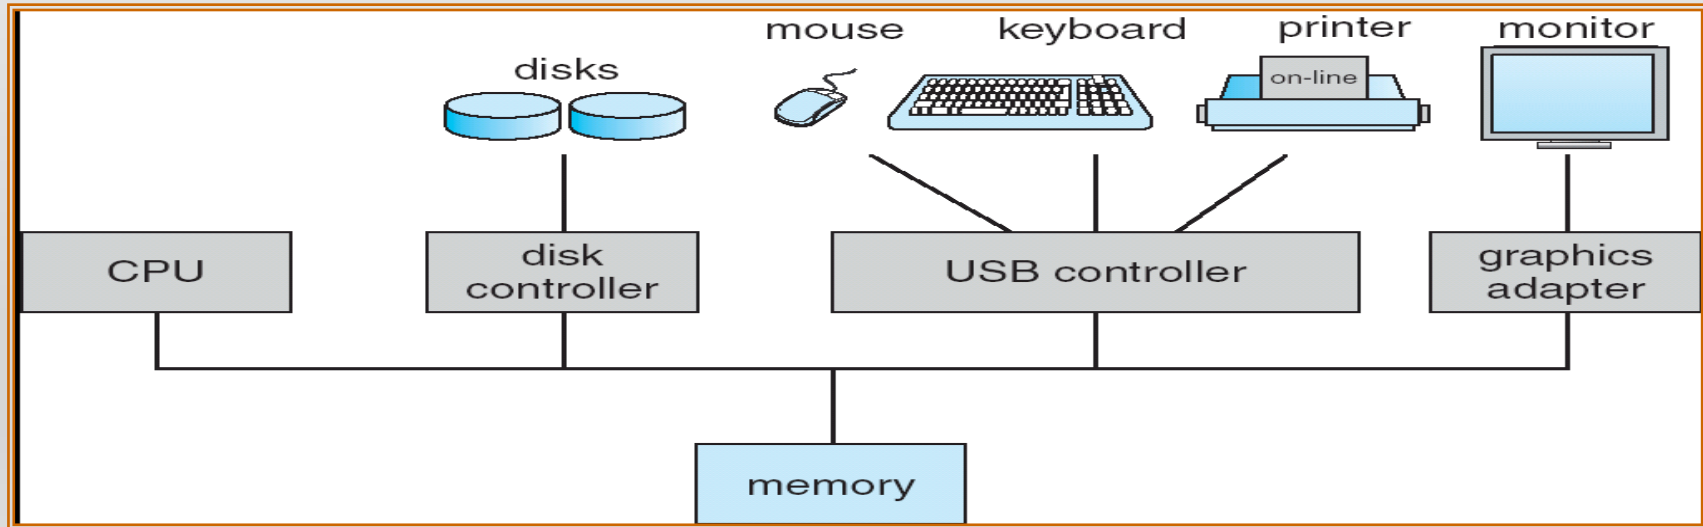
- x86 platform
 - Linux (CentOS, Redhat, openSUSE, Ubuntu, etc)
 - Windows (Windows10, XP, 2000, etc)
- PowerPC platform – Mac OS
- Smartphone Mobile OS
 - Android, iOS, Windows10 Mobile, Ubuntu Touch
- Embedded OS
 - Embedded Linux(Android, WebOS), Windows CE
 - Raspberry Pi, Xbox, etc

Review Slides (1)

- Definition of OS?
- Goals of OS?
- Importance of OS?

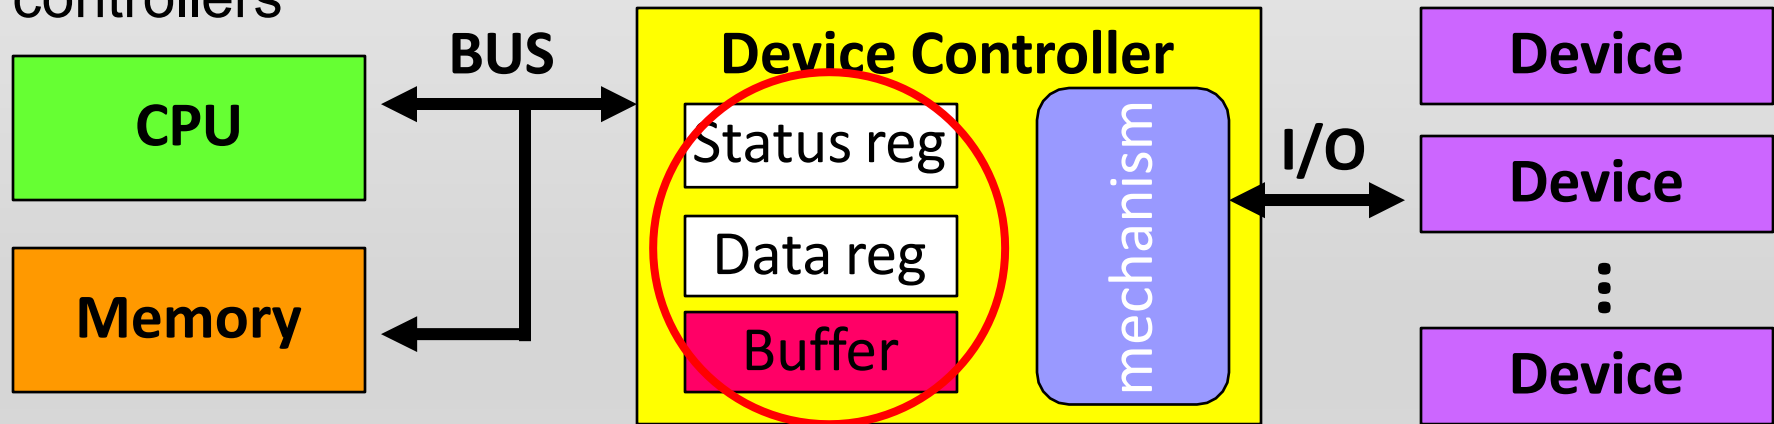
Computer-System Organization

- One or more CPUs, device controllers connect through **common bus** providing access to **shared memory**
- Goal: **Concurrent** execution of CPUs and devices competing for memory cycles



Computer-System Operations

- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- I/O is from the device to controller's local buffer
- CPU moves data from/to memory to/from local buffers in device controllers



Busy/wait output

- Simplest way to program device
 - Use instructions to test when device is ready

```
#define OUT_CHAR 0x1000      // device data register
#define OUT_STATUS 0x1001   // device status register
```

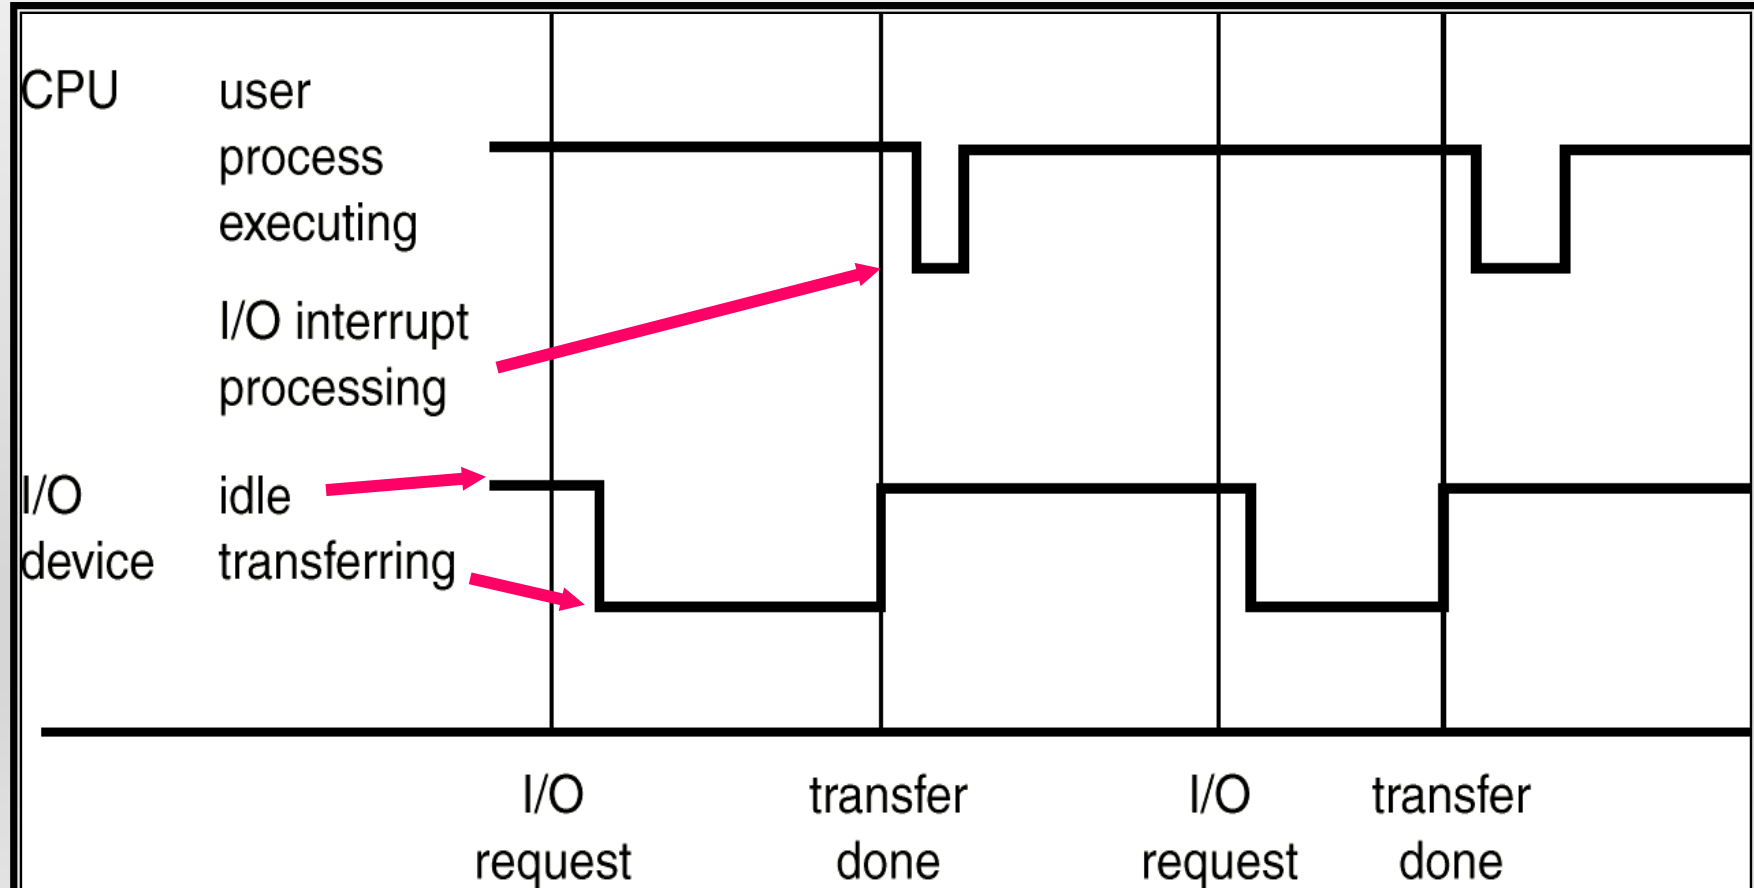
```
current_char = mystring;
while (*current_char != '\0') {
    poke(OUT_CHAR, *current_char);
    while (peek(OUT_STATUS) != 0); // busy waiting
    current_char++;
}
```


Interrupt I/O

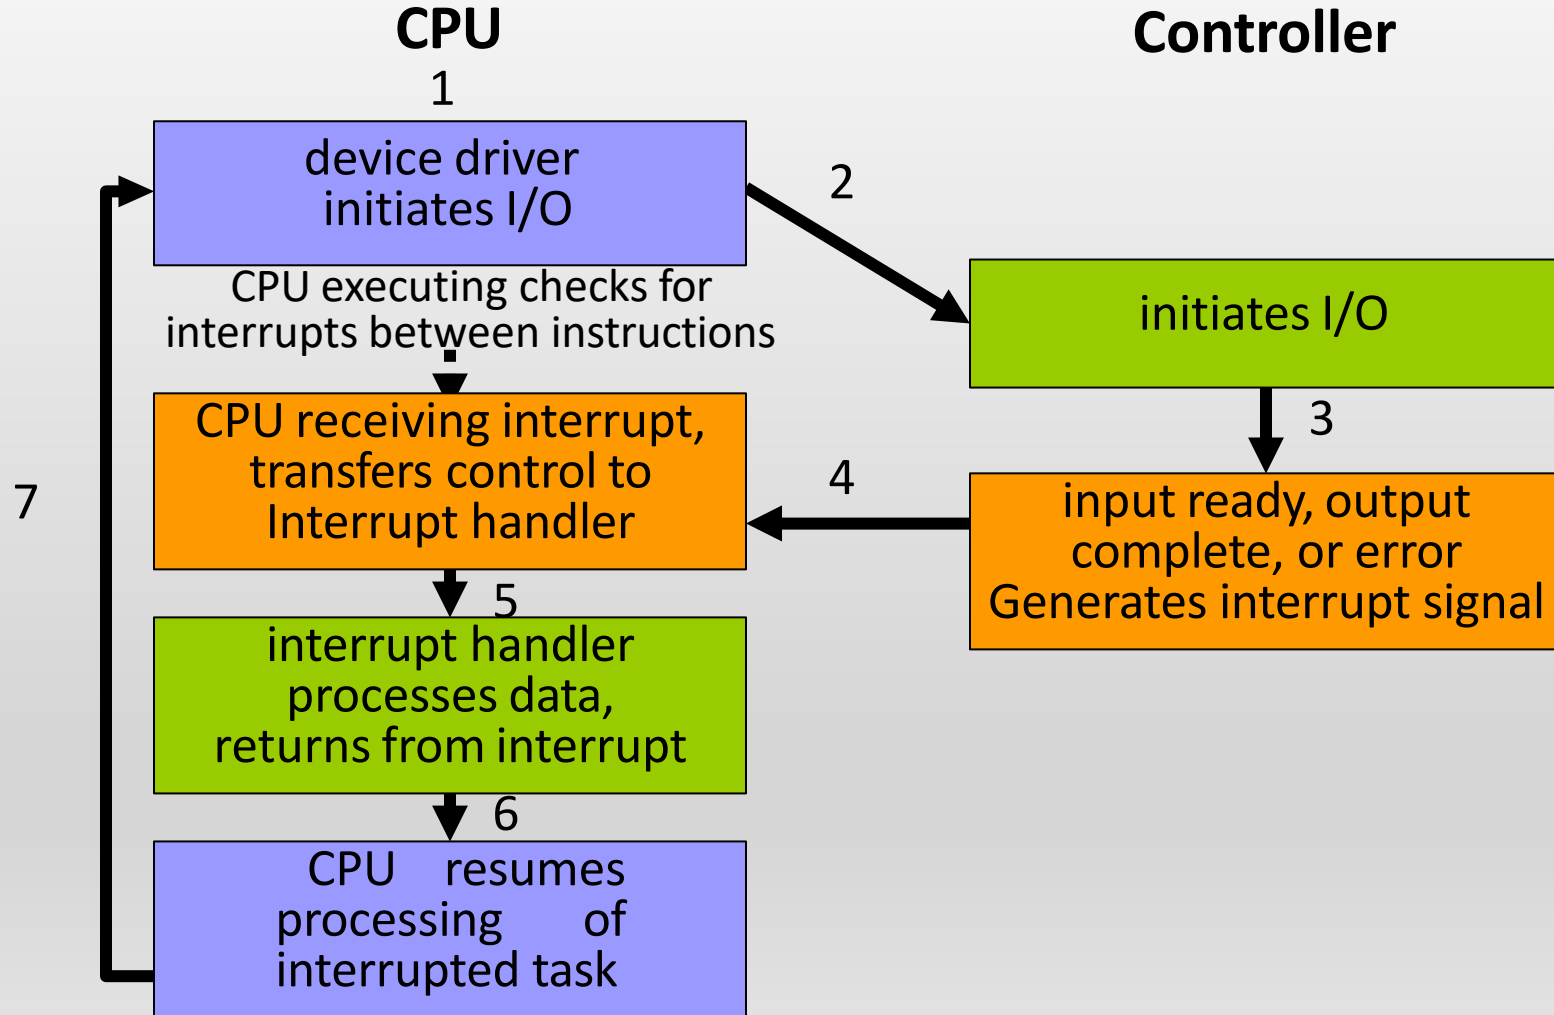
- Busy/wait is very inefficient
 - CPU can't do other work while testing device
 - Hard to do simultaneous I/O
- **Interrupts** allow a device to **change the flow of control in the CPU**
 - Causes subroutine call to handle device

Interrupt I/O Timeline

- Interrupt time line for I/O on a single process



Interrupt-Driven I/O

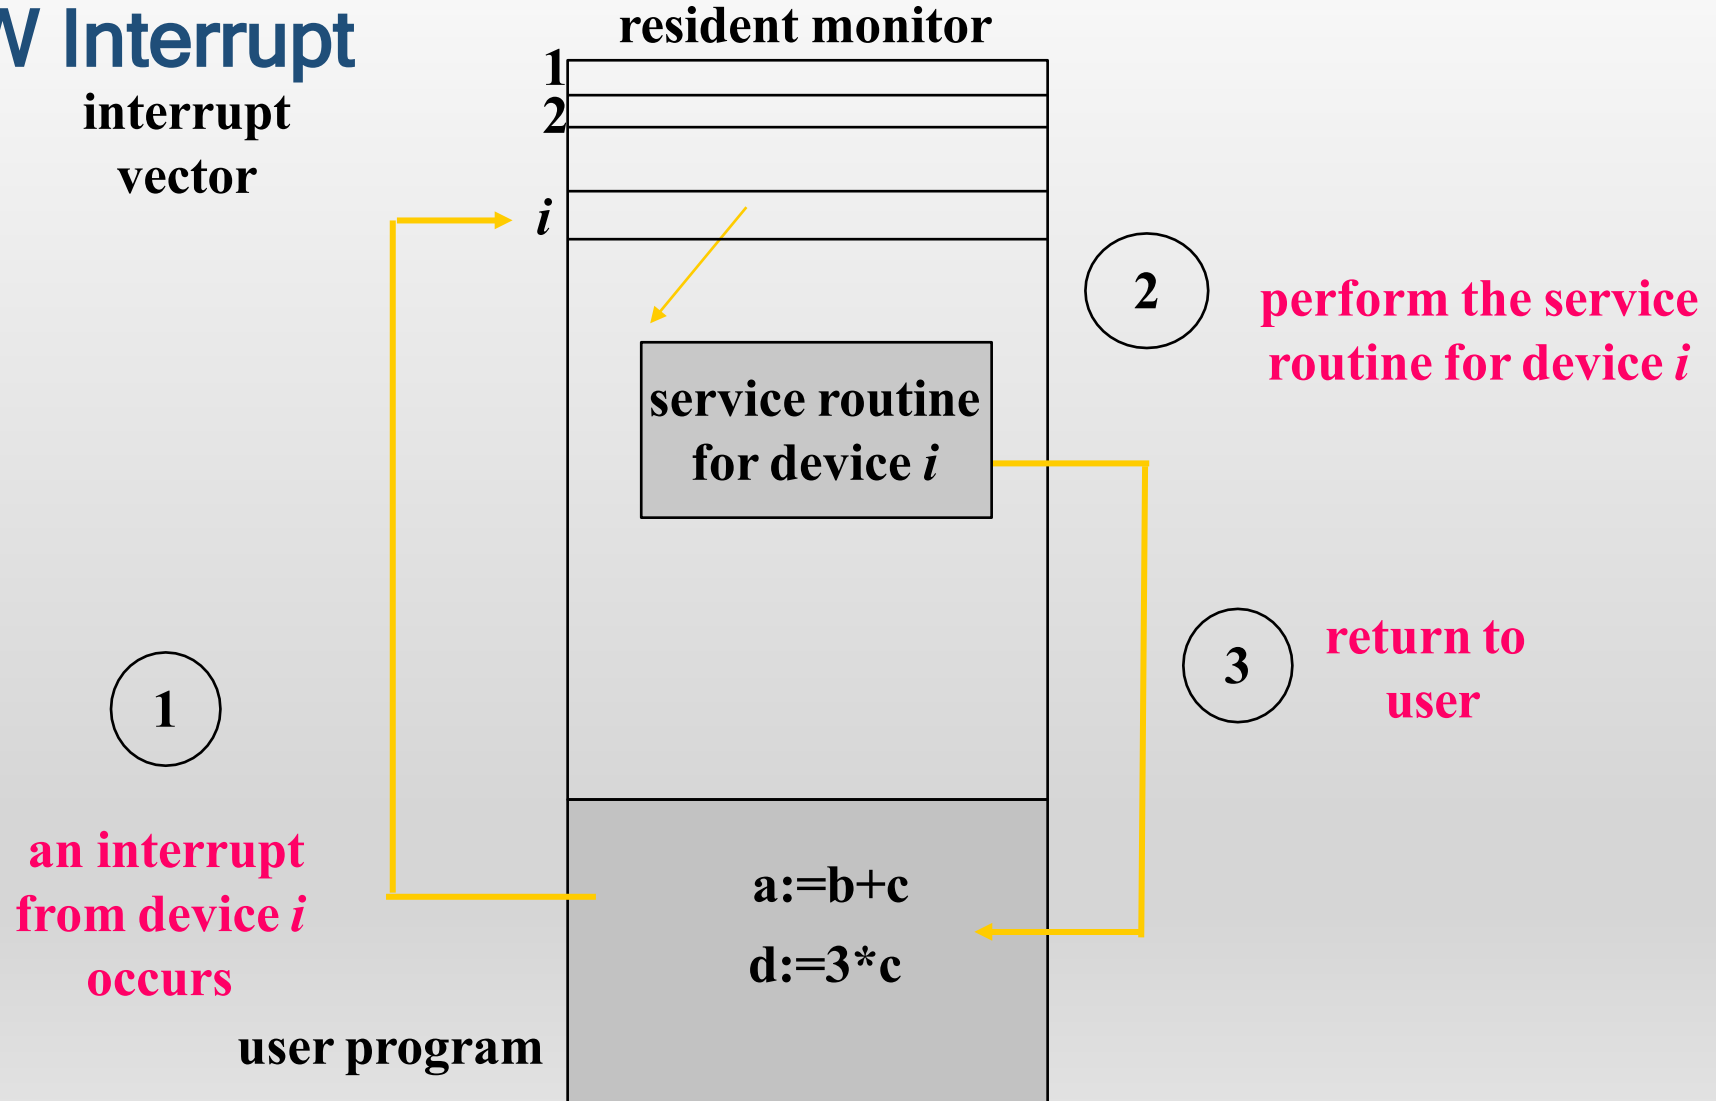


Interrupt

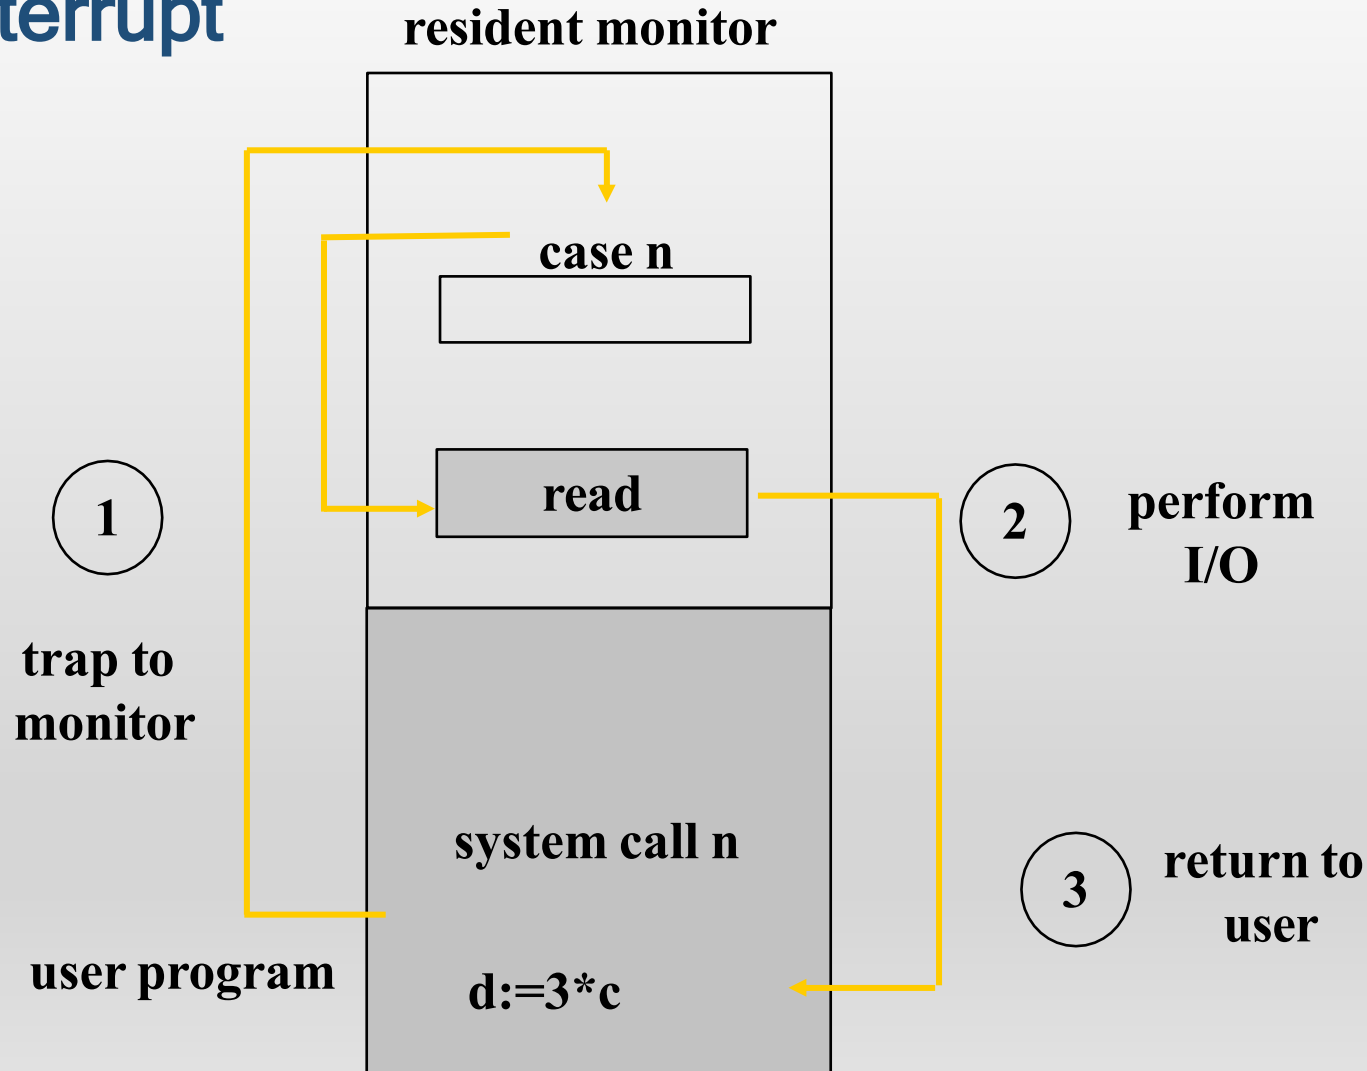
- Modern **OS are interrupt driven**
- The occurrence of an event is signaled by an interrupt from either hardware or software.
 - **Hardware** may trigger an interrupt at any time by sending a **signal** to CPU
 - **Software** may trigger an interrupt either by an **error** (division by zero or invalid memory access) or by a user request for an operating system **service (system call)**
- Software interrupt also called **trap**

HW Interrupt

interrupt
vector



SW Interrupt




Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the ***interrupt vector***, which contains the **addresses** (function pointer) of all the **service (i.e. interrupt handler) routines**
- Interrupt architecture must save the **address** of the **interrupted instruction**
- Incoming interrupts are ***disabled*** while another interrupt is being processed to prevent a lost interrupt

Review Slides (2)

- What is interrupt and how does it work?
- What is the difference between trap and interrupt?

Storage-Device Hierarchy



Storage-Device Hierarchy


- Storage systems organized in hierarchy.
 - **Speed** , Cost, Volatility
- **Main memory** – only large storage media that the **CPU can access directly**
 - RAM: Random Access Memory
- Secondary storage – extension of main memory that provides **large nonvolatile storage** capacity
 - Magnetic disk

RAM: Random-Access Memory

- DRAM (Dynamic RAM):
 - Need only **one transistor**
 - Consume **less power**
 - values must be periodically refreshed
 - Access Speed: **$\geq 30\text{ns}$**
- SRAM (Static RAM):
 - Need **six transistors**
 - Consume **more power**
 - Access Speed: **$10\text{ns}\sim 30\text{ns}$**
 - usage: **cache memory**

Disk Mechanism

- Speed of magnetic disk
 - Transfer time = data size / transfer rate
 - Positioning time (random access time)
 - seek time (cylinder) + rotational latency (sector)




Performance of Various Levels of Storage

| Level | 1 | 2 | 3 | 4 |
|---------------------------|---|-------------------------------|------------------|------------------|
| Name | registers | cache | main memory | disk storage |
| Typical size | < 1 KB | > 16 MB | > 16 GB | > 100 GB |
| Implementation technology | custom memory with multiple ports, CMOS | on-chip or off-chip CMOS SRAM | CMOS DRAM | magnetic disk |
| Access time (ns) | 0.25 – 0.5 | 0.5 – 25 | 80 – 250 | 5,000.000 |
| Bandwidth (MB/sec) | 20,000 – 100,000 | 5000 – 10,000 | 1000 – 5000 | 20 – 150 |
| Managed by | compiler | hardware | operating system | operating system |
| Backed by | cache | main memory | disk | CD or tape |


Caching

- Information in use **copied** from **slower** to **faster** storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there



Coherency and Consistency Issue

- The same data may appear in different levels
 - Issue: Change the **copy in register** make it inconsistent with other copies
- Single task accessing:
 - No problem, always use the **Highest level** copy
- Multi-task accessing:
 - Need to obtain the most recent value
- Distributed system:
 - Difficult b.c. copies are on different computers



Review Slides (3)

- Why storage hierarchy?
- Caching? involved issues?

Hardware Protection


- Dual-Mode Operation
- I/O Protection
- Memory Protection
- CPU Protection

Dual-Mode Operation

- What to protect?
 - Sharing system resources requires OS to ensure that an incorrect program cannot cause **other programs** to execute incorrectly
- Provide **hardware support** to differentiate between at least two modes of operations
 - **User mode** – execution done on behalf of a user
 - **Monitor mode** (also **kernel mode** or system mode) – execution done on behalf of **operating system**

Dual-Mode Operation (Cont'd)

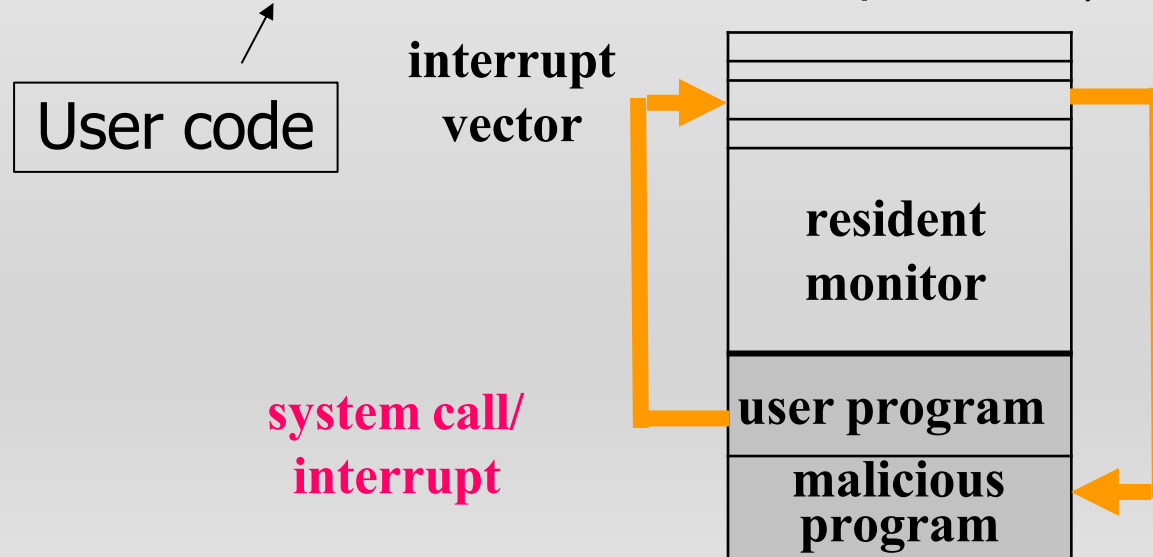
- Mode bit added to computer hardware to indicate the current mode: kernel (0) or user (1)
- When an interrupt/trap or fault occurs, hardware switches to monitor mode



- **Privileged instructions**
 - Executed only in **monitor mode**
 - Requested by users (system calls)

I/O Protection

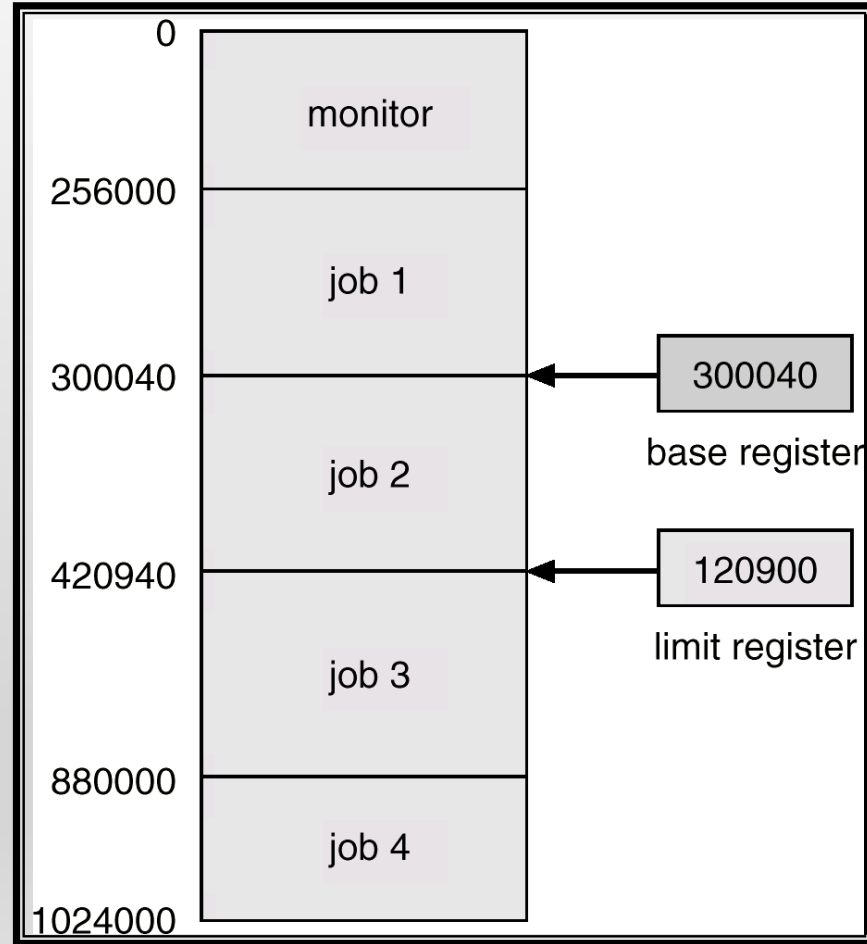
- All I/O instructions are privileged instructions
 - Any I/O device is shared between users
- Must ensure that a user program could never gain control of the computer in monitor mode (i.e., a user program that, as part of its execution, stores a **new address** in the interrupt vector)



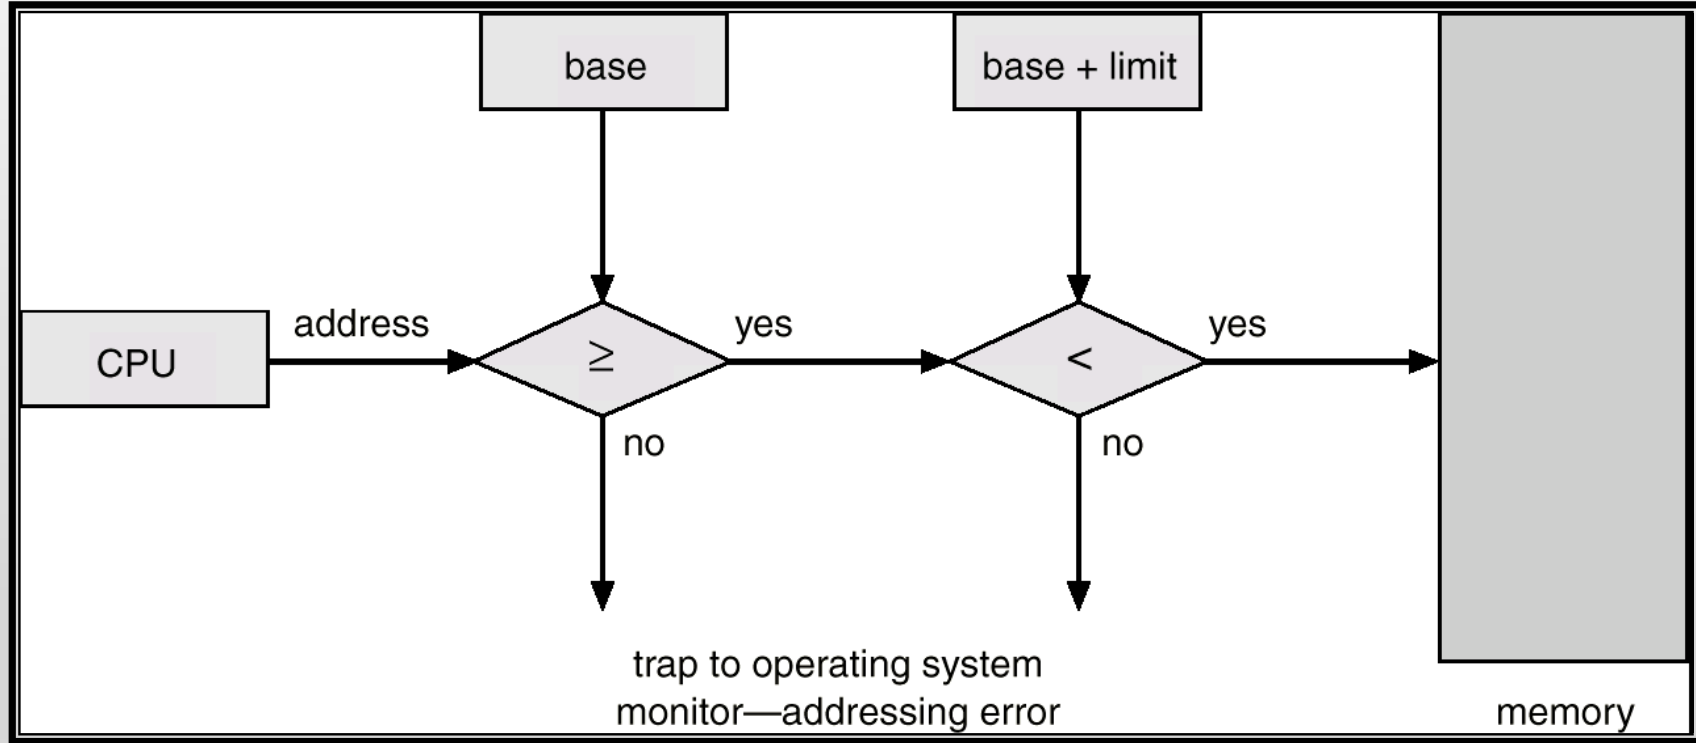
Memory Protection

- Protect
 - Interrupt vector and the interrupt service routines
 - Data access and over-write from other programs
- HW support: two registers for legal address determination:
 - **Base register** – holds the smallest legal physical memory address
 - **Limit register** – contains the size of the range
- Memory outside the defined range is protected

Use of Base and Limit Register



Hardware Address Protection



CPU Protection

- Prevent user program from not returning control
 - getting stuck in an infinite loop
 - not calling system services
- HW support: **Timer**—interrupts computer after specified period
 - Timer is decremented every clock tick
 - When timer reaches the value 0, an interrupt occurs
- Timer commonly used to implement **time sharing**
- **Load-timer** is a privileged instruction

Review Slides (4)

- Dual-mode Operation?
- CPU protection?
- Memory protection?

Reading Material & HW

- Chap 1
- Problem set
 - 1.8: What is the purpose of interrupt? How does an interrupt differ from trap? Can traps be generated intentionally by a user program? If so, for what purpose?
 - 1.10: some computer systems do not provide a privileged mode of operation in hardware. Is it possible to construct a secure operating system for these computer systems? Give arguments both that it is and that it is not possible.
 - Why dual mode operation can protect the system?

Q & A

Thank you for your attention