



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

Operating System

Lecture 03: Processes Concept

Shuo-Han Chen 陳碩漢

shch@nycu.edu.tw

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

Course Schedule

W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Lec07: Memory Management	V	
6	Oct. 9	Lec08: Virtual Memory Management	V	HW02
7	Oct. 16	Lec04: Process Scheduling	V	
8	Oct. 23	School Midterm Exam		
9	Oct. 30	Lec05: Process Synchronization	V	
10	Nov. 6	Lec06: Deadlocks	V	HW03
11	Nov. 13	Lec09: File System Interface	V	
12	Nov. 20	School Event – No class		
13	Nov. 27	Lec10: File System Implementation	V	HW04
14	Dec. 4	Lec11: Mass Storage System	V	
15	Dec. 11	Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

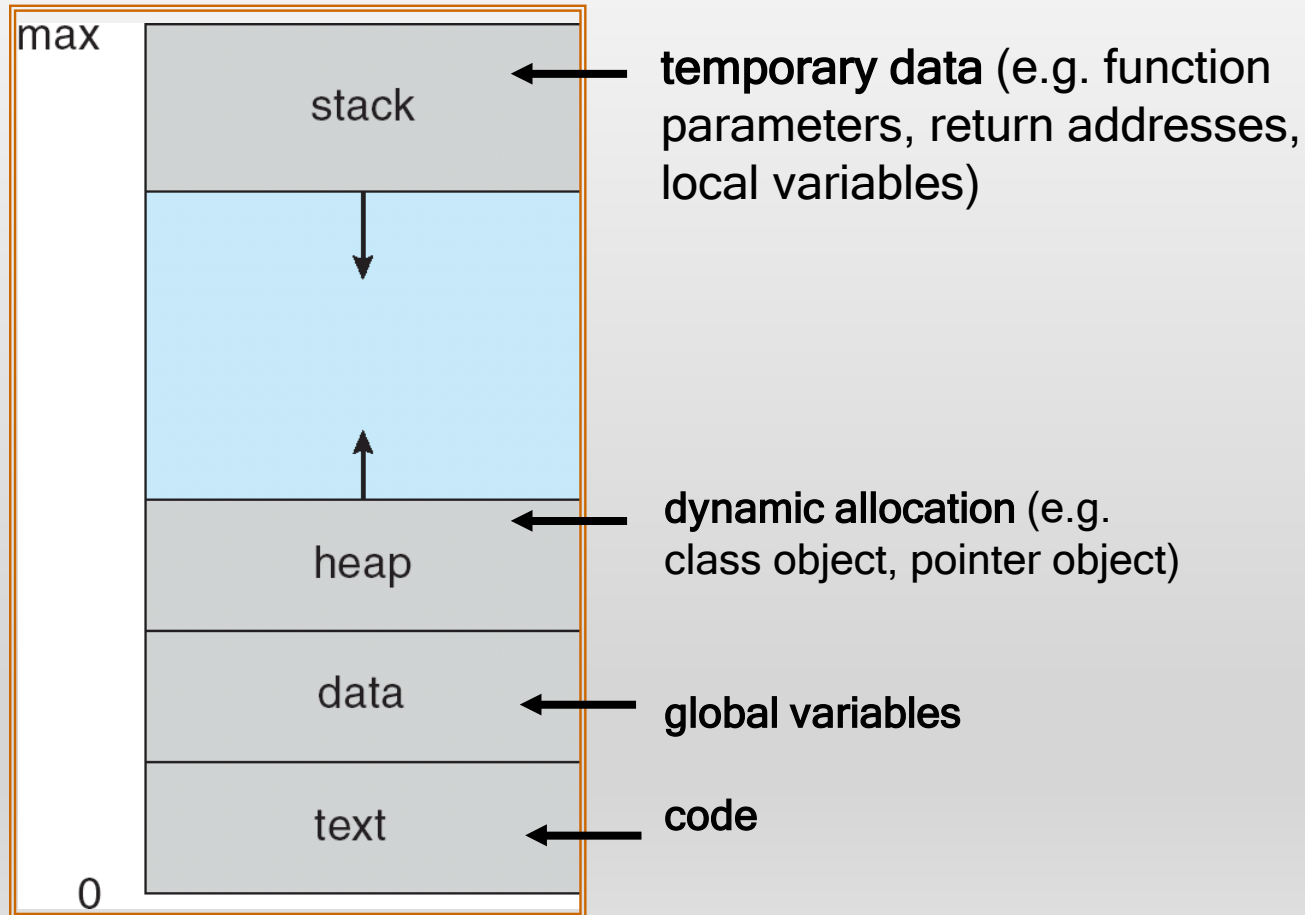
Outline

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication

Process Concept


- An operating system concurrently executes a variety of programs (e.g. Web browser, text editor, etc)
 - **Program** – passive entity: **binary stored in disk**
 - **Process** – active entity: **a program in execution in memory**
- A process includes:
 - **Code** segment (**text section**)
 - **Data section** – global variables
 - **Stack** – temporary local variables and functions
 - **Heap** – dynamic allocated variables or classes
 - Current activity (**program counter**, register contents)
 - A set of associated **resources** (e.g. open file handlers)

Process in Memory



Threads

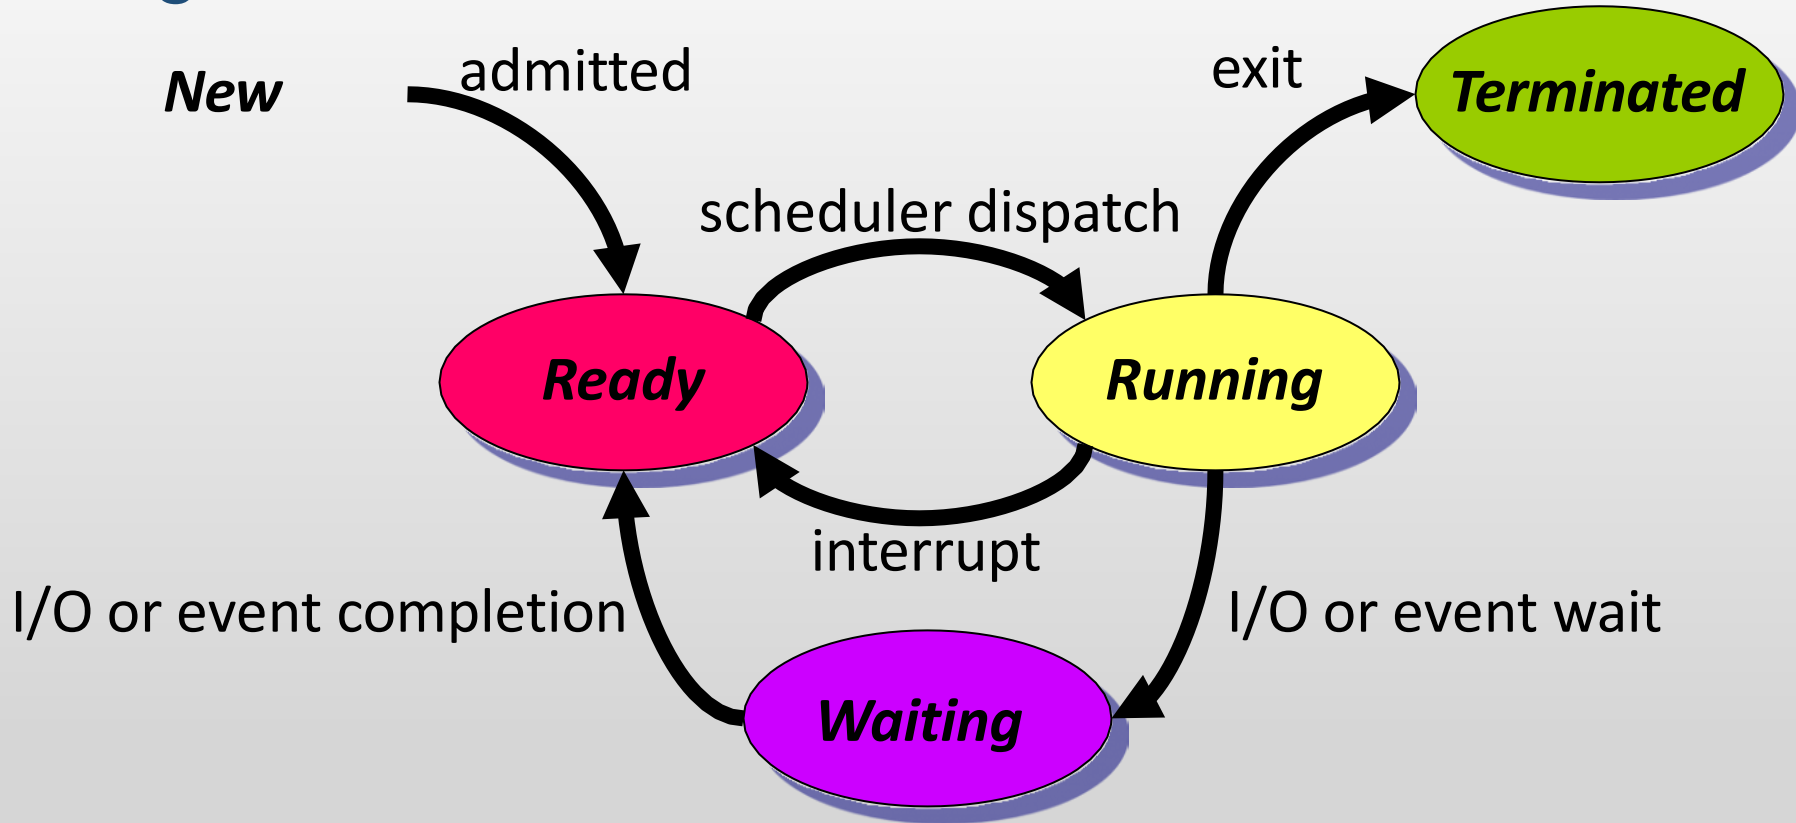
- A.k.a **lightweight process**: basic unit of CPU utilization
- All threads **belonging to the same process** share
 - **code** section, **data** section, and **OS resources** (e.g. open files and signals)
- But each thread has its own
 - **thread ID**, **program counter**, **register set**, and a **stack**



Process State

- States
 - **New**: the process is being created
 - **Ready**: the process is in the memory waiting to be assigned to a processor
 - **Running**: instructions are being executed by CPU
 - **Waiting**: the process is waiting for events to occur
 - **Terminated**: the process has finished execution

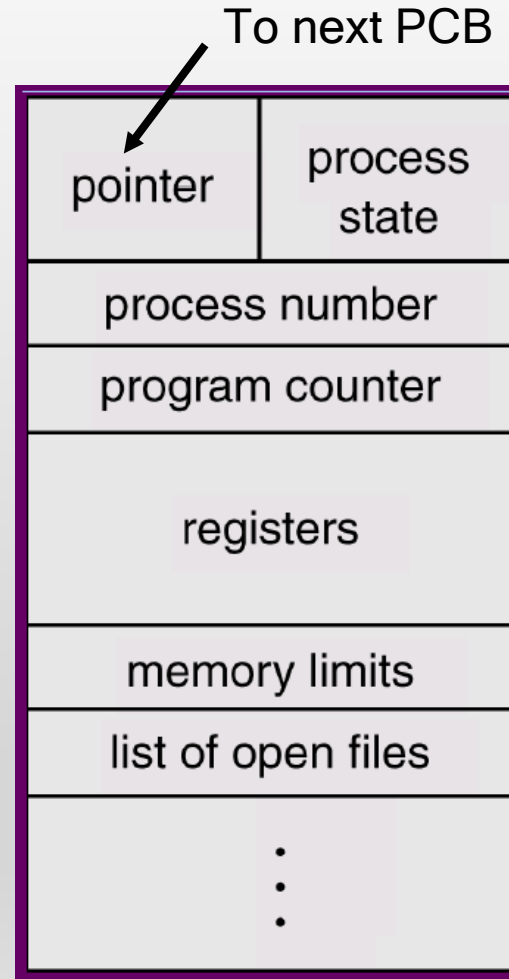
Diagram of Process State



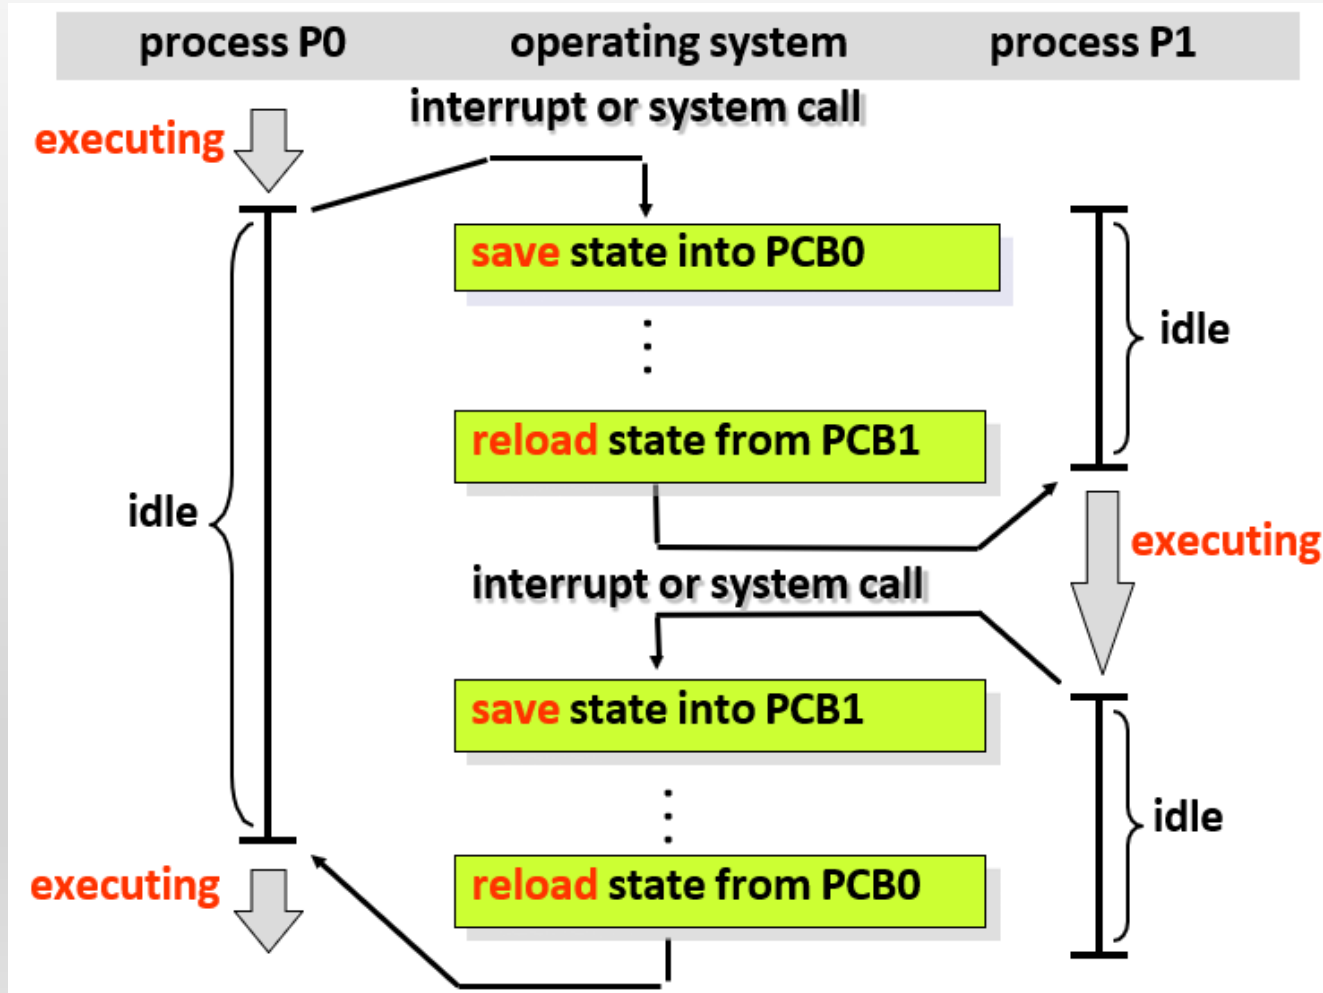
- Only one process is running on any processor at any instant
- However, many processes may be ready or waiting

Process Control Block (PCB)

- Info. associated with each process
- Process state
- Program counter
- CPU registers
- CPU scheduling information (e.g. priority)
- Memory-management information (e.g. base/limit register)
- I/O status information
- Accounting information



Context Switch



Context Switch

- **Context Switch**: Kernel saves the state of the old process and loads the saved state for the new process
- Context-switch time is purely **overhead**
- Switch time (about 1~1000 ms) depends on
 - memory speed
 - number of registers
 - existence of special instructions
 - a single instruction to save/load all registers
 - hardware support
 - **multiple sets of registers** (Sun UltraSPARC - a context switch means changing register file pointer)

Review Slides (1)

- What's the definition of a process?
- What's the difference between process and thread?
- What's PCB? its contents?
 - Process state
 - Program counter
 - CPU registers
- The kinds of process state?
 - New, Ready, Running, Waiting, Terminated
- What's context switch?

Outline

- Process Concept
- **Process Scheduling**
- Operations on Processes
- Interprocess Communication

Process Scheduling

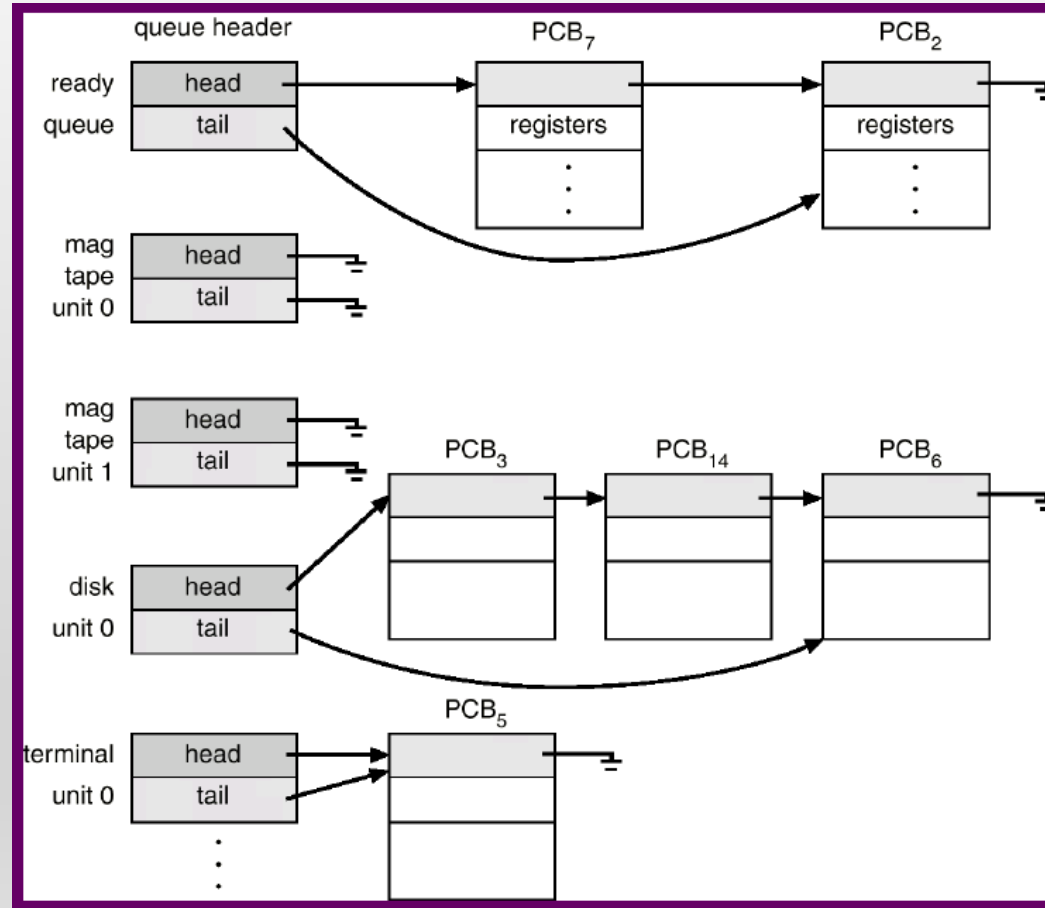
- **Multiprogramming**: CPU runs process at all times to **maximize CPU utilization**
- **Time sharing**: switch CPU frequently such that **users** can **interact** with each program while it is running
- Processes will have to wait until the CPU is free and can be re-scheduled

Process Scheduling Queues

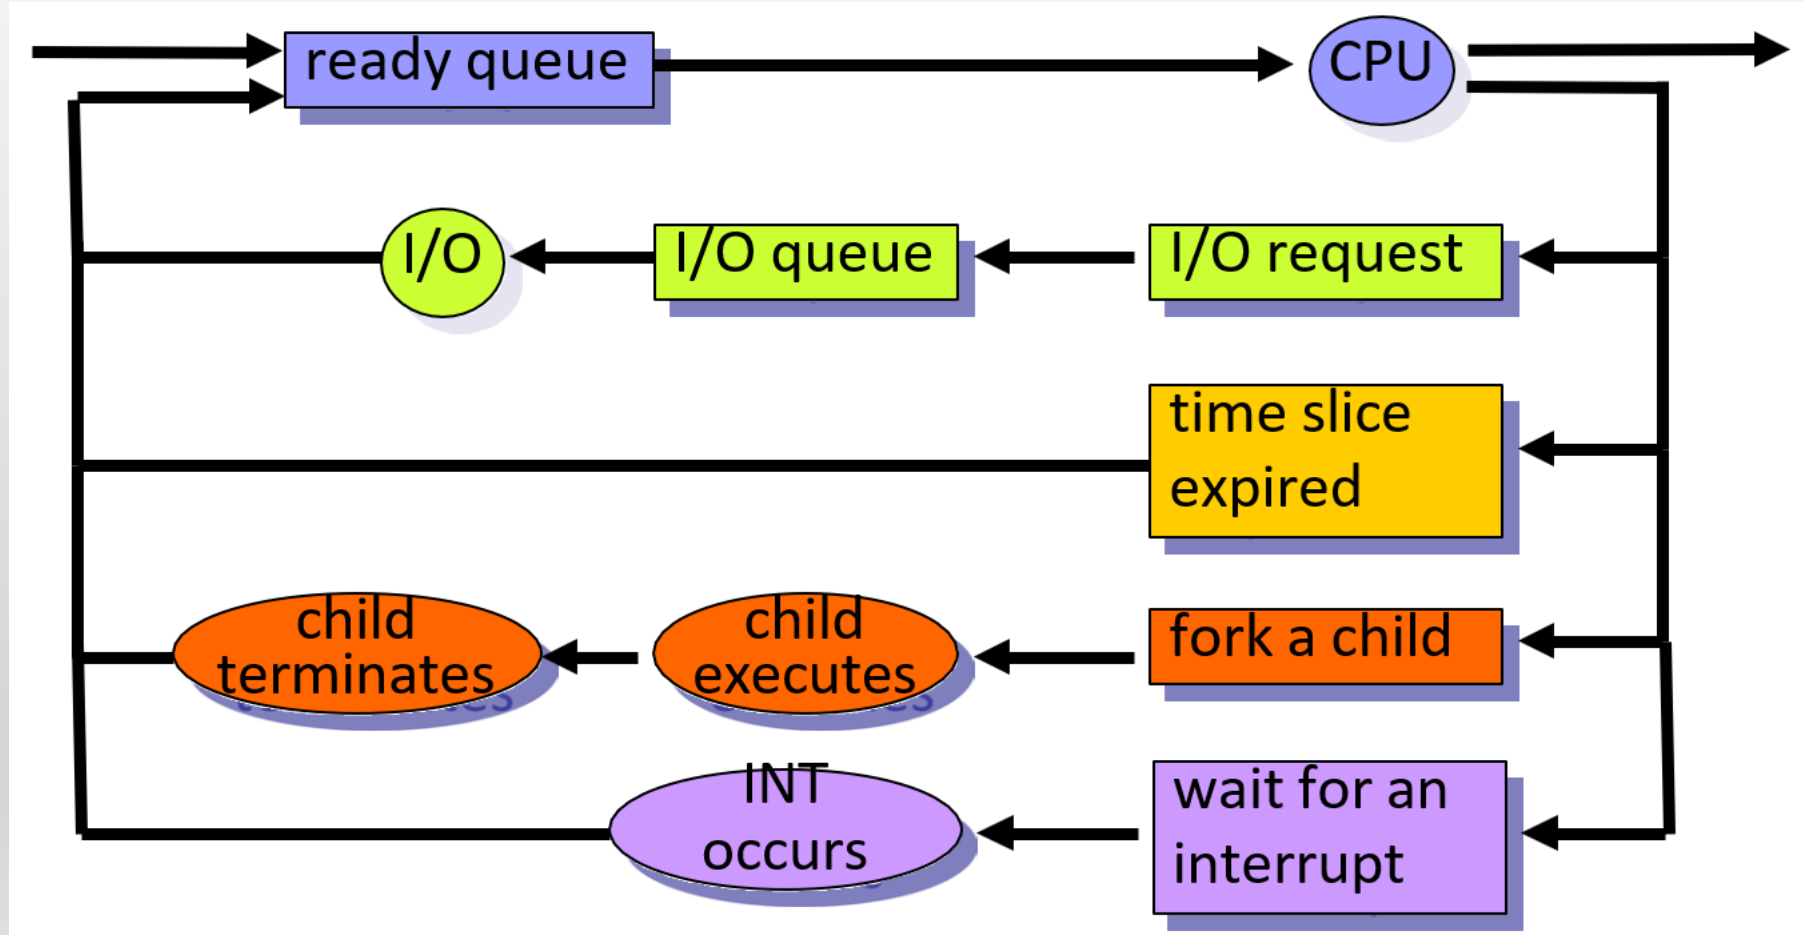
- Processes migrate between the various queues (i.e. switch among states)
- **Job queue** (New State) - set of **all processes** in the system
- **Ready queue** (Ready State) - set of all processes residing in main memory, **ready and waiting to execute**
- **Device queue** (Wait State)- set of processes **waiting for an I/O device**

Process Scheduling Queues

I/O queue

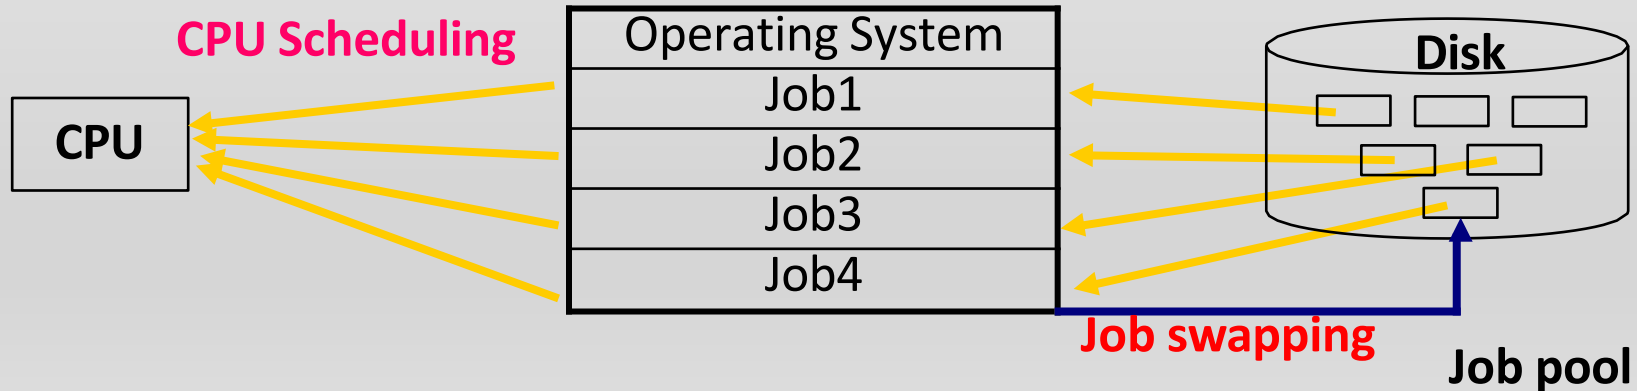


Process Scheduling Diagram



Schedulers

- **Short-term** scheduler (**CPU scheduler**)- selects which process should be executed and **allocated CPU** (Ready state -> Run state)
- **Long-term** scheduler (**job scheduler**) - selects which processes should be **loaded into memory** and brought into the ready queue (**New state -> Ready state**)
- **Medium-term** scheduler - selects which processes should be **swapped in/out memory** (Ready state -> Wait state)

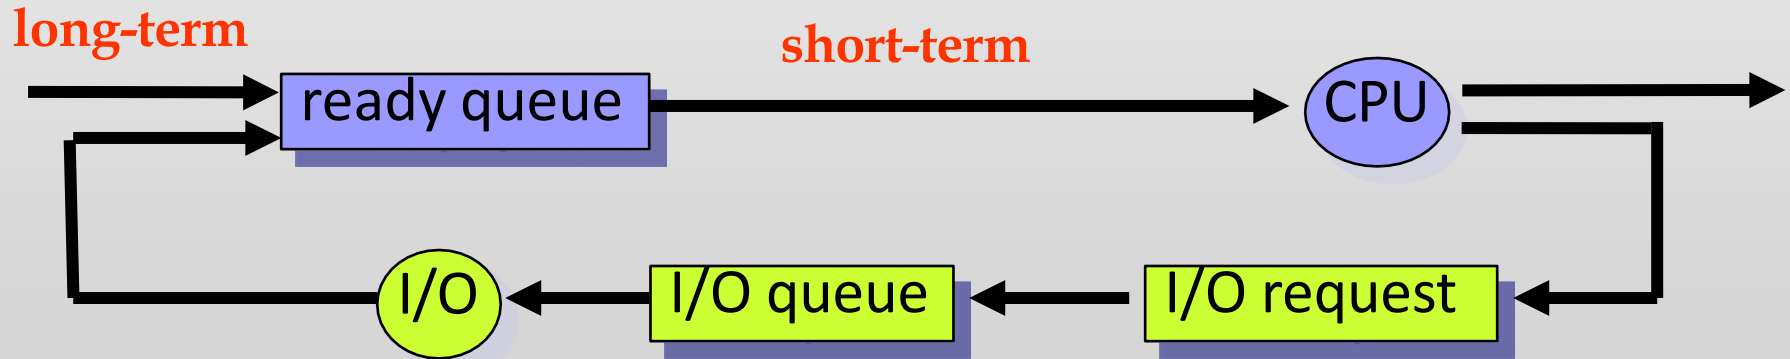


Long-Term Scheduler

- Control **degree of multiprogramming**
- Execute less frequently (e.g. invoked only when a process leaves the system or **once several minutes**)
- Select a **good mix of CPU-bound & I/O-bound** processes to increase system overall performance
- UNIX/NT: no long-term scheduler
 - Created process placed in memory for short-term scheduler
 - Multiprogramming degree is bounded by hardware limitation (e.g., # of terminals) or on the self-adjusting nature of users

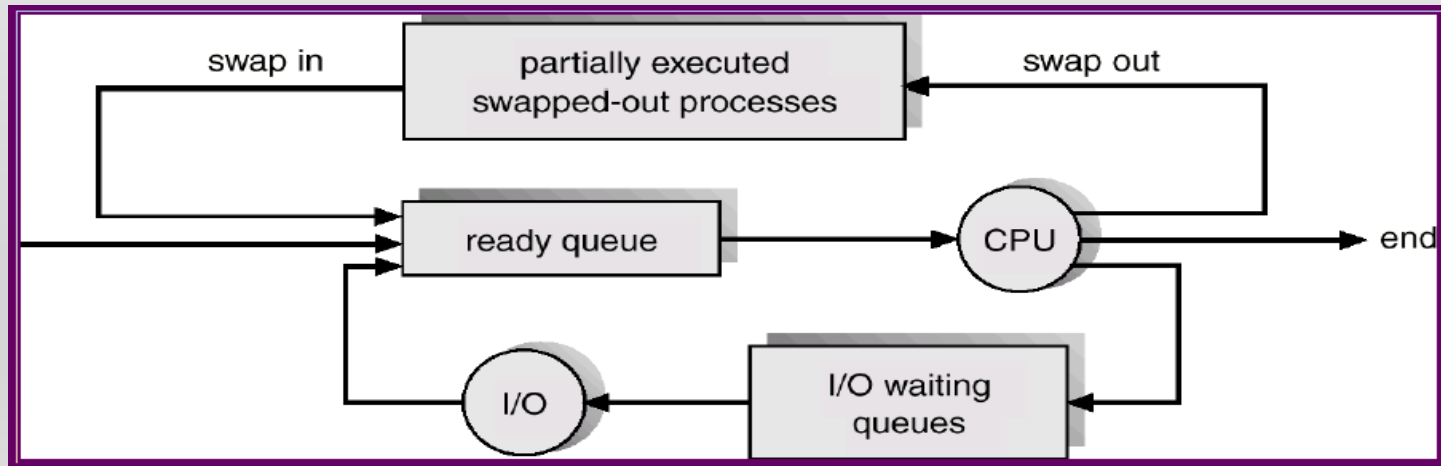
Short-Term Scheduler

- Execute quite frequently (e.g. **once per 100ms**)
- Must be efficient:
 - if 10 ms for picking a job, 100 ms for such a pick,
 - $\rightarrow \text{overhead} = 10 / 110 = 9\%$



Medium-Term Scheduler

- **swap out**: removing processes from memory to reduce the degree of multiprogramming
- **swap in**: reintroducing swap-out processes into memory
- Purpose: improve process mix ,_free up memory_
- Most modern OS doesn't have medium-term scheduler because having sufficient physical memory or using virtual memory

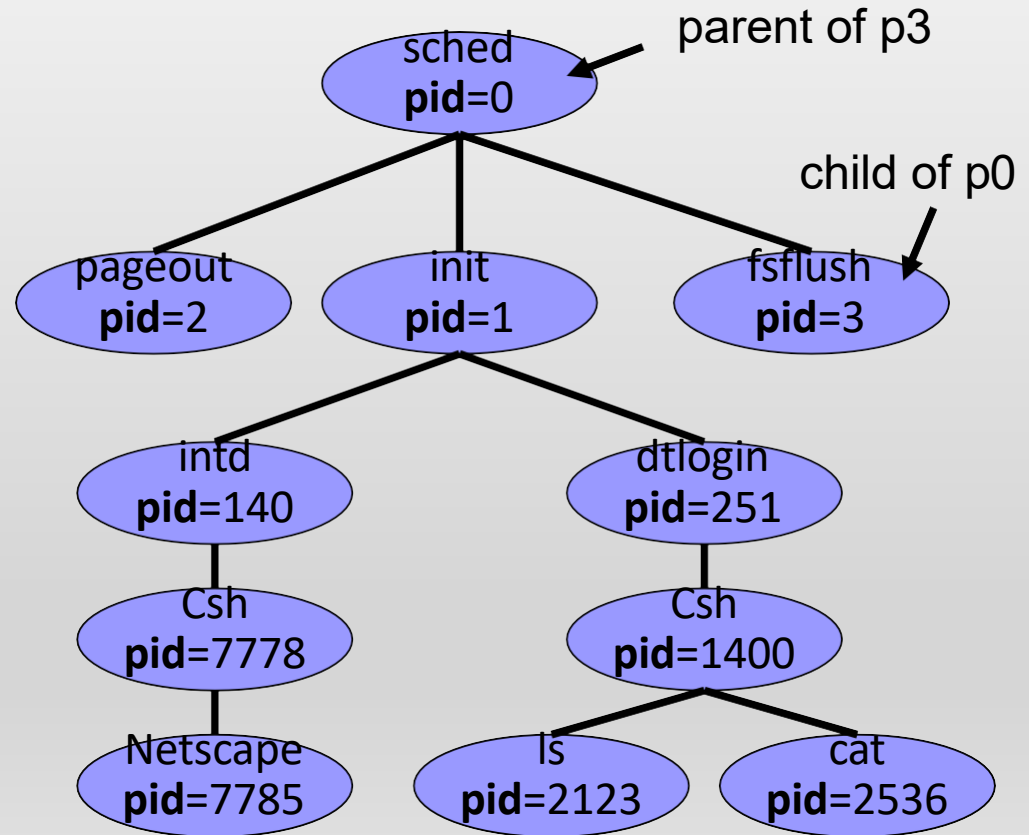


Outline

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication

Tree of Processes

- Each process is identified by a **unique** processor identifier (**pid**)



UNIX: “**ps -ael**” will list complete info of all active processes

Process Creation


- Resource sharing
 - Parent and child processes share **all** resources
 - Child process shares **subset** of parent's resources
 - Parent and child share **no** resources
- Two possibilities of execution
 - Parent and children **execute concurrently**
 - Parent **waits until children terminate**
- Two possibilities of address space
 - **Child duplicate of parent**, communication via sharing variables
 - **Child has a program loaded into it**, communication via message passing

UNIX/Linux Process Creation


- **fork** system call
 - Create a new (child) process
 - The new process **duplicates the address space** of its parent
 - Child & Parent **execute concurrently** after fork
 - Child: return value of fork is 0
 - Parent: return value of fork is PID of the child process
- **execvp** system call
 - **Load a new binary file** into memory - **destroying the old code**
- **wait** system call
 - The parent waits for **one of its child processes** to complete

UNIX/Linux Process Creation


- Memory space of fork():
 - Old implementation: A's child is an **exact copy** of parent
 - Current implementation: use **copy-on-write** technique to store **differences in A's child address space**



Originally



After A does
an **fork**



After the child
does an **execvp**

UNIX/Linux Example

```
#include <stdio.h> void main( )  
{  
    int A;  
    /* fork another process */  
    A = fork( );  
  
    if (A == 0) { /* child process */  
        printf("this is from child  
process\n"); execlp("/bin/ls", "ls",  
NULL);  
    } else { /* parent process */  
        printf("this is from parent  
process\n");  
        int pid = wait(&status);  
        printf("Child %d completes", pid);  
    }  
}
```

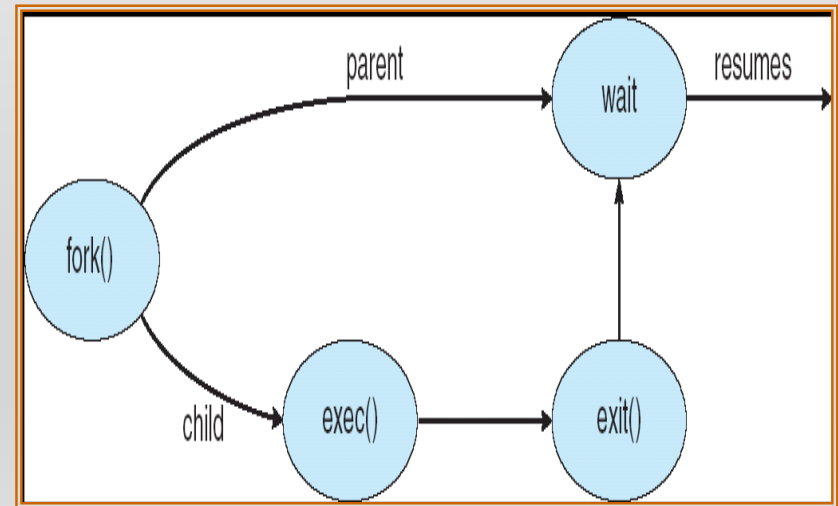
Output:

this is from child process

this is from parent process

a.out hello.c readme.txt


Child 32185 completes
process ends 32185



Example Quiz:

- How many processes are created?

```
#include <stdio.h>
#include <unistd.h>
int main() {
    for (int i=0; i<3; i++){
        fork();
    }
    return 0;
}
```



Process Termination

- Terminate when the last statement is executed or `exit()` is called
 - All resources of the process, including physical & virtual memory, open files, I/O buffers, are **deallocated by the OS**
- Parent may terminate execution of children processes by specifying its PID (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
- Cascading termination:
 - killing (exiting) parent -> killing (exiting) all its children

Review Slides (2)

- What's long-term scheduler? features?
- What's short-term scheduler? features?
- What's medium-term scheduler? features?
- What's the different between duplicate address space and load program?
Their commands?

Outline

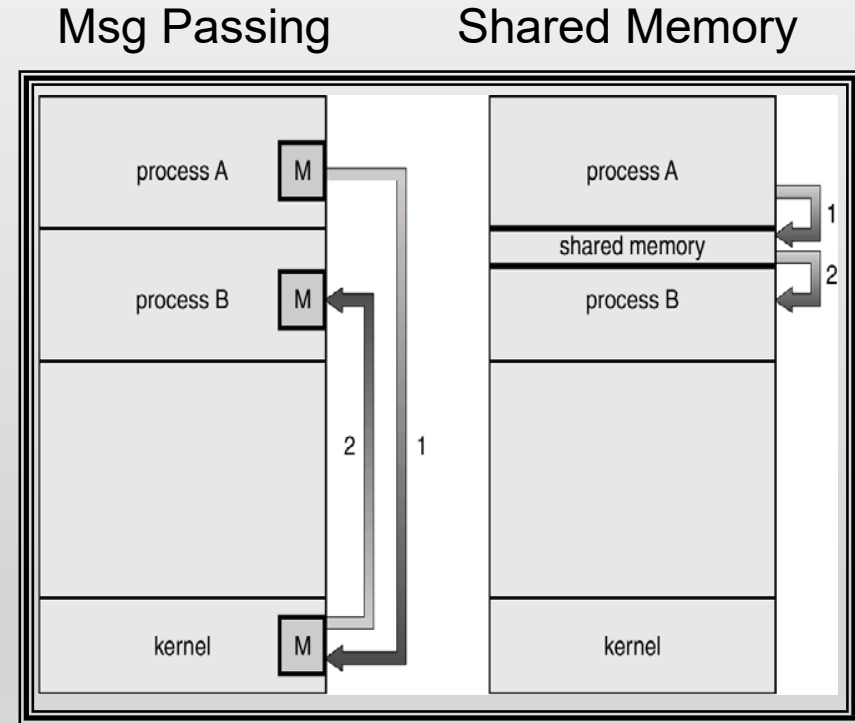
- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication

Interprocess Communication

- **IPC**: a set of methods for the exchange of data among multiple threads in one or more processes
- **Independent process**: cannot affect or be affected by other processes
- **Cooperating process**: otherwise
- Purposes
 - information sharing
 - computation speedup (not always true...)
 - convenience (performs several tasks at one time)
 - modularity

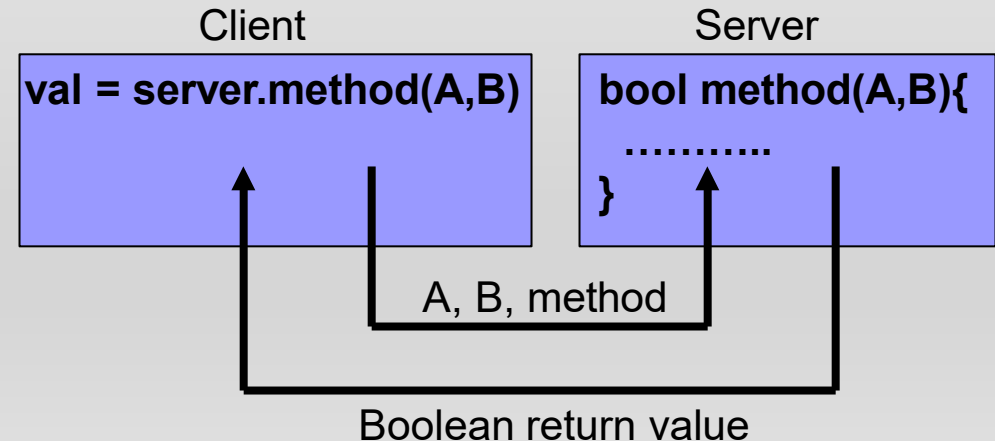

Communication Methods

- Shared memory:
 - Require more careful **user synchronization**
 - Implemented by memory access: faster speed
 - **Use memory address to access data**
- Message passing:
 - No conflict: **more efficient for small data**
 - **Use send/rcv message**
 - Implemented by **system call**: slower speed



Communication Methods

- Sockets:
 - A network connection identified by **IP & port**
 - Exchange **unstructured stream of bytes**
- Remote Procedure Calls:
 - Cause a **procedure** to execute in another address space
 - Parameters and return values are passed by message



Interprocess Communication


- Shared Memory
- Message Passing
- Socket
- Remote Procedure Calls

Shared Memory

- Processes are responsible for...
 - Establishing a region of shared memory
 - Typically, a shared-memory region resides in the address space of the process creating the shared-memory segment
 - Participating processes **must agree to remove memory access constraint** from OS
 - Determining the form of the data and the location
 - Ensuring data are not written simultaneously by processes

Consumer & Producer Problem

- **Producer** process produces information that is consumed by a **Consumer** process
- Buffer as a circular array with size B
 - next free: **in**
 - first available: **out**
 - empty: **in = out**
 - full: **$(in+1) \% B = out$**
- The solution allows at most (B-1) item in the buffer
 - Otherwise, cannot tell the buffer is full or empty



Shared-Memory Solution

```
/*producer*/
```

```
while (1) {
```

```
    while (((in + 1) % BUFFER_SIZE) == out)
```

```
        ; //wait if buffer is full
```

```
    buffer[in] = nextProduced;
```

```
    in = (in + 1) % BUFFER_SIZE;
```

```
}
```

“in” only modified by producer

```
/*consumer*/
```

```
while (1) {
```


```
    while (in == out); //wait if buffer is empty
```

```
    nextConsumed = buffer[out];
```


```
    out = (out + 1) % BUFFER_SIZE;
```

```
}
```

“out” only modified by consumer



```
/* global data structure */  
#define BUFSIZE 10  
item buffer[BUFSIZE];  
int in = out = 0;
```



Interprocess Communication

- Shared Memory
- Message Passing
- Socket
- Remote Procedure Calls

Message-Passing System

- Mechanism for processes to **communicate** and **synchronize** their actions
- IPC facility provides two operations:
 - **Send**(message) - message size fixed or variable
 - **Receive**(message)
- Message system - processes communicate **without resorting to shared variables**
- To communicate, processes need to
 - Establish a **communication link**
 - Exchange a message via **send/receive**

Message-Passing System

- Implementation of communication link
 - physical (e.g., shared memory, HW bus, or network)
 - logical (e.g., **logical properties**)
 - **Direct or indirect communication**
 - Symmetric or asymmetric communication
 - **Blocking or non-blocking**
 - Automatic or explicit buffering
 - Send by copy or send by reference
 - Fixed-sized or variable-sized messages

Direct communication

- Processes must **name each other explicitly**:
 - **Send (P, message)** - send a message to proc P
 - **Receive (Q, message)** - receive a message from process Q
- Properties of communication link
 - Links are **established automatically**
 - **One-to-One** relationship between links and processes
 - The link may be unidirectional, but is usually bi- directional

Direct communication

- Solution for producer-consumer problem:

```
/*producer*/  
while (1) {  
    send (consumer, nextProduced);  
}  
/*consumer*/  
while (1) {  
    receive (producer, nextConsumed);  
}
```


- **limited modularity**: if the **name** of a process is **changed**, all old names should be found

Indirect communication

- Messages are directed and received from **mailboxes** (also referred to as ports)
 - **Each mailbox has a unique ID**
 - Processes can communicate if they share a mailbox
 - **Send (A, message)** - send a message to mailbox A
 - **Receive (A, message)** - receive a message from mailbox A
- Properties of communication link
 - Link established only **if processes share a common mailbox**
 - **Many-to-Many** relationship between links and processes
 - Link may be unidirectional or bi-directional
 - Mailbox can be owned either by OS or processes

Indirect Communication


- Mailbox sharing



- Solutions
- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily a single receiver. Sender is notified who the receiver was

Synchronization

- Message passing may be either **blocking** (synchronous) or **non-blocking** (asynchronous)
 - **Blocking send**: sender is blocked until the message is received by receiver or by the mailbox
 - **Nonblocking send**: sender sends the message and resumes operation
 - **Blocking receive**: receiver is blocked until the message is available
 - **Nonblocking receive**: receiver receives a valid message or a null
- Buffer implementation
 - **Zero** capacity: blocking send/receive
 - **Bounded** capacity: if full, sender will be blocked
 - **Unbounded** capacity: sender never blocks




Interprocess Communication

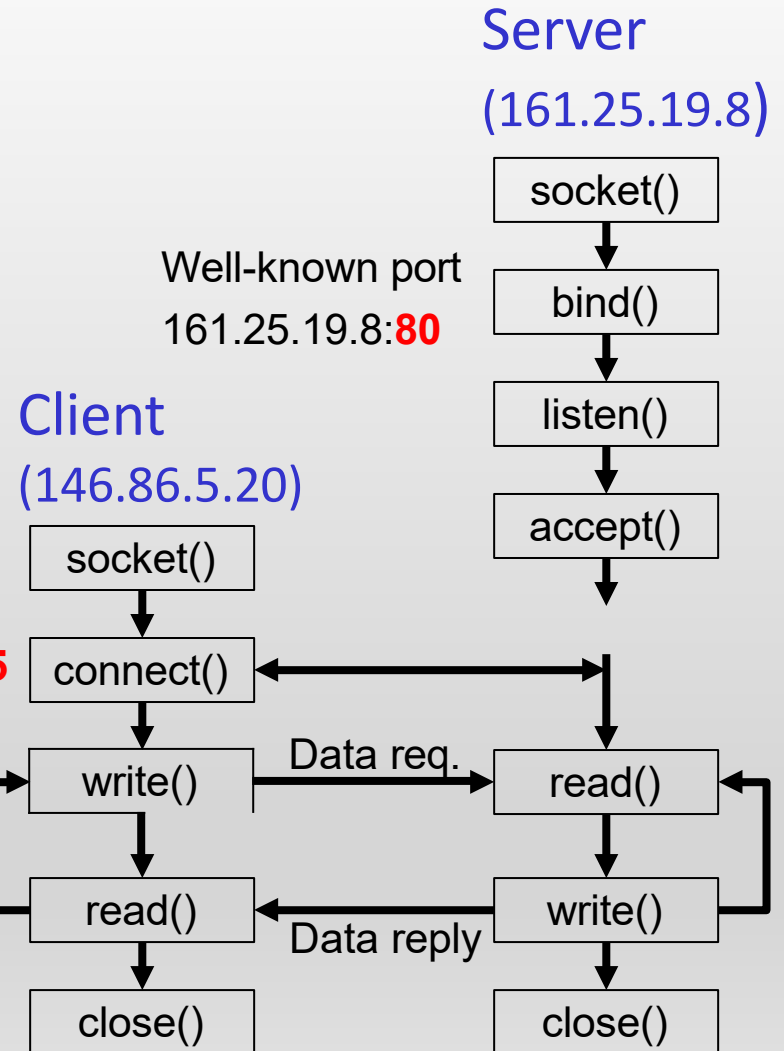
- Shared Memory
- Message Passing
- Socket
- Remote Procedure Calls

Sockets

- A socket is identified by a concatenation of **IP address** and **port number**
- Communication consists between a **pair of sockets**
- Use 127.0.0.1 to **refer itself**




Assign port
146.86.5.20:1625



Sockets

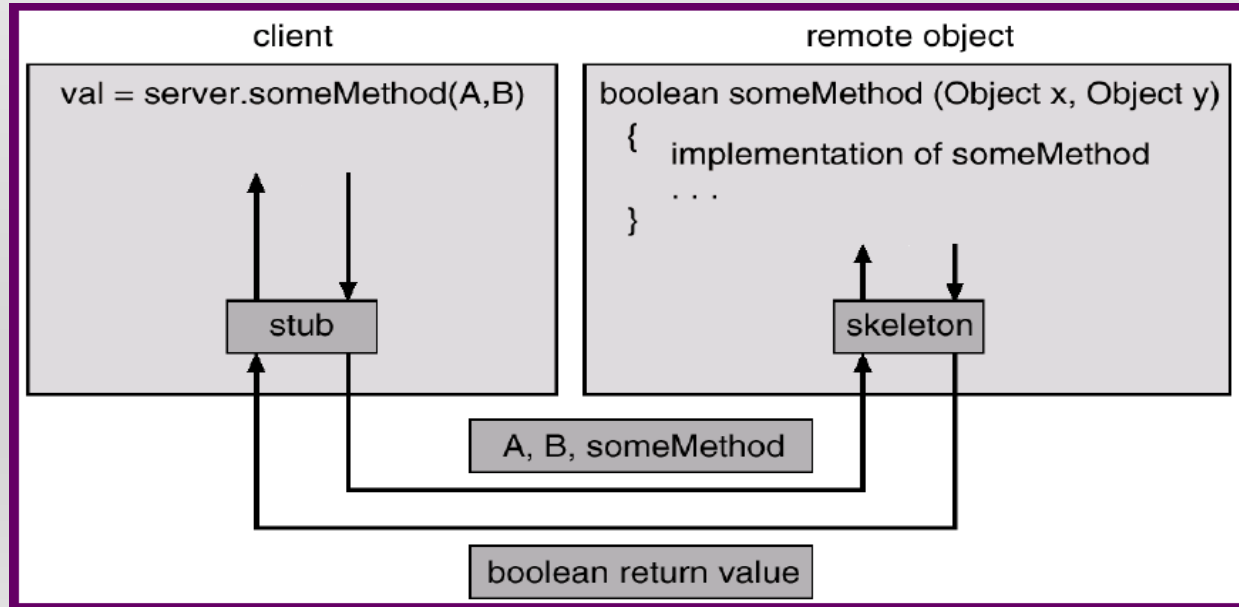
- Considered as a low-level form of communication **unstructured stream of bytes** to be exchanged
- **Data parsing** responsibility falls **upon the** server and the client applications

HTTP example:



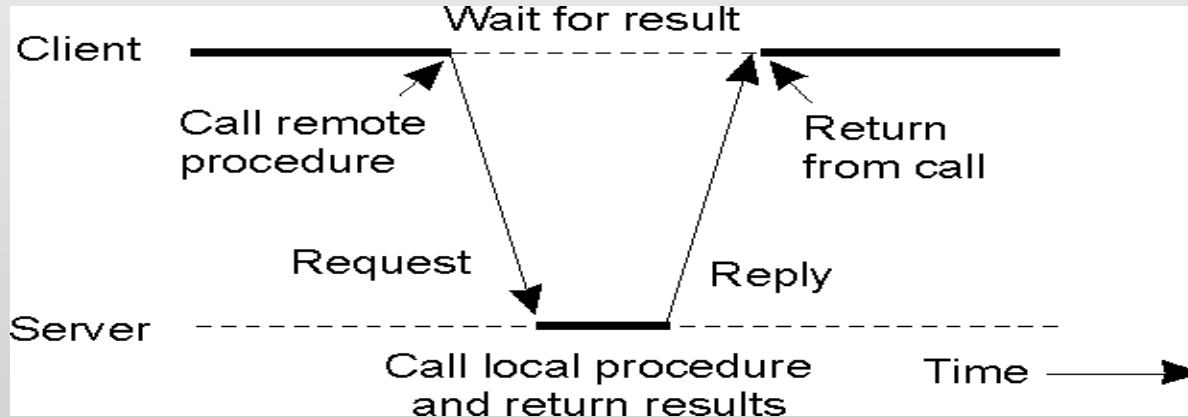
Remote Procedure Calls: RPC

- Remote procedure call (RPC) abstracts procedure calls between processes **on networked systems**
 - allows programs to call procedures located on other machines (and other processes)
- **Stubs** - client-side proxy for the actual procedure on the server



Client and Server Stubs

- Client stub:
 - Packs parameters into a message (i.e. **parameter marshaling**)
 - Calls OS to send directly to the server
 - Waits for result-return from the server



Server stub:

- Receives a call from a client
- Unpacks the parameters
- Calls the corresponding procedure
- Returns results to the caller

Review Slides (3)

- Shared memory vs. Message-passing system?
- Direct vs. Indirect message-passing system?
- Blocking vs. Non-Blocking?
- Socket vs. RPC?

Reading Material & HW

- Chap 3
- HW (Problem set)
 - 3.1
 - 3.2
 - 3.5
 - 3.7
 - 3.10

Example: POSIX Shared Memory

```
{  
    /* allocate a R/W shared memory segment */  
    char* segment_id = shmget(IPC_PRIVATE, 4096, S_IRUSR | S_IWUSR);  
    /* attach the shared memory segment */  
    char* shared_memory = (char*) shmat(segment_id, NULL, 0);  
    /* write a message to the shared memory segment */  
    sprintf(shared_memory, "Write to shared memory");  
    /* print out the string from the shared memory segment */  
    printf("%s\n", shared_memory);  
    /* detach the shared memory segment */  
    shmdt(shared_memory);  
    /* remove the shared memory segment */  
    shmctl(shared_memory, IPC_RMID, NULL);  
}
```

Diagram illustrating the parameters of the `shmget` and `shmat` functions:


- `shmget`:
 - `IPC_PRIVATE`: Mem.location
 - `4096`: size
 - `S_IRUSR | S_IWUSR`: R/W mode
- `shmat`:
 - `segment_id`: Mem.location
 - `NULL`: Mem.location
 - `0`: R/W mode

Example: Mach Message Passing

- Mach operating system
 - developed at CMU
 - **microkernel** design
 - most communications are carried out by **messages** and **mailboxes** (aka **ports**)
 - Problem: performance (data coping)
- When each task (process) is created
 - **kernel** & **notify** mailboxes also created
 - kernel mailbox: channel between OS & task
 - notify mailbox: OS sends event notification to

Mach Mailbox

- port-allocate: system call to create a mailbox
 - default buffer size: 8 messages
 - **FIFO** queueing
 - a message: one fixed-size **header** + variable-length **data** portion
 - implementing both blocking- & non-blocking send/receive



RPC Problems

- Data representations -> integer, floating?
- Different address spaces -> pointer?
- Communication error -> duplicate or missing calls

RPC Problems: Data Representation Issue

- Problem
 - IBM mainframes use EBCDIC char code and IBM PC uses ASCII code
 - Integer: one's complement and 2's complement
 - Floating-point numbers
 - Little endian and big endian
- Solution
 - **External data representation (XDR)**

RPC Problems: Address Space Issue

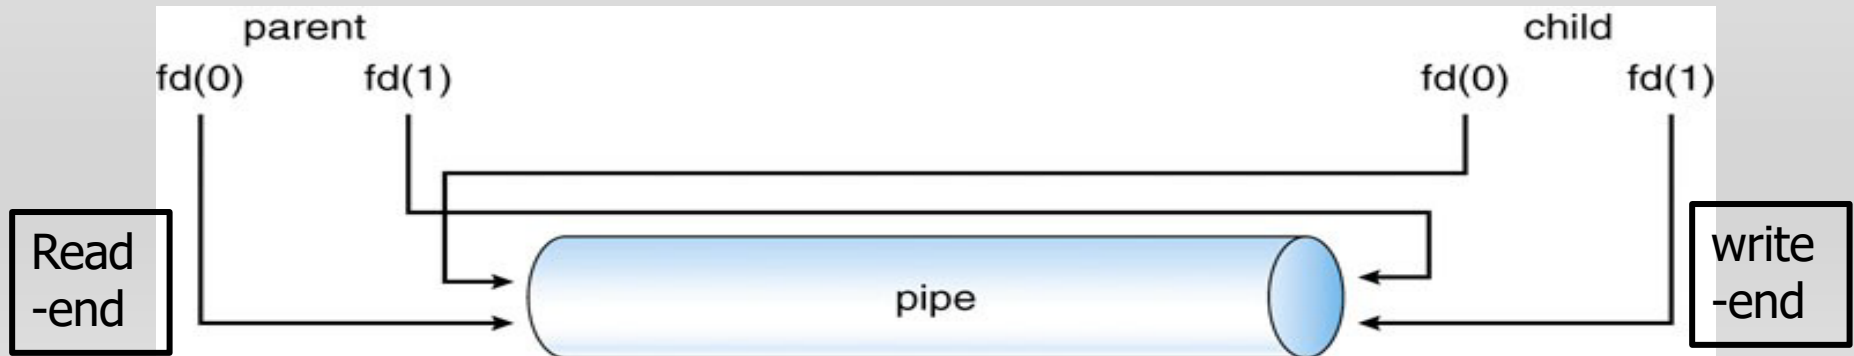
- A pointer is only meaningful in its address space
- Solutions
 - No pointer usage in RPC calls
 - Copy the entire pointed area (such as arrays or strings)
 - Only suitable for bounded and known areas

RPC Problems: Communication Issue

- RPCs may fail, or be duplicated and execute more than once, as a result of common **network errors**
- **at most once**: prevent duplicate calls
 - Implemented by attaching a **timestamp** to each message
 - The server must keep a history large enough to ensure that repeated messages are detected
- **exact once**: prevent missing calls
 - The server must **acknowledge to the client** that the RPC call was received and executed
 - The client must **resend each RPC call periodically** until the server receives the ACK

Pipes

- One of the 1st IPC mechanism in early UNIX systems
- Pipe is a special type of **file**
- Issues in implementing
 - uni- or bi-directional?
 - half or full duplex? (travel in both directions simultaneously)
 - Must a relationship (parent - child) exist?
 - Over a network, or reside on the same machine?



Ordinary Pipes

- Also called **anonymous pipes** in Windows
- Requires a **parent-child relationship** between the communicating processes
 - Implemented as a **special file** on Unix (via `fork()`, a **child process inherits open files from its parent**)
 - Can only be used between processes on the same machine
- **Unidirectional**: two pipes must be used for two-way communication

UNIX:

```
int fd[2];  
pipe(fd);
```

Windows:

```
CreatePipe(&ReadHandle, &WriteHandle, &sa, 0)
```

Named Pipes


- No parent-child relationship is required
- Several processes can use it for communications
 - It may have several writers
- Continue to exist after communicating processes exit
- In Unix:
 - Also called **FIFO**
 - Communicating processes have to be on the **same machine**
- In Windows:
 - bi-directional
 - Communicating processes can be on **different machine**

UNIX/Linux: Fork

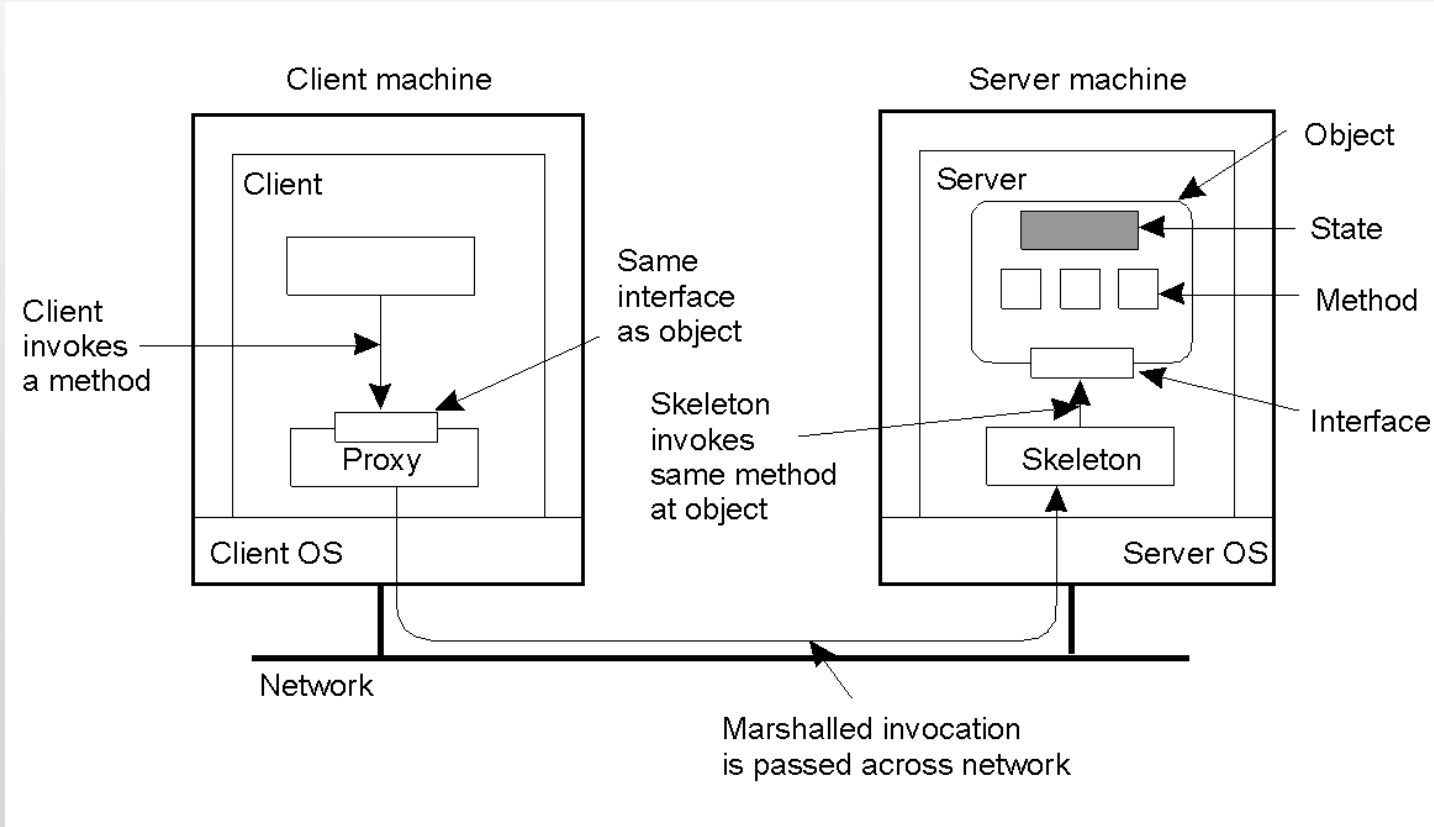
- Inherited from the parent:
 - process credentials
 - environment
 - **stack**
 - **memory**
 - **open file descriptors**
 - **signal handling settings**
 - scheduler class
 - process group ID
 - session ID
 - current working directory
 - root directory
 - file mode creation mask (umask)
 - resource limits
 - controlling terminal
- Unique to the child:
 - **process ID**
 - **different parent process ID**
 - Own copy of file descriptors and directory streams.
 - process, text, data and other memory locks are NOT inherited.
 - process times, in the tms struct
 - resource utilizations are set to 0
 - **pending signals** initialized to the empty set
 - timers created by timer_create not inherited
 - **asynchronous input or output operations** not inherited

Remote Method Invocation

- RMI is a Java mechanism similar to RPC
- RMI allows a Java program on one machine to invoke a **method** on a **remote object** instead of a function



Distributed Objects



A remote object with client-side proxy

Static & Dynamic RMI

- RMI = Remote Method Invocation
 - Invoke an object's method through proxy
- Static invocation
 - `objectname.methodname(para)`
 - If interfaces change, apps must be recompiled
- Dynamic invocation
 - `invoke(object, method, inpars, outpars)`

Q & A

Thank you for your attention