



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

*Operating System*

# Lecture 06-1: Deadlocks

Shuo-Han Chen 陳碩漢

[shch@nycu.edu.tw](mailto:shch@nycu.edu.tw)

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

# Course Schedule

W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01 Due 10/5
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Typhoon – No class	V	
6	Oct. 9	Lec07: Memory Management	V	
7	Oct. 16	Lec08: Virtual Memory Management	V	HW02 Due 11/2
8	Oct. 23	Lec04: Multithreaded Programming	V	
9	Oct. 30	Midterm Exam		
10	Nov. 6	Lec05: Process Scheduling	V	Let's take a breath
11	Nov. 13	Lec06: Process Synchronization & Deadlocks	X	HW03
12	Nov. 20	School Event – No class	V	
13	Nov. 27	Lec09: File System Interface	V	
14	Dec. 4	Lec10: File System Implementation	V	HW04
15	Dec. 11	Lec11: Mass Storage System & Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

# Overview

- System Model
- Deadlock Characterization
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

# Deadlock Problem

- A set of blocked processes each **holding** some resources and **waiting** to acquire a resource held by another process in the set
- Ex1: 2 processes and 2 tape drivers
  - Each process holds a tape drive
  - Each process requests another tape drive
- Ex2: 2 processes, and semaphores A & B
  - P1 (hold B, wait A): **wait(A)**, signal(B)
  - P2 (hold A, wait B): **wait(B)** , signal(A)

# Necessary Conditions


- Mutual exclusion:
  - only 1 process at a time can use a resource
- Hold & Wait:
  - a process holding some resources and is waiting for another resource
- No preemption:
  - a resource can be only released by a process voluntarily
- Circular wait:
  - there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0 \rightarrow P_1 \rightarrow P_2 \rightarrow \dots \rightarrow P_n \rightarrow P_0$
- -> All four conditions must hold for possible deadlock!

# System Model

- Resources types  $R_1, R_2, \dots, R_m$ 
  - E.g. CPU, memory pages, I/O devices
- Each resource type  $R_i$  has  $W_i$  instances
  - E.g. a computer has 2 CPUs
- Each process utilizes a resource as follows:
  - Request  $\rightarrow$  use  $\rightarrow$  release


# Resource-Allocation Graph

- 3 processes, P1 ~ P3
- 4 resources, R1 ~ R4
  - R1 and R3 each has one instance
  - R2 has **two instances**
  - R4 has **three instances**
- Request edges:
  - *P1 -> R1*: P1 requests R1
- Assignment edges:
  - *R2 -> P1*: One instance of R2 is allocated to P1
- -> P1 is **hold on** an instance of **R2** and **waiting for** an instance of **R1**



# Resource-Allocation Graph w/ Deadlock


- If the graph contains a **cycle**, a deadlock **may exist**
- In the example:
  - P1 is waiting for P2
  - P2 is waiting for P3
  - -> P1 is also waiting for P3
  - Since P3 is waiting for P1 or P2, and they both waiting for P3
  - -> **deadlock!**





# RA Graph w/ Cycle but NO Deadlock

- If the graph contains a **cycle**, a deadlock **may exist**
- In the example:
  - P1 is waiting for P2 or P3
  - P3 is waiting for P1 or P4
  - Since P2 and P4 wait no one
  - -> **no deadlock between P1 & P3!**



# Deadlock Detection

- If graph contains **no cycle** → **no deadlock**
  - **Circular wait** cannot be held
- If graph contains a cycle:
  - if **one instance per resource type** → **deadlock**
  - if **multiple instances per resource type** → **possibility** of deadlock

# Handling Deadlocks

- Ensure the system will **never enter a deadlock state**
  - deadlock **prevention**: ensure that at least one of the **4 necessary conditions** cannot hold
  - deadlock **avoidance**: **dynamically** examines the resource-allocation state before allocation
- Allow to **enter a deadlock state** and **then recover**
  - deadlock **detection**
  - deadlock **recovery**
- **Ignore the problem** and pretend that deadlocks never occur in the system
  - used by most operating systems, including UNIX.

# Review Slides ( I )

- deadlock necessary conditions?
  - mutual exclusion
  - hold & wait
  - no preemption
  - circular wait
- resource-allocation graph?
  - cycle in RAG → deadlock?
- deadlock handling types?
  - deadlock prevention
  - deadlock avoidance
  - deadlock recovery
  - ignore the problem

# Deadlock Prevention & Deadlock Avoidance

# Deadlock Prevention

- Mutual exclusion (ME): do not require ME on sharable resources
  - e.g. there is no need to ensure ME on read-only files
  - Some resources are not shareable, however (e.g. printer)
- Hold & Wait:
  - When a process requests a resource, it does not hold any resource
  - Pre-allocate all resources before executing
  - resource utilization is low; starvation is possible

# Deadlock Prevention (con't)

- No preemption
  - When a process is waiting on a resource, all its holding resources are preempted
    - e.g. P1 request R1, which is allocated to P2, which in turn is waiting on R2. ( $P1 \rightarrow R1 \rightarrow P2 \rightarrow R2$ )
    - R1 can be preempted and reallocated to P1
  - Applied to resources whose states can be easily saved and restored later
    - e.g. CPU registers & memory
  - It cannot easily be applied to other resources
    - e.g. printers & tape drives

# Deadlock Prevention (con't)

- Circular wait
    - impose a total ordering of all resources types
    - a process requests resources in an increasing order
      - Let  $R=\{R_0, R_1, \dots, R_N\}$  be the set of resource types
      - When request  $R_k$ , should release all  $R_i, i \geq k$
    - Example:
      - $F(\text{tape drive}) = 1, F(\text{disk drive}) = 5, F(\text{printer}) = 12$
      - A process must request tape and disk drive before printer
    - proof: counter-example does not exist
      - $P_0 (R_0) \rightarrow R_1, P_1 (R_1) \rightarrow R_2, \dots, P_N(R_N) \rightarrow R_0$
      - Conflict:  $R_0 < R_1 < R_2 < \dots < R_N < R_0$
- PN hold  $R_N$ , wait  $R_0$




# Avoidance Algorithms

- **Single instance** of a resource type
  - resource-allocation graph (RAG) algorithm based on circle detection
- **Multiple instances** of a resource type
  - banker's algorithm based on safe sequence detection


# Resource-Allocation Graph (RAG) Algorithm

- **Request edge**:  $P_i \rightarrow R_j$ 
  - Process  $P_i$  is **waiting** for resource  $R_j$
- **Assignment edge**:  $R_j \rightarrow P_i$ 
  - Resource  $R_j$  is **allocated** and held by process  $P_i$
- **Claim edge**:  $P_i \rightarrow R_j$ 
  - process  $P_i$  may **request**  $R_j$  in the future
- **Claim edge** converts to **request edge**
  - When a **resource is requested** by process
- **Assignment edge** converts to a **claim edge**
  - When a **resource is released** by a process



# Resource-Allocation Graph (RAG) Algorithm

- Resources **must be claimed a priori** in the system
- Grant a request** only if **NO cycle created**
- Check for safety using a **cycle-detection algorithm**,  $O(n^2)$
- Example: R2 cannot be allocated to P2




# Avoidance Algorithms

- Single instance of a resource type
  - resource-allocation graph (RAG) algorithm based on circle detection
- Multiple instances of a resource type
  - banker's algorithm based on safe sequence detection

# Deadlock Avoidance

- **safe state**: a system is in a safe state if there exists **a sequence of allocations** to satisfy requests by all processes
  - This sequence of allocations is called **safe sequence**
- **safe state** -> no deadlock
- **unsafe state** -> **possibility** of deadlock
- **deadlock avoidance** -> **ensure that a system never enters an unsafe state**



# Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t0:

Hint from processes → Max Needs    Current Holding

P0	10	5
P1	4	2
P2	9	2

→ <P1, P0, P2> is a safe sequence

# Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t0:

	<u>Max Needs</u>	<u>Current Holding</u>	<u>Available</u>
P0	10	5	
P1	4	2	3
P2	9	2	

➔  $\langle P1, P0, P2 \rangle$  is a safe sequence

1. P1 satisfies its allocation with 3 available resources

# Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t0:

	<u>Max Needs</u>	<u>Current Holding</u>	<u>Available</u>
P0	10	5	5
P1	4	0	
P2	9	2	

➔  $\langle P1, P0, P2 \rangle$  is a safe sequence

1. P1 satisfies its allocation with 3 available resources
2. P0 satisfies its allocation with 5 available resources



# Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at  $t_0$ :

	<u>Max Needs</u>	<u>Current Holding</u>	<u>Available</u>
P0	10	0	
P1	4	0	
P2	9	2	10

➔  $\langle P1, P0, P2 \rangle$  is a safe sequence

1. P1 satisfies its allocation with 3 available resources
2. P0 satisfies its allocation with 5 available resources
3. P2 satisfies its allocation with 10 available resources

# Un-Safe State w/o Safe Sequence

- Assuming at t1:

	<u>Max Needs</u>	<u>Current Holding</u>	<u>Available</u>
P0	10	5	
P1	4	2	2
P2	9	<del>2</del> 3	

*if P2 requests & is allocated 1 more tape drive*

➔ No safe sequence exist...

➔ this allocation enters the system into an unsafe state

- A request is only granted if the allocation leaves the system in a safe state

# Banker's Algorithm

- Use for **multiple instances** of each resource type
- Banker algorithm:
  - Use a general **safety algorithm** to **pre-determine** if any **safe sequence** exists after allocation
  - Only proceed the allocation if safe sequence exists
- Safety algorithm:
  1. Assume processes need **maximum** resources
  2. Find a process that can be satisfied by free resources
  3. Free the resource usage of the process
  4. Repeat to step 2 until all processes are satisfied

# Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:3, B:3, C:2

	<u>Max</u>	<u>Allocation</u>	<u>Need(Max.-Alloc.)</u>
	A B C	A B C	A B C
P0	7 5 3	0 1 0	7 4 3
P1	3 2 2	2 0 0	1 2 2
P2	9 0 2	3 0 2	6 0 0
P3	2 2 2	2 1 1	0 1 1
P4	4 3 3	0 0 2	4 3 1

- Safe sequence: P1

# Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:5, B:3, C:2

	<u>Max</u>	<u>Allocation</u>	<u>Need(Max.-Alloc.)</u>
	A B C	A B C	A B C
P0	7 5 3	0 1 0	7 4 3
P1	3 2 2	2 0 0	1 2 2
P2	9 0 2	3 0 2	6 0 0
P3	2 2 2	2 1 1	0 1 1
P4	4 3 3	0 0 2	4 3 1

- Safe sequence: P1, P3

# Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:7, B:4, C:3

	<u>Max</u>	<u>Allocation</u>	<u>Need(Max.-Alloc.)</u>
	A B C	A B C	A B C
P0	7 5 3	0 1 0	7 4 3
P1	3 2 2	2 0 0	1 2 2
P2	9 0 2	3 0 2	6 0 0
P3	2 2 2	2 1 1	0 1 1
P4	4 3 3	0 0 2	4 3 1

- Safe sequence: P1, P3, P4

# Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:7, B:4, C:5

	<u>Max</u>	<u>Allocation</u>	<u>Need(Max.-Alloc.)</u>
	A B C	A B C	A B C
P0	7 5 3	0 1 0	7 4 3
P1	3 2 2	2 0 0	1 2 2
P2	9 0 2	3 0 2	6 0 0
P3	2 2 2	2 1 1	0 1 1
P4	4 3 3	0 0 2	4 3 1

- Safe sequence: P1, P3, P4, P2

# Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:10, B:4, C:7

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
P1	3	2	2	2	0	0	1	2	2
P2	9	0	2	3	0	2	6	0	0
P3	2	2	2	2	1	1	0	1	1
P4	4	3	3	0	0	2	4	3	1

- Safe sequence: P1, P3, P4, P2, P0



# Banker's Algorithm Example

- Total instances: A:10, B:5, C:7
- Available instances: A:3, B:3, C:2

	<u>Max</u>	<u>Allocation</u>	<u>Need(Max-Alloc)</u>
	A B C	A B C	A B C
P0	7 5 3	0 1 0	7 4 3
→ P1	3 2 2	2 0 0	1 2 2
P2	9 0 2	3 0 2	6 0 0
P3	2 2 2	2 1 1	0 1 1
→ P4	4 3 3	0 0 2	4 3 1

- If Request (P1) = (1, 0, 2): P1 allocation → 3, 0, 2
  - Enter another safe state (Safe sequence: P1, P3, P4, P0, P2)
- If Request (P4) = (3, 3, 0): P4 allocation → 3, 3, 2
  - enter into an unsafe state (no safe sequence can be found!)

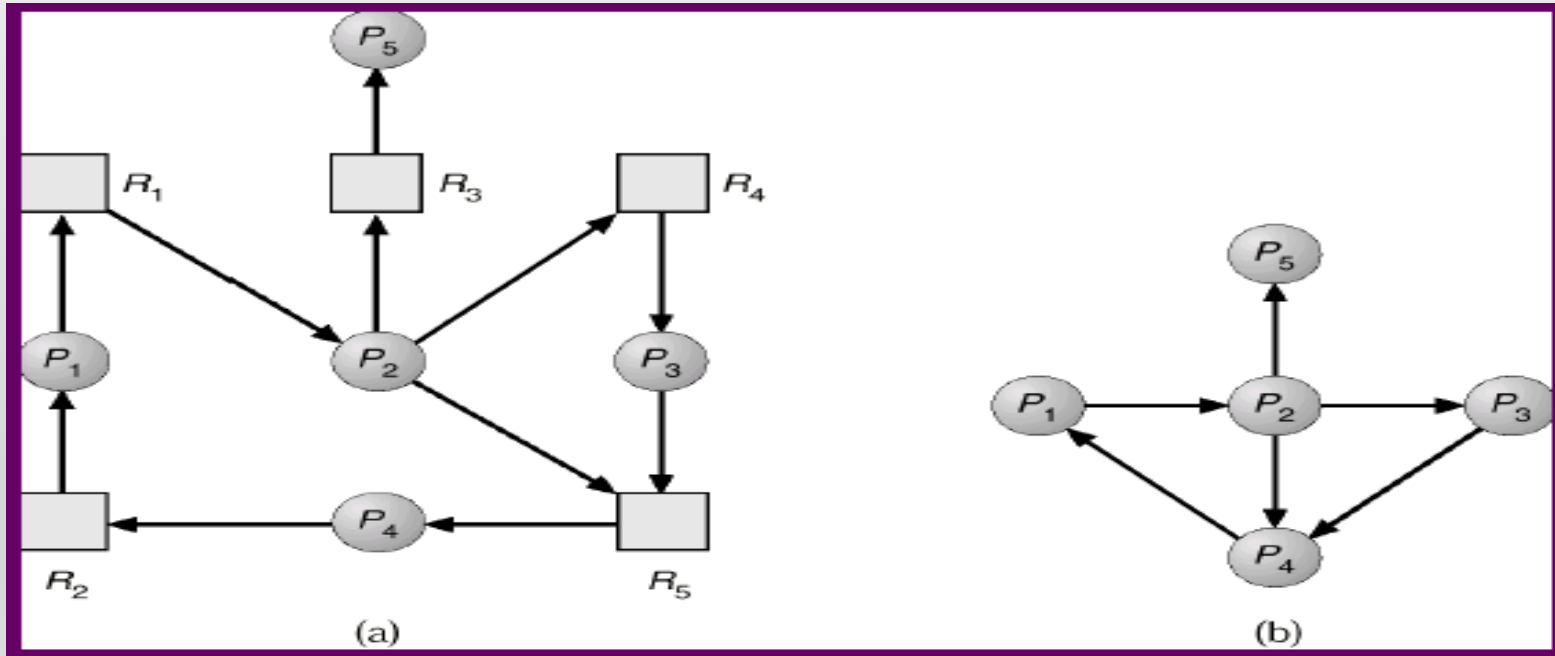
# Review Slides ( II )

- deadlock prevention methods?
  - mutual exclusion
  - hold & wait
  - no preemption
  - circular wait
- deadlock avoidance methods?
  - safe state definition?
  - safe sequence?
  - claim edge?

# Deadlock Detection & Deadlock Recovery

# Deadlock Detection

- **Single instance** of each resource type
  - convert request/assignment edges into wait-for graph
  - deadlock exists if there is a cycle in the wait-for graph



Resource-Allocation Graph

Corresponding wait-for graph

# Multiple-Instance for Each Resource Type

- Total instances: A:7, B:2, C:6
- Available instances: A:0, B:0, C:0

	<u>Allocation</u>			<u>Request</u>		
	A	B	C	A	B	C
P0	0	1	0	0	0	0
P1	2	0	0	2	0	2
P2	3	0	3	0	0	0
P3	2	1	1	1	0	0
P4	0	0	2	0	0	2

- The system is in a safe state  $\Rightarrow$   $\langle P0, P2, P3, P1, P4 \rangle$   
 $\Rightarrow$  no deadlock
- If P2 request =  $\langle 0, 0, 1 \rangle \Rightarrow$  no safe sequence can be found  
 $\Rightarrow$  the system is deadlocked

# Deadlock Recovery

- Process termination
  - abort all deadlocked processes
  - abort 1 process at a time until the deadlock cycle is eliminated
    - which process should we abort first?
- Resource preemption
  - select a victim: which one to preempt?
  - rollback: partial rollback or total rollback?
  - starvation: can the same process be preempted always?

# Reading Material & HW

- Chap 7
- Problem Set
  - 7.6, 7.7, 7.8, 7.9, 7.12, 7.13

*Q & A*

*Thank you for your attention*