



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

*Operating System*

# Lecture 06: Process Synchronization

Shuo-Han Chen 陳碩漢

[shch@nycu.edu.tw](mailto:shch@nycu.edu.tw)

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

# Course Schedule

W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01 Due 10/5
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Typhoon – No class	V	
6	Oct. 9	Lec07: Memory Management	V	
7	Oct. 16	Lec08: Virtual Memory Management	V	HW02 Due 11/2
8	Oct. 23	Lec04: Multithreaded Programming	V	
9	Oct. 30	Midterm Exam		
10	Nov. 6	Lec05: Process Scheduling	V	Let's take a breath
11	Nov. 13	Lec06: Process Synchronization & Deadlocks	X	HW03
12	Nov. 20	School Event – No class		
13	Nov. 27	Lec09: File System Interface	V	
14	Dec. 4	Lec10: File System Implementation	V	HW04
15	Dec. 11	Lec11: Mass Storage System & Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

# Overview

- Background
- Critical Section
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Atomic Transactions

# Background

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanism to ensure the orderly execution of cooperating processes

# Consumer & Producer Problem

- Determine whether buffer is **empty** or **full**
  - Previously: use **in**, **out** position
  - Now: use **count** value

```
/*producer*/  
while (1) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE) ;  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
/*consumer*/  
while (1) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```

# Concurrent Operations on counter

- The statement “counter++” may be implemented in machine language as:

```
move ax, counter
```

```
add  ax, 1
```

```
move counter, ax
```

- The statement “counter--” may be implemented as:

```
move    bx, counter
```

```
sub     bx, 1
```

```
move    counter, bx
```

# Instruction Interleaving

- Assume counter is initially 5. One interleaving of statement is:

producer: move ax, counter      ➔ ax = 5

producer: add ax, 1      ➔ ax = 6

*context switch*

consumer: move bx, counter      ➔ bx = 5

consumer: sub bx, 1      ➔ bx = 4

*context switch*

producer: move counter, ax      ➔ counter = 6

*context switch*

consumer: move counter, bx      ➔ counter = 4

- The value of counter may be either 4, 5, or 6, where the correct result should be 5

# Race Condition

- **Race condition**: the situation where several processes **access and manipulate shared data concurrently**. The final value of the shared data depends upon which process finishes last
- To prevent race condition, concurrent processes must be **synchronized**
  - On a single-processor machine, we could **disable interrupt** or use **non-preemptive CPU scheduling**
- Commonly described as **critical section problem**




# Critical Section

# The Critical-Section Problem

- Purpose: a protocol for processes to cooperate
- Problem description:
  - N processes are competing to use some shared data
  - Each process has a code segment, called critical section, in which the shared data is accessed
  - Ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section -> mutually exclusive

# The Critical-Section Problem

- General code section structure
  - Only one process can be in a critical section



# Critical Section Requirements

1. **Mutual Exclusion**: if process P is executing in its CS, no other processes can be executing in their CS
  2. **Progress**: if no process is executing in its CS and there exist some processes that wish to enter their CS, these processes cannot be postponed indefinitely
  3. **Bounded Waiting**: A bound must exist on the number of times that other processes are allowed to enter their CS after a process has made a request to enter its CS
- -> How to design entry and exist section to satisfy the above requirement?

# Review Slides (1)


- Race condition?
- Critical-Section (CS) problem? 4 sections?
  - entry, CS, exit, remainder
- 3 requirements for solutions to CS problems?
  - mutual exclusion
  - progress
  - bounded waiting

# Critical Section Solutions & Synchronization Tools

- Software Solution
- Synchronization Hardware
- Semaphore
- Monitor

# Algorithm for Two Processes

- Only 2 processes, P0 and P1
- Shared variables
  - int **turn**; //initially turn = 0
  - **turn = i** -> P<sub>i</sub> can enter its critical section



→ Mutual exclusion? **Yes**    Progress? **No**  
Bounded-Wait? **Yes**

# Peterson's Solution for Two Processes

- Shared variables
  - int **turn**; //initially turn = 0
  - **turn = i -> Pi can enter its critical section**
  - Boolean **flag**[2]; //initially flag [0] = flag [1] = false
  - **Flag [i] = true -> Pi ready to enter its critical section**

//Pi:

do {

```
flag[ i ] = TRUE;
turn = j ;
while (flag [ j ] &&
      turn == j ) ;
```

**critical section**

```
flag [ i ] = FALSE ;
```

remainder section

} while (1) ;

Enter CS when **either**:

1. a process gets its turn
2. the other process is not ready



# Proof of Peterson's Solution

- Mutual exclusion:
  - If P0 CS -> flag[1] == false || turn == 0
  - If P1 CS -> flag[0] == false || turn == 1
- Assume both processes in CS -> flag[0] == flag[1] == true
  - -> turn==0 for P0 to enter, turn==1 for P1 to enter
  - However, "turn" will be either 0 or 1 because its value will be set for both processes, but only one value will last
  - Therefore, P0 ,P1 can't in CS at the same time!

```
/* process 0 */
do {
    flag[ 0 ] = TRUE;
    turn = 1 ;
    while (flag [ 1 ] && turn == 1 ) ;
    → critical section
    flag [ 0 ] = FALSE ;
    remainder section
} while (1);
```

```
/* process 1 */
do {
    flag[ 1 ] = TRUE;
    turn = 0 ;
    while (flag [ 0 ] && turn == 0 ) ;
    → critical section
    flag [ 1 ] = FALSE ;
    remainder section
} while (1);
```

# Proof of Peterson's Solution

- Progress (e.g., P0 wishes to enter its CS):
  1. If P1 is not ready  $\rightarrow \text{flag}[1] = \text{false} \rightarrow \text{P0 can enter}$
  2. If both are ready  $\rightarrow \text{flag}[0] == \text{flag}[1] == \text{true}$   
If  $\text{turn} == 0$  then P0 enters, otherwise P1 enters
- **Either cases, some waiting process can enter CS!**

```
/* process 0 */
do {
    flag[ 0 ] = TRUE;
    turn = 1 ;
    → while (flag [ 1 ] && turn == 1 ) ;
        critical section
    flag [ 0 ] = FALSE ;
        remainder section
} while (1);
```

(2)



(1)



```
/* process 1 */
do {
    flag[ 1 ] = TRUE;
    turn = 0 ;
    while (flag [ 0 ] && turn == 0 ) ;
        critical section
    flag [ 1 ] = FALSE ;
        remainder section
} while (1);
```

# Proof of Peterson's Solution

- Bounded waiting (e.g., P0 wishes to enter its CS):
  1. Once P1 exits CS  $\rightarrow$   $\text{flag}[1] == \text{false}$   $\rightarrow$  P0 can enter
  2. If P1 exits CS && reset  $\text{flag}[1] = \text{true}$   
 $\rightarrow$   $\text{turn} == 0$  (overwrite P0 setting)  $\rightarrow$  P0 can enter
- **P0 won't wait indefinitely!**

```
/* process 0 */
do {
    flag[ 0 ] = TRUE;
    turn = 1 ;
    → while (flag [ 1 ] && turn == 1 ) ;
        critical section
    flag [ 0 ] = FALSE ;
        remainder section
} while (1) ;
```

(2)



(1)



```
/* process 1 */
do {
    flag[ 1 ] = TRUE;
    turn = 0 ;
    while (flag [ 0 ] && turn == 0 ) ;
        critical section
    flag [ 1 ] = FALSE ;
        remainder section
} while (1) ;
```

# Producer/Consumer Problem

- Producer process

```
while (TRUE) {  
    entry-section( );  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
    computing();  
    exit-section( );  
}
```

- Consumer process

```
while (TRUE) {  
    entry-section( );  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    computing();  
    exit-section( );  
}
```

➔ Incorrect: deadlock, if consumer enters the CS first.

# Producer/Consumer Problem

- Producer process

```
while (TRUE) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter++;  
    computing();  
    exit-section( );  
}
```

- Consumer process

```
while (TRUE) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter--;  
    computing();  
    exit-section( );  
}
```

➔ Correct but poor performance

# Producer/Consumer Problem

- Producer process

```
while (TRUE) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter++;  
    exit-section( );  
    computing();  
}
```

- Consumer process

```
while (TRUE) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter--;  
    exit-section( );  
    computing();  
}
```

➔ Correct & Maximize concurrent performance

# Bakery Algorithm (n processes)

- Before enter its CS, each process receives a #
- Holder of the smallest # enters CS
- The numbering scheme always generates # in non-decreasing order; i.e., 1,2,3,3,4,5,5,5
- If processes  $P_i$  and  $P_j$  receive the same #, if  $i < j$ , then  $P_i$  is served first
- Notation:
  - $(a, b) < (c, d)$  if  $a < c$  or if  $a == c \ \&\& \ b < d$

# Bakery Algorithm (n processes)

//Process i: do {

Get ticket

choosing [ i ] = TRUE ;

num[ i ] = max(num[0],num[1],...,num[n-1]) + 1;

choosing [ i ] = FALSE ;

FCFS

for (j = 0; j < n; j++) {

while (choosing [ j ] ) ;

while ((num[ j ] != 0) &&

((num[ j ], j) < (num[ i ], i))) ;

}

Cannot compare when  
num is being modified

release  
ticket

num[ i ] = 0 ;

reminder section

} while (1) ;

- Bounded-waiting because processes enter CS on a First-Come, First Served basis



# Bakery Algorithm (n processes)

- Why cannot compare when num is being modified?
- Without locking...
  1. Let 5 be the current maximum number
  2. If P1 and P4 take number together, but P4 finishes before P1
    - **P1 = 0**; P4 = 6 -> P4 will enter the CS
  3. After P1 takes the number
    - **P1 = P4 = 6 -> P1 will enter the CS as well!!!**
- With locking...
  - P4 will have to wait until P1 finish taking the number
  - **Both P1 & P4 will have the new number “6” before comparison**

# Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of **type** `pthread_mutex_t` and initialized with `pthread_mutex_init()`
- A mutex is destroyed with `pthread_mutex_destory()`
- A critical section can then be protected using `pthread_mutex_lock()` and `pthread_mutex_unlock()`
- Example:

```
#include "pthread.h"           specify default
pthread_mutex  mutex;         attribute for the mutex
pthread_mutex_init (&mutex, NULL);
pthread_mutex_lock(&mutex);    // enter critical section

    Critical Section

pthread_mutex_unlock(&mutex);  // leave critical section
pthread_mutex_destory(&mutex);
```

# Condition Variables (CV)

- CV represent some **condition** that a thread can:
  - Wait on, until the condition occurs; or
  - Notify other waiting threads that the condition has occurred
- Three operations on condition variables:
  - **wait()** --- **Block** until another thread calls **signal()** or **broadcast()** on the CV
  - **signal()** --- Wake up **one thread** waiting on the CV
  - **broadcast()** --- Wake up **all threads** waiting on the CV
- In Pthread, CV **type** is a **pthread\_cond\_t**
  - Use **pthread\_cond\_init()** to initialize
  - **pthread\_cond\_wait (&theCV, &somelock)**
  - **pthread\_cond\_signal (&theCV)**
  - **pthread\_cond\_broadcast (&theCV)**

# Using Condition Variable

- Example:
  - A threads is designed to **take action when x=0**
  - Another thread is responsible for decrementing the counter

```
pthread_cond_t  cond;  
pthread_cond_init (cond, NULL);
```

```
pthread_mutex_t  mutex;  
pthread_mutex_init (mutex, NULL);
```

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- All condition variable operation MUST be performed while a mutex is locked!!!

# Using Condition Variable

```
action() {  
→ pthread_mutex_lock (&mutex)  
  while (x != 0)  
    pthread_cond_wait (cond, mutex);  
  pthread_mutex_unlock (&mutex);  
  take_action();  
}
```

```
→ counter() {  
  pthread_mutex_lock (&mutex)  
  x--;  
  if (x==0)  
    pthread_cond_signal (cond);  
  pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex

# Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex
2. Wait()

- Put the thread into **sleep & releases the lock**

1. Lock mutex

# Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        → pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex
2. Wait()

- Put the thread into **sleep & releases the lock**
- **Waked up**, but the thread is locked

1. Lock mutex
2. Signal()

# Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
        → pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex
2. Wait()
  - Put the thread into **sleep & releases the lock**
  - **Waked up**, but the thread is locked
  - **Re-acquire lock** and resume execution

1. Lock mutex
2. Signal()
3. Releases the lock



# Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex
2. Wait()
  - Put the thread into **sleep & releases the lock**
  - **Waked up**, but the thread is locked
  - **Re-acquire lock** and resume execution
3. Release the lock

1. Lock mutex
2. Signal()
3. Releases the lock

# Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- What really happens...

1. Lock mutex

2. Wait()

- Put the thread into **sleep & releases the lock**

- **Waked up**, but the thread is locked

- **Re-acquire lock** and resume execution

3. Release the lock

1. Lock mutex

2. Signal()

3. Releases the lock

Another reason why  
condition variable op.  
MUST within mutex lock

# ThreadPool Implementation

## Task structure

```
typedef struct {  
    void (*function)(void *);  
    void *argument;  
} threadpool_task_t;
```

## Threadpool structure

```
struct threadpool_t {  
    pthread_mutex_t lock;  
    pthread_cond_t notify;  
    pthread_t *threads;  
    threadpool_task_t *queue;  
    int thread_count;  
    int queue_size;  
    int head;  
    int tail;  
    int count;  
    int shutdown;  
    int started;  
};
```

## Allocate thread and task queue

```
/* Allocate thread and task queue */  
pool->threads = (pthread_t *) malloc(sizeof(pthread_t) * thread_count);  
pool->queue = (threadpool_task_t *) malloc(sizeof(threadpool_task_t) * queue_size);
```

# ThreadPool Implementation

```
static void *threadpool_thread(void *threadpool)
{
    threadpool_t *pool = (threadpool_t *)threadpool;
    threadpool_task_t task;

    for(;;) {
        /* Lock must be taken to wait on conditional variable */
        pthread_mutex_lock(&(pool->lock));

        /* Wait on condition variable, check for spurious wakeups.
           When returning from pthread_cond_wait(), we own the lock. */
        while((pool->count == 0) && (!pool->shutdown)) {
            pthread_cond_wait(&(pool->notify), &(pool->lock));
        }
    }
}
```

# ThreadPool Implementation

```
/* Grab our task */  
task.function = pool->queue[pool->head].function;  
task.argument = pool->queue[pool->head].argument;  
pool->head += 1;  
pool->head = (pool->head == pool->queue_size) ? 0 : pool->head;  
pool->count -= 1;  
  
/* Unlock */  
pthread_mutex_unlock(&(pool->lock));  
  
/* Get to work */  
(* (task.function))(task.argument);  
}
```

# Synchronization HW

# Hardware Support

- The CS problem occurs because the modification of a shared variable may be **interrupted**
- If disable interrupts when in CS...
  - not feasible in multiprocessor machine
  - clock interrupts cannot fire in any machine
- HW support solution: **atomic instructions**
  - atomic: **as one uninterruptible unit**
  - examples: **TestAndSet**(var), **Swap**(a,b)

# Atomic TestAndSet()

```
boolean TestAndSet ( bool  &lock) {  
    bool  value = lock ;  
    lock = TRUE ;  
    return value ;  
}
```

**execute atomically:**  
return the value of “lock”  
and set “lock” to TRUE

Mutual exclusion? **Yes** Progress? **Yes** Bounded-Wait? **No!**

Shared data: boolean lock; //initially lock = FALSE;

do { // P0

while (TestAndSet (lock) ) ;

critical section

lock = FALSE;

remainder section

} while (1) ;

do { // P1

while (TestAndSet (lock) ) ;

critical section

lock = FALSE;

remainder section

} while (1) ;

obtain lock

release lock



# Atomic Swap()

- Idea: enter CS if lock==false:

Shared data: boolean **lock**; //initially **lock = FALSE**;

do { // P0

```
key0 = TRUE;  
while (key0 == TRUE)  
    Swap (lock, key0);
```

critical section

```
lock = FALSE;
```

remainder section

} while (1);

do { // P1

```
key1 = TRUE;  
while (key1 == TRUE)  
    Swap (lock, key1);
```

critical section

```
lock = FALSE;
```

remainder section

} while (1);

Mutual exclusion? **Yes** Progress? **Yes** Bounded-Wait? **No!**

## Review Slide (2)

- Use software solution to solve CS?
  - Peterson's and Bakery algorithms
- Use HW support to solve CS?
  - TestAndTest(), Swap()

# Semaphores

# Semaphore

- A **tool** to generalize the synchronization problem (**easier to solve, but no guarantee for correctness**)
- More specifically...
  - a **record** of **how many units** of a particular resource are available
    - If #record = 1 -> **binary semaphore, mutex lock**
    - If #record > 1 -> **counting semaphore**
  - accessed only through 2 **atomic** ops: **wait** & **signal**
- **Spinlock** implementation:
  - Semaphore is an **integer variable**

```
wait (S) {                               signal (S) {
    while (S <= 0) ;                       S++;
    S--;                                   }
}
```

← busy waiting

# POSIX Semaphore

- Semaphore is part of **POSIX standard** BUT it is **not belonged to Pthread**
    - **It can be used with or without thread**
  - POSIX Semaphore routines:
    - **sem\_init**(sem\_t \*sem, int pshared, unsigned int value)
    - **sem\_wait**(sem\_t \*sem)
    - **sem\_post**(sem\_t \*sem)
    - **sem\_getvalue**(sem\_t \*sem, int \*valptr)
    - **sem\_destory**(sem\_t \*sem)
- Initial value of the semaphore
- Current value of the semaphore

- Example:

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem);
sem_wait(&sem);
// critical section
sem_post(&sem);
sem_destory(&sem);
```

# n-Process Critical Section Problem

- shared data:

```
semaphore mutex ; // initially mutex = 1
```

- Process  $P_i$ :

```
do {
```

```
    wait (mutex) ;    // pthread_mutex_lock(&mutex)
```

```
        critical section
```

```
    signal (mutex); // pthread_mutex_unlock(&mutex)
```

```
        remainder section
```

```
} while (1) ;
```

- Progress? Yes
- Bounded waiting? Depends on the implementation of wait()


# Non-busy waiting Implementation

- Semaphore is **data struct with a queue**
  - may use any queuing strategy (FIFO, FILO, etc)

```
typedef struct {  
    int value; // init to 0  
    struct process *L ;  
    // “PCB” queue  
} semaphore ;
```

E.g.,:

**Value = -3**



- wait() and signal()
  - use system calls: sleep() and wakeup()
  - must be executed atomically

```
void wait (semaphore S) {  
    S.value--; // subtract first  
    if (S.value < 0) {  
        add this process to S.L ;  
        sleep( );  
    }  
}
```

```
void signal (semaphore S) {  
    S.value++;  
    if (S.value <= 0) {  
        remove a process P from S.L ;  
        wakeup(P);  
    }  
}
```

# Atomic Operation

- How to ensure atomic wait & signal ops?
  - Single-processor: disable interrupts
  - Multi-processor:
    - HW support (e.g. Test-And-Set, Swap)
    - SW solution (Peterson's solution, Bakery algorithm)



# Semaphore with Critical Section

```
void wait (semaphore S) {  
    entry-section( );  
    S.value--;  
    if (S.value < 0) {  
        add this process to S.L ;  
        exit-section( );  
        sleep( );  
    }  
    else {  
        exit-section( );  
    }  
}
```

```
void signal (semaphore S) {  
    entry-section( );  
    S.value++;  
    if (S.value <= 0)  
        remove a process P from S.L;  
    exit-section( );  
    wakeup(P);  
}  
else {  
    exit-section( );  
}  
}
```

- Busy waiting for entry-section()?
  - limited to only the CS of wait & signal (~10 instructions) -> very short period of time

# Cooperation Synchronization

- P1 executes S1 ; 2 executes S2
  - S2 be executed only after S1 has completed
- Implementation:
  - shared var:  
semaphore **sync** ; // initially sync = 0

P1:

S1 ;

signal (**sync**) ;

P2:

wait (**sync**) ;

S2 ;

# A More Complicated Example

(Initially, all semaphores are 0)

begin

P1: S1; signal(**a**); signal(**b**);

P2: wait(**a**); S2; signal(**c**);

P3: wait(**b**); S3; signal(**d**);


P4: wait(**c**); S4; signal(**e**); signal(**f**);

P5: wait(**e**); S5; signal(**g**);

P6: wait(**f**); wait(**d**); S6; signal(**h**);


P7: wait(**g**); wait(**h**); S7;

end



# Deadlocks & Starvation

- **Deadlocks**: 2 processes are waiting indefinitely for each other to release resources
- **Starvation**: example: LIFO queue in semaphore process queue



# Review Slide (3)

- What's semaphore? 2 operations?
- What's busy-waiting (spinlock) semaphore?
- What's non-busy-waiting (non-spinlock) semaphore?
- How to ensure atomic wait & signal ops?
- Deadlock? starvation?

# Classical Synchronization Problems

# Listing & Purpose

- Purpose: **used for testing newly proposed synchronization scheme**
- Bounded-Buffer (Producer-Consumer) Problem
- Reader-Writers Problem
- Dining-Philosopher Problem

# Bounded-Buffer Problem

- A pool of  $n$  buffers, each capable of holding one item
- Producer:
  - grab an empty buffer
  - place an item into the buffer
  - waits if no empty buffer is available
- Consumer:
  - grab a buffer and retracts the item
  - place the buffer back to the free pool
  - waits if all buffers are empty



# Readers-Writers Problem

- A set of shared data objects
- A group of processes
  - reader processes (read shared objects)
  - writer processes (update shared objects)
  - a writer process has exclusive access to a shared object
- Different variations involving priority
  - **first RW problem**: no reader will be kept waiting unless a writer is updating a shared object
  - **second RW problem**: once a writer is ready, it performs the updates as soon as the shared object is released
    - writer has higher priority than reader
    - once a writer is ready, no new reader may start reading

# First Reader-Writer Algorithm

// mutual exclusion for write

semaphore wrt=1

// mutual exclusion for readcount

semaphore mutex=1

int readcount=0;

Writer(){

while(TRUE){

wait(wrt);

// Writer Code

signal(wrt);

}

}

- Readers share a single wrt lock
- Writer may have starvation problem

Acquire write lock  
if reads haven't

Reader(){

while(TRUE){

wait(mutex);

readcount++;

if(readcount==1)

wait(wrt);

signal(mutex);

// Reader Code

wait(mutex);

readcount--;

if(readcount==0)

signal(wrt);


signal(mutex);

}

}

# Dining-Philosophers Problem

- 5 persons sitting on 5 chairs with 5 chopsticks
- A person is either thinking or eating
  - thinking: no interaction with the rest 4 persons
  - eating: need 2 chopsticks at hand
  - a person picks up 1 chopstick at a time
  - done eating: put down both chopsticks
- deadlock problem
  - one chopstick as one semaphore
- starvation problem



# Monitors

# Motivation

- Although semaphores provide a convenient and effective synchronization mechanism, its correctness is depending on the programmer
  - All processes access a shared data object must execute `wait()` and `signal()` in the right order and right place
  - This may not be true because honest programming error or uncooperative programmer

# Monitor --- A high-level language construct

- The representation of a **monitor type** consists of
  - declarations of **variables** whose values define the state of an instance of the type
  - **Procedures/functions** that implement operations on the type
- The monitor type is similar to a **class in O.O. language**
  - A procedure within a monitor can access only **local variables** and the formal **parameters**
  - The local variables of a monitor can be used only by the local procedures
- But, the monitor ensures that **only one process at a time can be active within the monitor**
- Similar idea is incorporated to many prog. language:
  - concurrent pascal, C# and Java

# Monitor


- High-level synchronization construct that allows the safe sharing of an abstract data type among

concurrent processes

## Syntax

```
monitor monitor-name {  
    // shared variable declarations  
    procedure body  $P_1$  (...) {  
        ...  
    }  
    procedure body  $P_2$  (...) {  
        ...  
    }  
    procedure body  $P_n$  (...) {  
        ...  
    }  
    initialization code {  
    }  
}
```

## Schematic View



# Monitor Condition Variables

- To allow a process to **wait within** the monitor, a **condition variable** must be declared, as

**condition x, y;**

- Condition variable can only be used with the operations **wait()** and **signal()**

- **x.wait();**

**means that the process invoking this operation is suspended until another process invokes**


- **x.signal();**

**resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect**

**(In contrast, signal always change the state of a semaphore)**



# Monitor With Condition Variables



# Dining Philosophers Example

```
monitor dp {  
    enum {thinking, hungry, eating} state[5]; //current state  
    condition self[5]; //delay eating if can't obtain chopsticks  
    void pickup(int i)           // pickup chopsticks  
    void putdown(int i)         // putdown chopsticks  
    void test(int i)             // try to eat  
    void init() {  
        for (int i = 0; i < 5; i++)  
            state[i] = thinking;  
    }  
}
```

```

void pickup(int i) { state[i] =
    hungry; test(i); //try to
    eat
    if (state[i] != eating)
        self[i].wait(); //wait to eat
}

```

//try to let  $P_i$  eat (if it is hungry)

```

void test(int i) {
    if ( (state[(i + 4) % 5] != eating) &&(state[(i + 1) % 5] != eating) && (state[i]
        == hungry) ) {
        //No neighbors are eating and  $P_i$  is hungry
        state[i] = eating;
        self[i].signal();
    }
}

```


```

void putdown(int i) {
    state[i] = thinking;
    // check if neighbors
    // are waiting to eat
    test((i+4) % 5);
    test((i+1) % 5);
}

```

If  $P_i$  is suspended, resume it  
 If  $P_i$  is not suspended, **no effect**


# An illustration



P1:  
DiningPhilosophers.pickup(1)  
eat  
DiningPhilosophers.putdown(1)

P2:  
DiningPhilosophers.pickup(2)  
eat  
DiningPhilosophers.putdown(2)

# An illustration




```
void pickup(int i) {  
    state[i] = hungry;  
    test(i); //try to eat  
    if (state[i] != eating)  
        self[i].wait(); //wait to eat  
}
```

P1:  
→ DiningPhilosophers.pickup(1)  
eat  
DiningPhilosophers.putdown(1)

P2:  
DiningPhilosophers.pickup(2)  
eat  
DiningPhilosophers.putdown(2)

# An illustration




```
void pickup(int i) {  
    state[i] = hungry;  
    test(i); //try to eat  
    if (state[i] != eating)  
        self[i].wait(); //wait to eat  
}
```

P1:  
→ DiningPhilosophers.pickup(1)  
eat  
DiningPhilosophers.putdown(1)

P2:  
DiningPhilosophers.pickup(2)  
eat  
DiningPhilosophers.putdown(2)

# An illustration



**hungry**  $\rightarrow$  self[2].wait

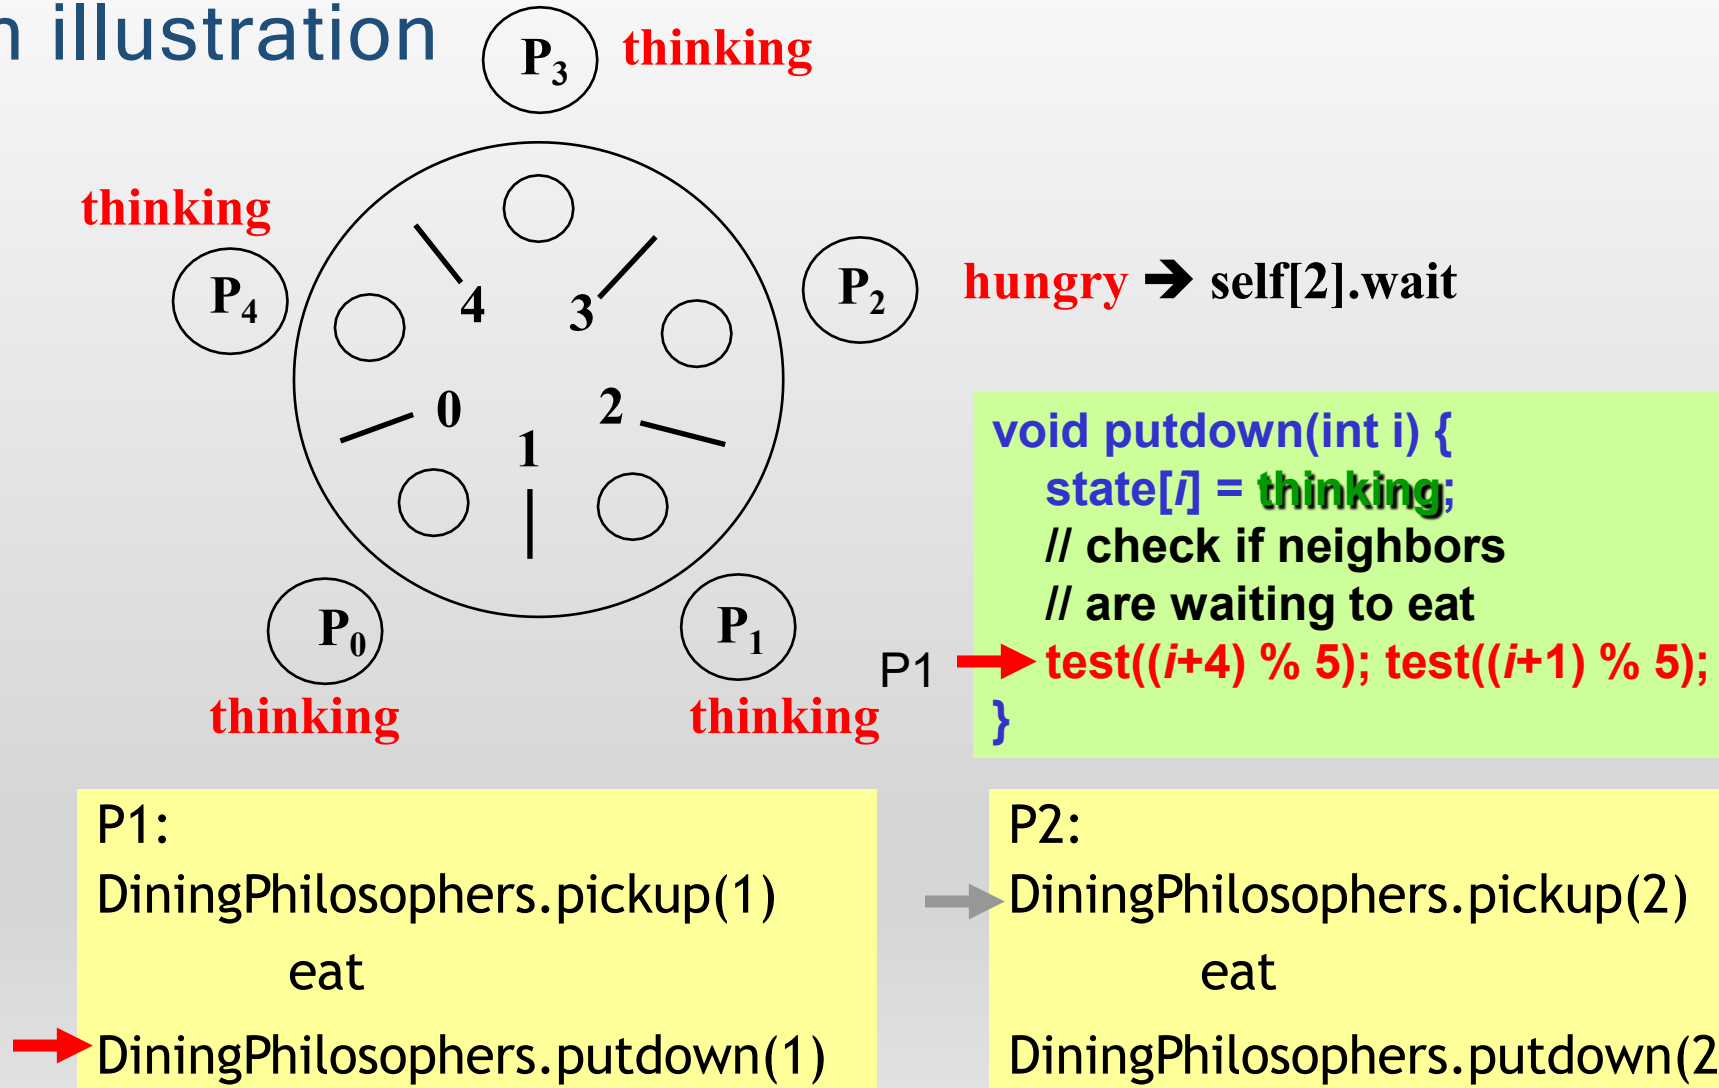
```
void pickup(int i) {  
    state[i] = hungry;  
    test(i); //try to eat  
    if (state[i] != eating)  
        self[i].wait(); //wait to eat  
}
```

P2  $\rightarrow$

P1:  
DiningPhilosophers.pickup(1)  
eat  
DiningPhilosophers.putdown(1)

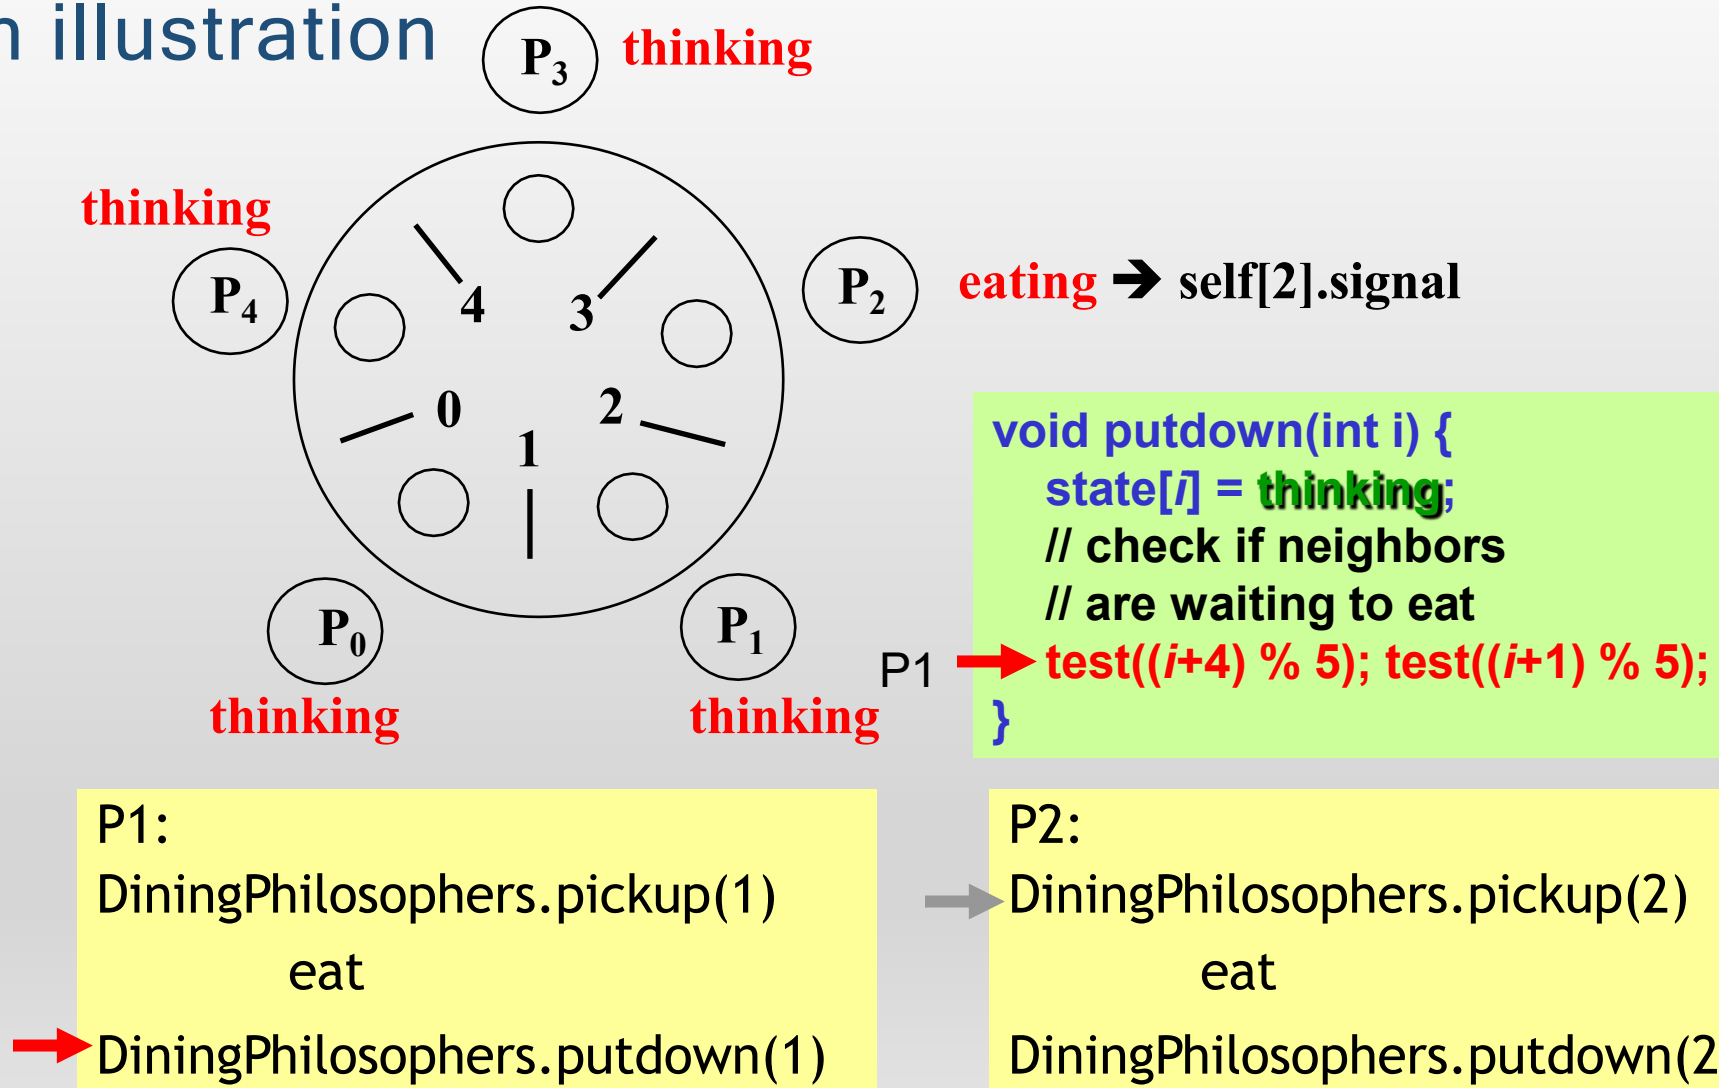
P2:  
 $\rightarrow$  DiningPhilosophers.pickup(2)  
eat  
DiningPhilosophers.putdown(2)

# An illustration






# An illustration



# An illustration



P1:  
DiningPhilosophers.pickup(1)  
eat

→ DiningPhilosophers.putdown(1)

P2:  
DiningPhilosophers.pickup(2)  
eat

→ DiningPhilosophers.putdown(2)

# Synchronized Tools in JAVA

- Synchronized Methods (Monitor)
  - Synchronized method uses the method receiver as a lock
  - Two invocations of synchronized **methods cannot interleave on the same object**
  - When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object block until the first thread exits the object

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() { c++; }  
    public synchronized void decrement() { c--; }  
    public synchronized int value() { return c; }  
}
```

# Synchronized Tools in JAVA

- Synchronized Statement (Mutex Lock)
  - Synchronized blocks uses the **expression** as a lock
  - A synchronized Statement can only be executed once the thread has obtained a **lock for the object or the class that has been referred to in the statement**
  - useful for improving concurrency **with fine-grained**

```
public void run()
{
    synchronized(p1)
    {
        int i = 10; // statement without locking requirement
        p1.display(s1);
    }
}
```

# Review Slides (4)

- Bounded-buffer problem?
- Reader-Writer problem?
- Dining Philosopher problem?
- What is monitor and why need monitor?

# Atomic Transactions

# System Model

- Transaction: a collection of instructions  
(or instructions) that performs a single logic function
- Atomic Transaction: operations happen as a single logical unit of work, in its entirety, or not at all
- Atomic transaction is particular a concern for database system
  - Strong interest to use DB techniques in OS

# File I/O Example

- Transaction is a series of **read** and **write** operations
- Terminated by **commit** (transaction successful) or **abort** (transaction failed) operation
- Aborted transaction must be **rolled back** to undo any changes it performed
- It is part of the responsibility of the system to ensure this property



# Log-Based Recovery

- **Record** to stable storage information about all **modifications by a transaction**
  - **Stable storage**: never lost its stored data
- **Write-ahead logging**: Each log record describes single **transaction write operation**
  - Transaction name
  - Data item name
  - Old & new values
  - Special events: <Ti starts>, <Ti commits>
- Log is used to **reconstruct the state of the data** items modified by the transactions
  - Use **undo (Ti)**, **redo(Ti)** to recover data

# Checkpoints

- When failure occurs, must consult the log to **determine which transactions must be re-done**
  - Searching process is time consuming
  - Redone may not be necessary for all transactions
- Use **checkpoints** to reduce the above overhead:
  - Output all **log records** to stable storage
  - Output all **modified data** to stable storage
  - Output a log record **<checkpoint>** to stable storage

# Review Slides (5)

- What is atomic transaction?
- Purpose of commit, abort, rolled-back?
- How to use log and checkpoints?

# Reading Material & HW

- Chap 6
- HWs
  - 6.1, 6.2, 6.3, 6.4, 6.5, 6.6, 6.9, 6.14, 6.20

# Backup

# Case Study:

- Solaris 2
- Windows XP

# Solaris 2 Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses adaptive mutexes for efficiency when protecting data from short code segments.
  - Mutex and semaphore always serialize data accesses
- Uses condition variables and readers-writers locks when longer sections of code need access to data.
  - Efficient for data that is accessed frequently, but in a read- only manner

# Solaris 2 Adaptive Mutex

- Multiprocessor system
  - Data locked (i.e. in use)
    - Locking thread is running -> requesting thread spins on the mutex (**spinlock**)
    - Locking thread is not in run state -> requesting thread blocks on the mutex (**waiting lock**)
- Uniprocessor system
  - Requesting thread always blocks



# Solaris 2 Turnstile

- Uses **turnstiles** to order the list of threads waiting to acquire either an **adaptive mutex** or **reader-writer lock**
  - A turnstile is a queue structure containing threads blocked on a lock
- To prevent a priority inversion, turnstiles are organized according to a ***priority-inheritance protocol***
  - Temporarily inherit the priority of the high-priority thread (blocked on this lock)



# XP Synchronization

- Use **interrupt masks** to protect access to global resources on **uniprocessor systems (disable interrupt)**
- Uses **spinlocks** on **multiprocessor system**
- Dispatcher objects: either in signaled or nonsignaled state
  - Signaled: object is available immediately
  - Nonsignaled: object is not available
  - Thread queue associated with **each object**
  - **WaitForSingleObject** or **WaitForMultipleObjects**

Q & A

*Thank you for your attention*