



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

Operating System

Lecture 07: Memory Management

Shuo-Han Chen 陳碩漢


shch@nycu.edu.tw

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

Questionnaire of HW01

- Second questionnaire of this semester
- There are some Dcard discussions and feedback on HW01
 - Please let me know more (before 10/10 23:59)
- Our response
 - Create a channel for discussion so that everyone can see the QA
 - Provide in-depth HW explanation to a certain level (may not be up to your expectation)
 - Final grade is for sure to be adjusted
 - Possible act: Relax the late submission penalty



Course Schedule

| W | Date | Lecture | Online | Homework |
|----|----------|---|--------|---------------|
| 1 | Sept. 4 | Lec00: Course Overview & Historical Prospective | | |
| 2 | Sept. 11 | Lec01: Introduction | V | |
| 3 | Sept. 18 | Lec02: OS Structure | V | HW01 Due 10/5 |
| 4 | Sept. 25 | Lec03: Processes Concept | X | |
| 5 | Oct. 2 | Typhoon – No class | V | |
| 6 | Oct. 9 | Lec07: Memory Management | V | HW02 |
| 7 | Oct. 16 | Lec08: Virtual Memory Management | V | |
| 8 | Oct. 23 | Lec04: Process Scheduling | V | |
| 9 | Oct. 30 | School Midterm Exam | | |
| 10 | Nov. 6 | Lec05: Process Synchronization | V | HW03 |
| 11 | Nov. 13 | Lec06: Deadlocks | V | |
| 12 | Nov. 20 | School Event – No class | | |
| 13 | Nov. 27 | Lec09: File System Interface | V | HW04 |
| 14 | Dec. 4 | Lec10: File System Implementation | V | |
| 15 | Dec. 11 | Lec11: Mass Storage System & Lec12: IO Systems | V | |
| 16 | Dec. 18 | School Final Exam | | |

Overview

- Background
- Swapping
- Contiguous Allocation
- Paging
- Segmentation
- Segmentation with Paging

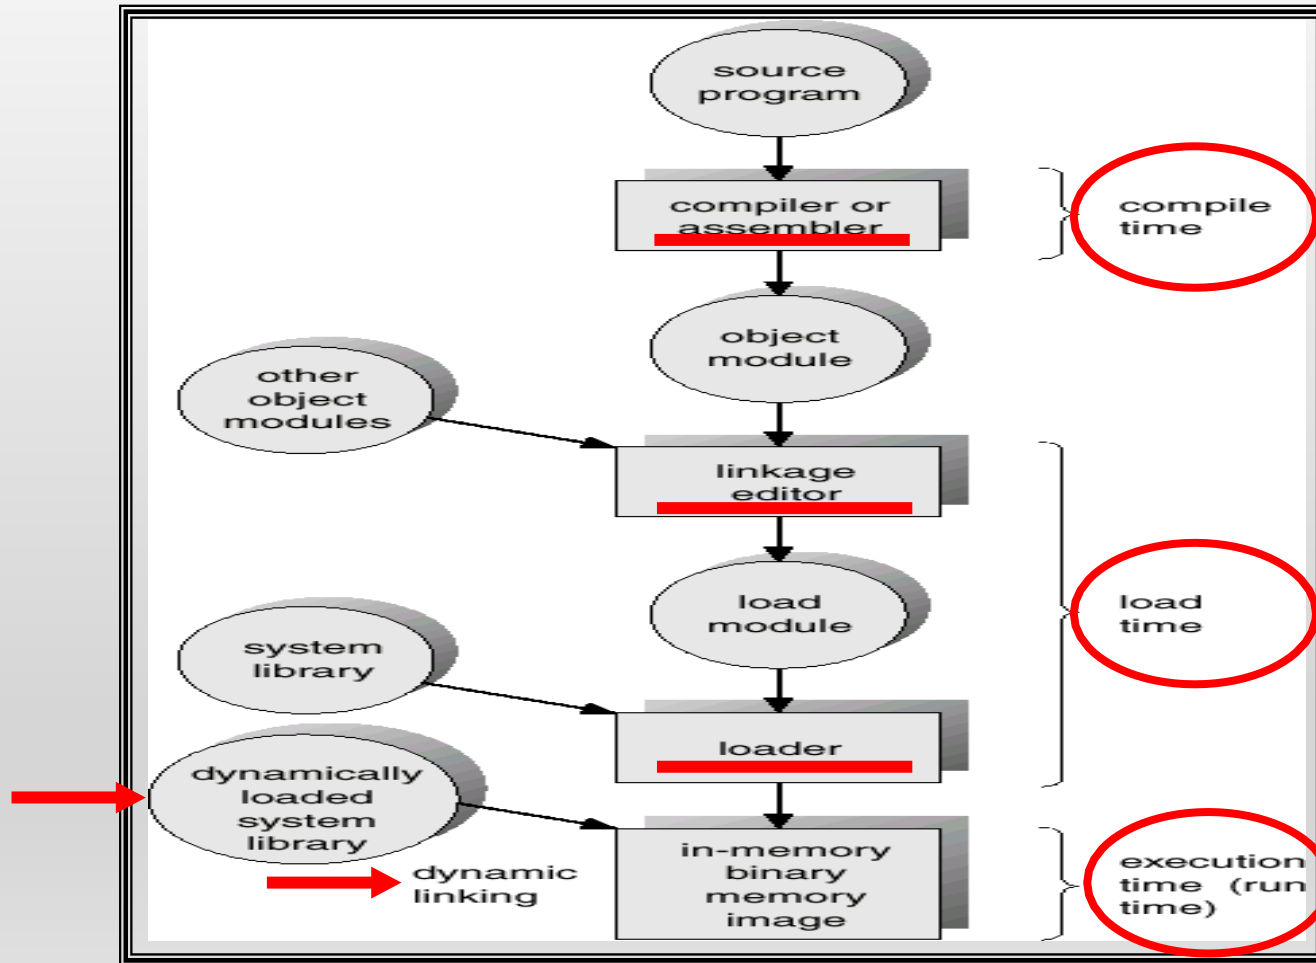
Background

- Main memory and registers are the only storage CPU can access directly
- Collection of **processes** are waiting on disk to be **brought into memory and be executed**
- **Multiple programs are brought into memory** to improve resource utilization and response time to users
- A process may **be moved between disk and memory** during its execution

Outline

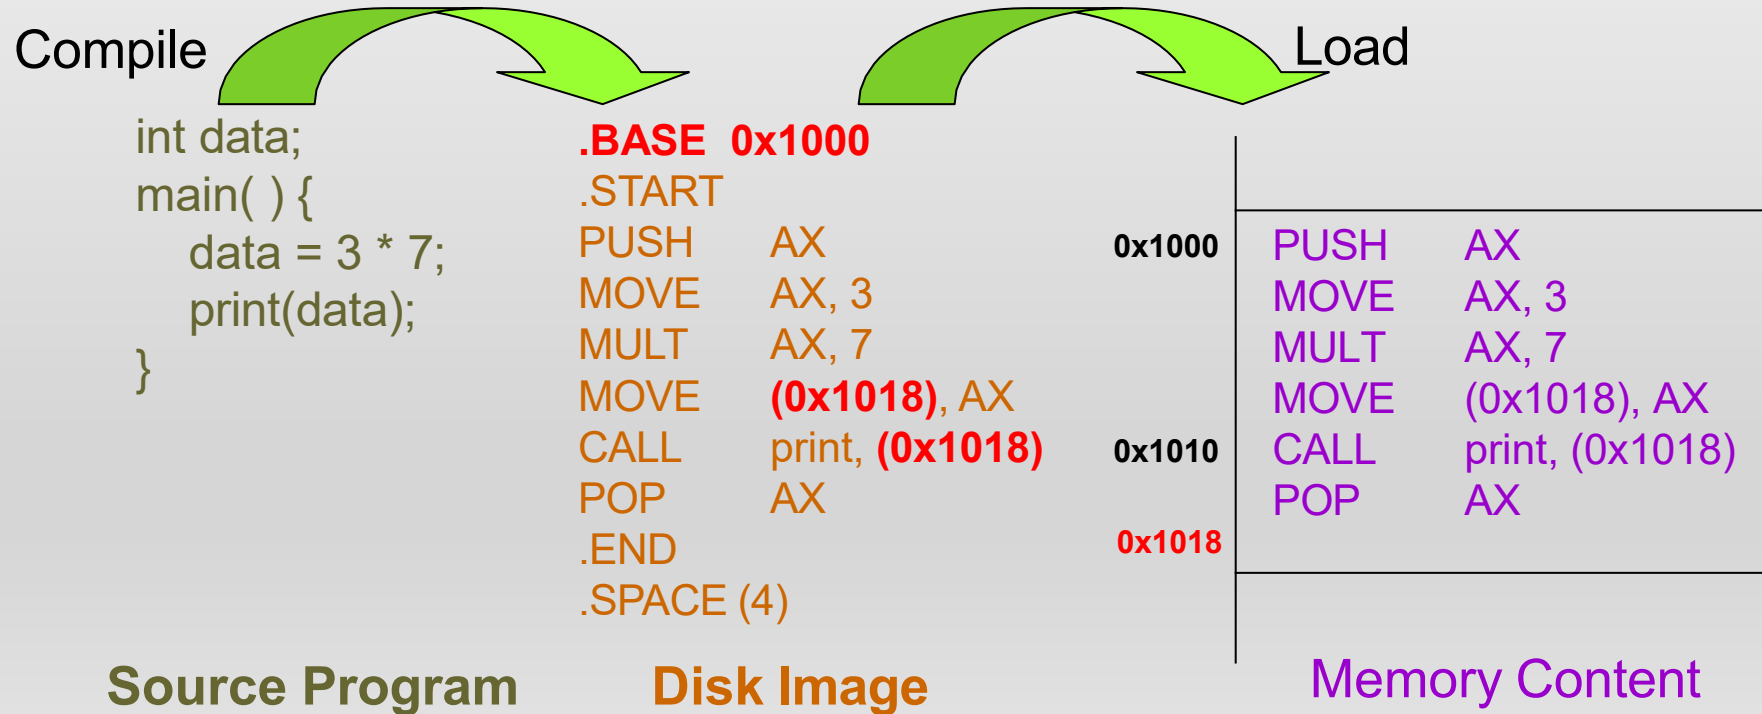
- How to refer memory in a program?
 - address binding
- How to load a program into memory ?
 - static/dynamic loading and linking
- How to move a program between mem. & disk?
 - swap
- How to allocate memory?
 - paging, segment

Multistep Processing of a User Program



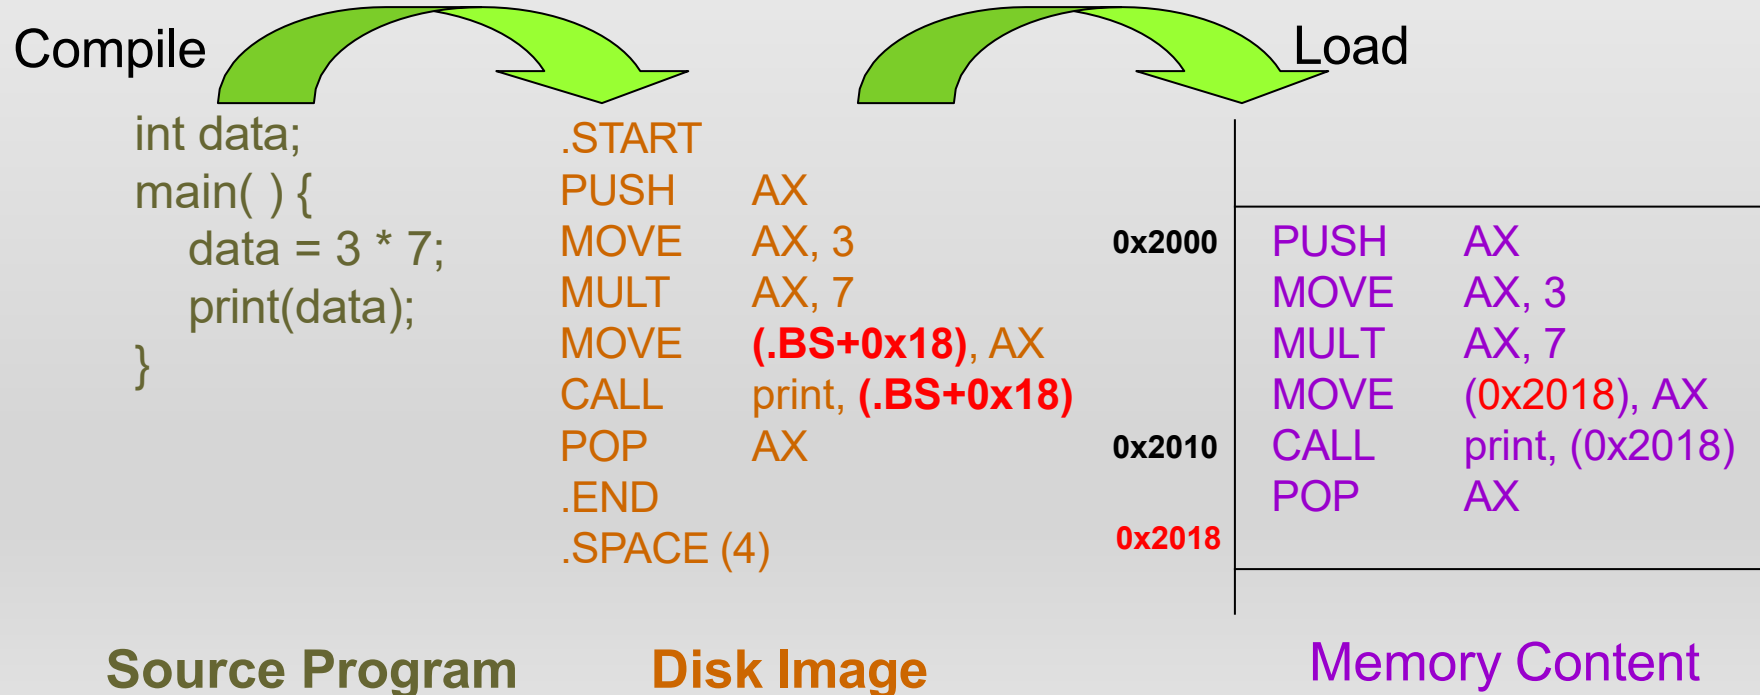
Address Binding - Compile Time

- Program is written as symbolic code
- Compiler translates symbolic code into *absolute code*
- If starting location changes -> **recompile**
- Example: MS-DOS .COM format binary




Address Binding - Load Time

- Compiler translates symbolic code into *relocatable code*
- **Relocatable code**:
 - Machine language that can be run from any memory location
- If starting location -> **reload the code**



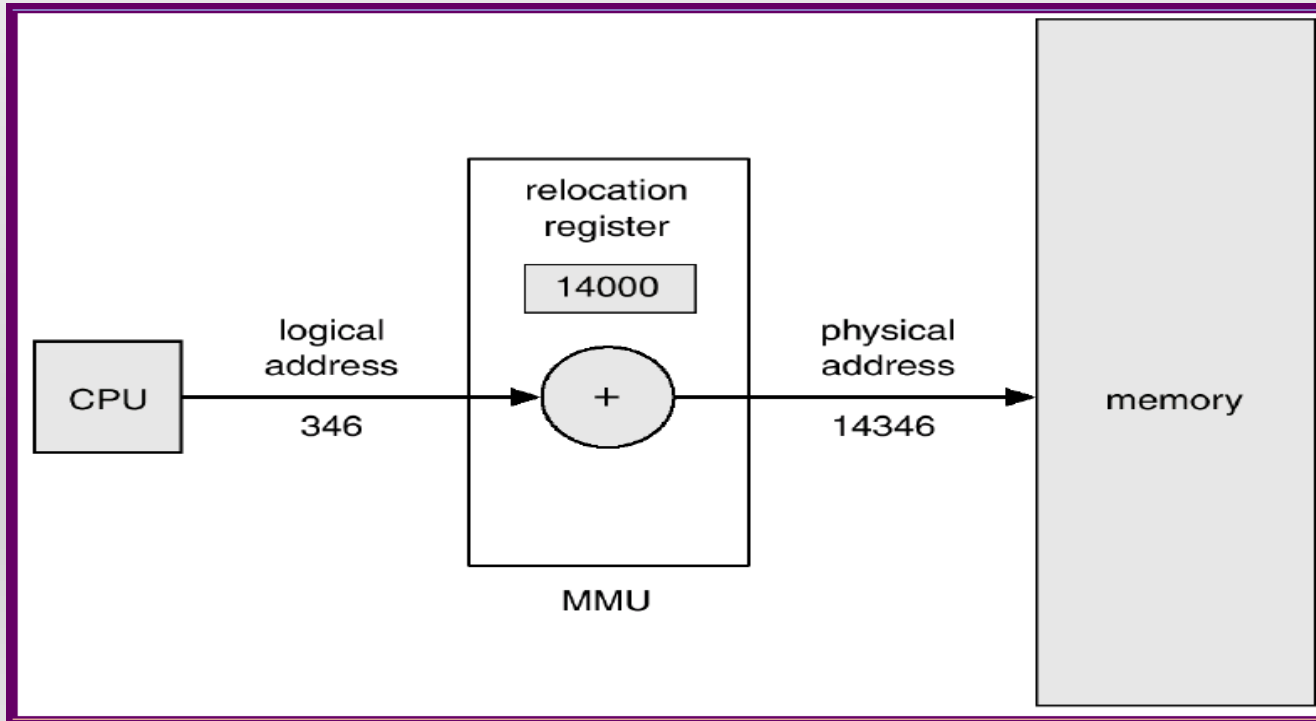
Address Binding - Execution Time

- Compiler translates symbolic code into *logical-address (i.e. virtual-address) code*
- Special **hardware (i.e. MMU)** is needed for this scheme
- Most general-purpose OS use this method



Memory-Management Unit (MMU)

- Hardware device that **maps virtual to physical address**
- The value in the **relocation register** is added to every address generated by a user process at the time it is sent to memory



Logical vs. Physical Address

- Logical address – generated by CPU
 - a.k.a. virtual address
- Physical address – seen by the memory module
- compile-time & load-time address binding
 - logical addr = physical addr
- Execution-time address binding
 - logical addr \neq physical addr
- The user program deals with logical addresses; it never sees the real physical addresses

Outline

- How to refer memory in a program?
 - address binding
- How to load a program into memory?
 - static/dynamic loading and linking
- How to move a program between mem. & disk?
 - swap
- How to allocate memory?
 - paging, segment

Dynamic Loading

- The entire program must be in memory for it to execute?
- No, we can use **dynamic-loading**
 - A routine is loaded into memory when it is called
- **Better memory-space utilization**
 - unused routine is never loaded
 - Particularly useful when large amounts of code are infrequently used (e.g., **error handling code**)
- **No special support from OS** is required implemented through program (library, API calls)

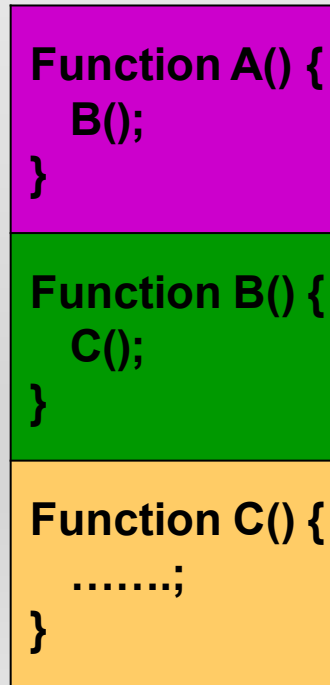
Dynamic Loading Example in C

- `dlopen()`: opens a library and prepares it for use
- `desym()`: looks up the value of a symbol in a given (opened) library.
- `dlclose()`: closes a DL library

```
#include <dlfcn.h>
int main() {
    double (*cosine)(double);
    void* handle = dlopen ("/lib/libm.so.6", RTLD_LAZY);
    cosine = dlsym(handle, "cos");
    printf ("%f\n", (*cosine)(2.0));
    dlclose(handle);
}
```


Dynamic Loading

Disk image

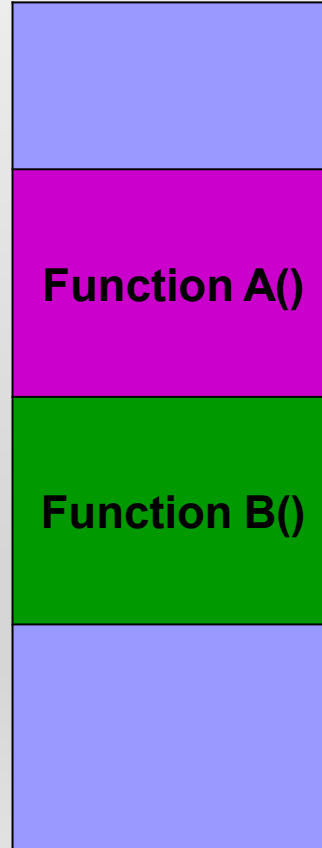


Memory content

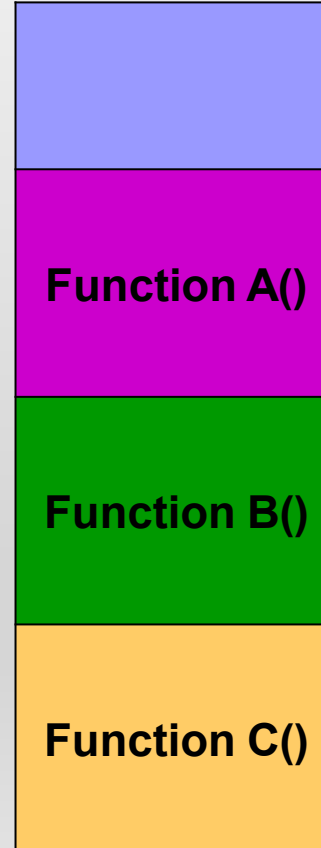
Init



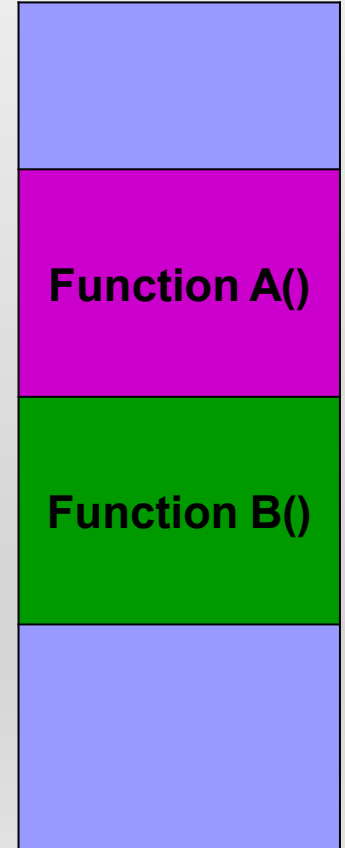
After B() called



After C() called

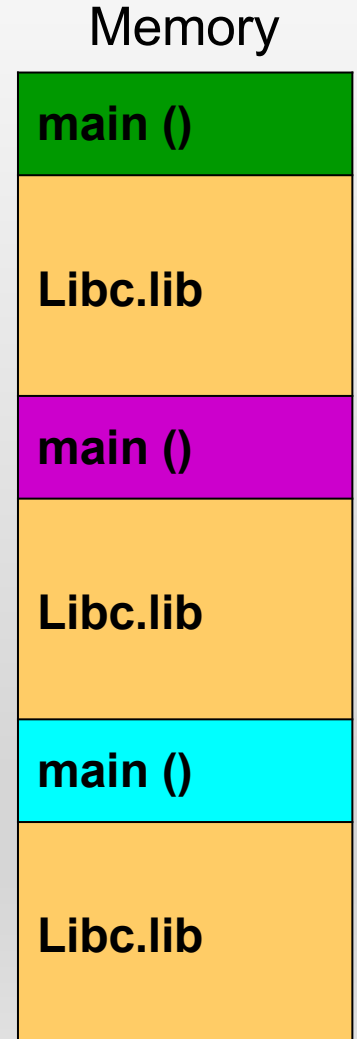
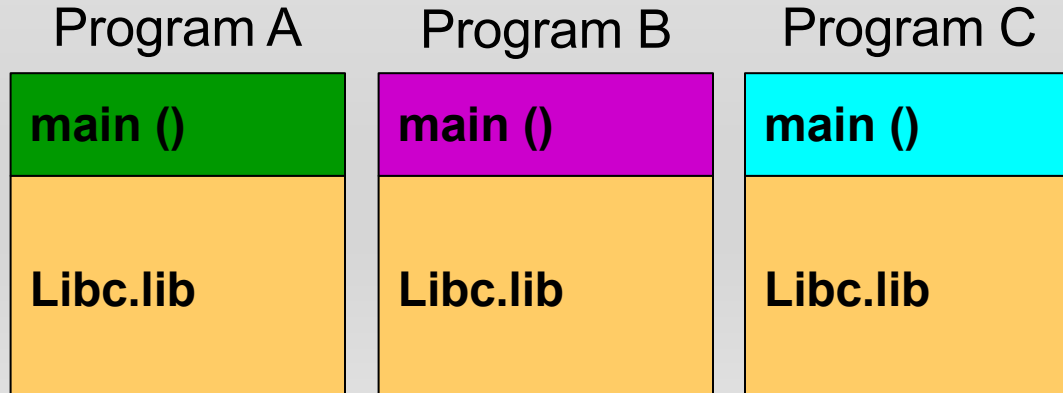


After C() ends




Static Linking

- *Static linking* : libraries are combined by the **loader** into the program in-memory image
 - Waste memory: **duplicate code**
 - Faster during execution time
- ***Static linking + Dynamic loading**
 - Still can't prevent duplicated code



Dynamic Linking

- *Dynamic linking*: Linking postponed **until execution time**
 - **Only one code copy** in memory and **shared by everyone**
 - A stub is included in the program in-memory image for each lib reference
 - Stub call -> check if the referred lib is in memory -
> if not, load the lib -> execute the lib
 - DLL (Dynamic link library) on Windows



Review Slides (1)

- 3 types of address binding?
 - compile-time
 - load-time
 - execution-time
- logical address? physical address?
 - virtual -> physical mapping?
- dynamic loading? static loading?
- dynamic linking? static linking?

Outline

- How to refer memory in a program?
 - address binding
- How to load a program into memory?
 - static/dynamic loading and linking
- How to move a program between mem. & disk?
 - swap
- How to allocate memory?
 - paging, segment

Swapping


- A process can be swapped out of memory to a **backing store**, and later brought back into memory for continuous execution
 - Also used by **midterm scheduling**, different from context switch
- **Backing store** – a chunk of **disk**, **separated from file system**, to provide direct access to these memory images
- Why Swap a process:
 - Free up memory
 - Roll out, roll in: swap lower-priority process with a higher one

Swapping (cont'd)

- Swap back memory location
 - If binding is done at compile/load time
 - -> swap back memory address must be the **same**
 - If binding is done at execution time
 - -> swap back memory address can be **different**
- A process to be swapped == **must be idle**
 - Imagine a process that is waiting for I/O is swapped
 - Solutions:
 - Never swap a process with pending I/O
 - I/O operations are done through OS buffers (i.e. a memory space not belongs to any user processes)

Process Swapping to Backing Store

- A major part of swap time is transfer time; **total transfer time is directly proportional to the amount of memory swapped**



Outline

- How to refer memory in a program?
 - address binding
- How to load a program into memory?
 - static/dynamic loading and linking
- How to move a program between mem. & disk?
 - swap
- How to allocate memory?
 - paging, segment

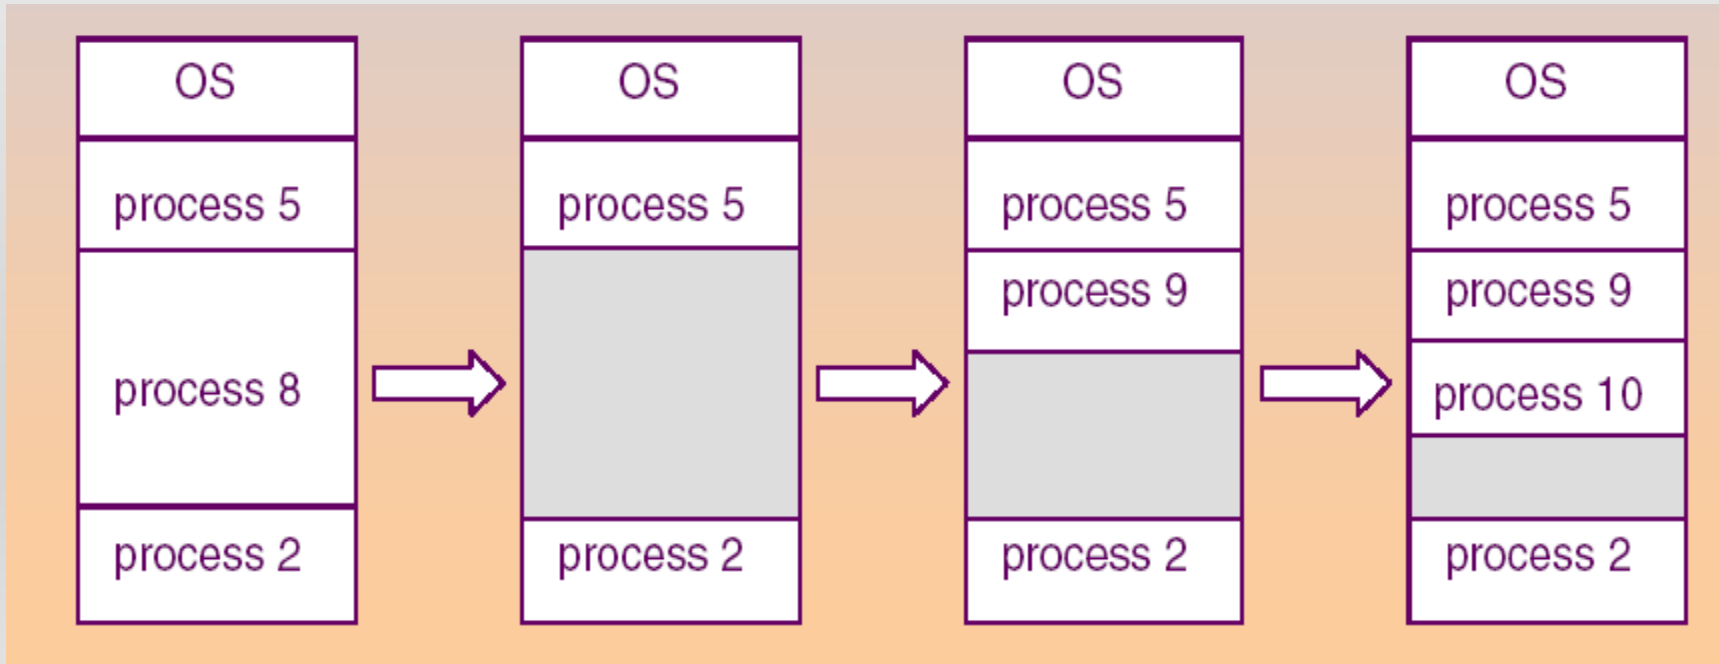
Contiguous Memory Allocation

Memory Allocation

- **Fixed-partition** allocation:
 - Each process loads into one partition of fixed-size
 - **Degree of multi-programming** is bounded by the number of partitions
- **Variable-size** partition
 - Hole: block of contiguous free memory
 - Holes of various sizes are scattered in memory

Multiple Partition (Variable-Size) Method

- When a process arrives, it is allocated a hole large enough to accommodate it
- The OS maintains info. on each **in-use** and **free-hole**
- A freed hole can be merged with another hole to form a larger hole



Dynamic Storage Allocation Problem

- How to satisfy a request of size n from a list of free holes
- **First-fit** – allocate the 1st hole that fits
- **Best-fit** – allocate the smallest hole that fits
 - Must search through the whole list
- **Worst-fit** – allocate the largest hole
 - Must also search through the whole list
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Fragmentation

- *External fragmentation*


- Total free memory space is big enough to satisfy a request, but is not contiguous
- Occur in variable-size allocation

- *Internal fragmentation*


- Memory that is internal to a partition but is not being used
- Occur in fixed-partition allocation

- Solution: *compaction*

- Shuffle the memory contents to place all free memory together in one large block at **execution time**
- Only if the binding is done at execution time



Compaction



Review Slides (2)

- Swapping?
- Contiguous memory allocation?
 - fixed-size memory allocation?
 - variable-size memory allocation?
 - first-fit, best-fit, worst-fit?
- External & internal fragmentation?
 - compaction?


Non-Contiguous Memory Allocation – Paging

Paging Concept

- Method:
 - Divide **physical memory** into fixed-sized blocks called **frames**
 - Divide **logical address space** into blocks of the **same size** called **pages**
 - To run a program of **n** pages, need to find n free frames and load the program
 - **keep track of free frames**
 - Set up a **page table** to translate logical to physical addresses
- Benefit:
 - Allow the **physical address space** of a process to be **noncontiguous**
 - Avoid external fragmentation
 - Limited internal fragmentation
 - Provide **shared memory/pages**

Paging Example

- Page table:
 - Each entry maps to the **base address of a page** in physical memory
 - A structure maintained by OS **for each process**
 - Page table includes only pages owned by a process
 - A process cannot access memory outside its space

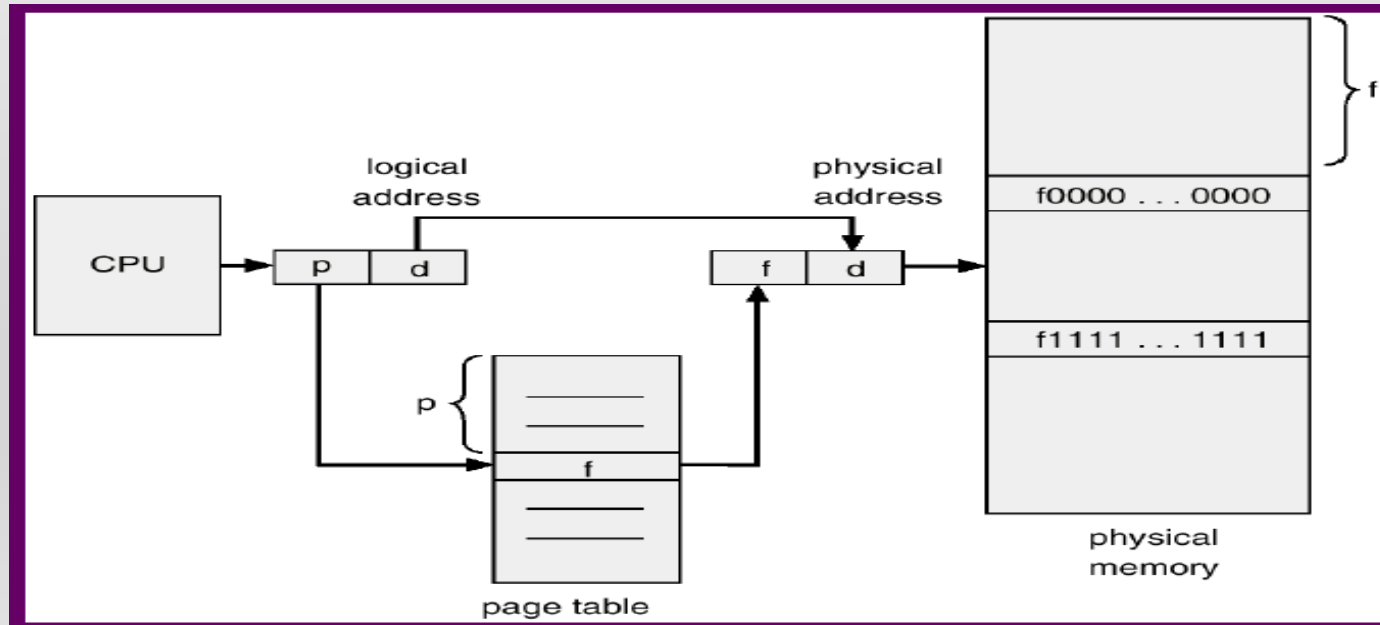


Address Translation Scheme

- Logical address is divided into two parts:
 - **Page number (p)**
 - used as an **index into a page table** which contains **base address of each page** in physical memory
 - N bits means a **process** can allocate **at most 2^N pages**
 - **$\rightarrow 2^N \times$ page size memory size**
 - **Page offset (d)**
 - combined with the base address to define the physical memory address that is sent to the memory unit
 - N bits means the **page size** is 2^N
 - Ex. 4 KiB = 4096 B = 2^{12}
- Physical addr = page **base addr** + page **offset**

Address Translation Architecture


- If Page size is 1KB(2^{10}) & Page 2 maps to frame 5
- Given 13 bits logical address: ($p=2, d=20$), what is physical addr.?
 - $5 \times (1\text{KB}) + 20 = 1,010,000,000,000 + 0,000,010,100$
 - $= 1,010,000,010,100$



Address Translation

- Total number of pages does not need to be the same as the total number of frames
 - Total # pages determines the logical memory size of a process
 - Total # frames depending on the size of physical memory
- E.g.: Given 32 bits logical address, 36 bits physical address, and 4KB page size, what does it mean?
 - Page table size: $2^{32} / 2^{12} = 2^{20}$ entries
 - Max program memory: $2^{32} = 4$ GiB
 - Total physical memory size: $2^{36} = 64$ GiB
 - Number of bits for page number: 2^{20} pages -> 20bits
 - Number of bits for frame number: 2^{24} frames -> 24bits
 - Number of bits for page offset: 4KB page size = 2^{12} bytes -> 12

Free Frames



Page / Frame Size

- The page (frame) size is defined by the hardware
 - Typically a power of 2
 - Ranging from 512 bytes to 16MB / page
 - 4KB / 8KB page is commonly used
- Internal fragmentation?
 - Larger page size -> More space waste
- But page sizes have grown over time
 - memory, process, and data sets have become larger
 - Better I/O performance (during page fault)
 - The page table is smaller

Paging Summary

- Paging helps separate **user's view of memory** and the actual **physical memory**
- User view's memory: one single contiguous space
 - Actually, user's memory is scatter out in physical memory
- OS maintains a copy of the **page table for each process**
- OS maintains a **frame table** for managing physical memory
 - One entry for each physical frame
 - Indicate whether a frame is free or allocated
 - If allocated, to which page of which process or processes

Implementation of Page Table

- Page table is kept in memory
- Page-table base register (PTBR)
 - The physical memory address of the page table
 - The PTBR value is stored in PCB (Process Control Block)
 - Changing the value of PTBR during Context-switch
- With PTBR, each memory reference results in 2 memory reads
 - One for the page table and one for the real address
- The 2-access problem can be solved by
 - Translation Look-aside Buffers (TLB) (HW) which is implemented by Associative memory (HW)

Associative Memory

- All memory entries can be accessed at the same time
 - Each entry corresponds to an associative register
- But number of **entries are limited**
 - Typical number of entries: 64 ~ 1024

■ Associative memory – parallel search

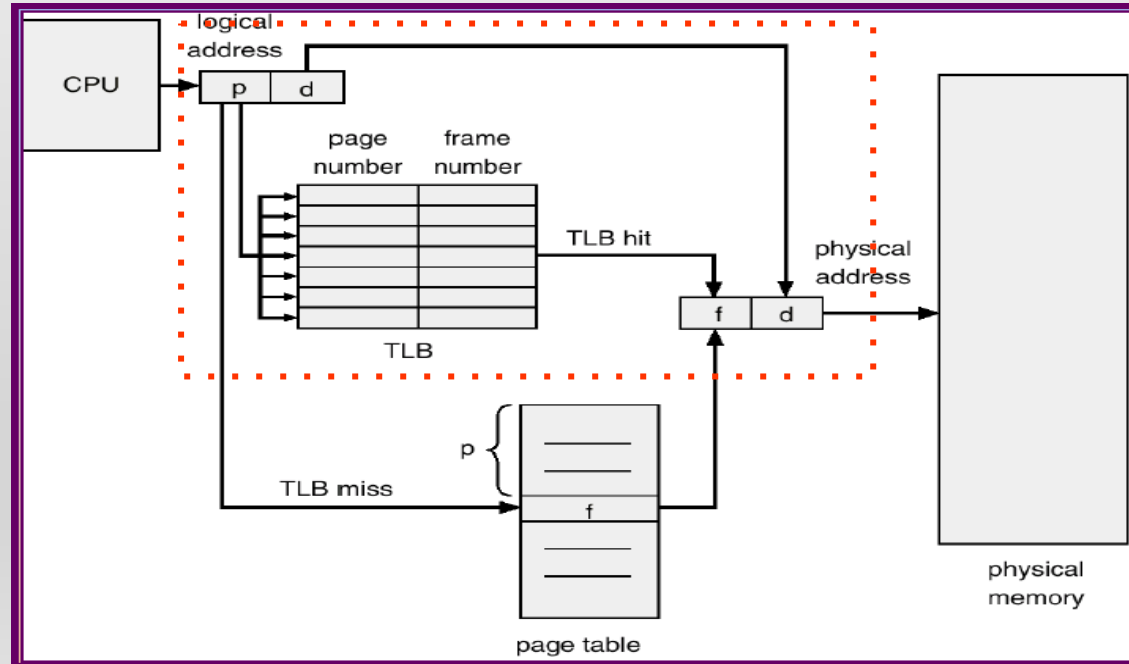
| Page # | Frame # |
|--------|---------|
| | |
| | |
| | |
| | |

Address translation (A' , A'')

- ◆ If A' is in associative register, get frame # out.
- ◆ Otherwise get frame # from page table in memory

Translation Look-aside Buffer (TLB)

- A cache for page table shared by all processes
- TLB must be flushed after a context switch
 - Otherwise, TLB entry must have a PID field (address-space identifiers (ASIDs))



Effective Memory-Access Time

- 20 ns for TLB search
- 100 ns for memory access
- Effective Memory-Access Time (EMAT)
 - 70% TLB hit-ratio:
 - $EMAT = 0.70 \times (20 + 100) + (1-0.70) \times (20+100+100) = 150 \text{ ns}$
 - 98% TLB hit-ratio
 - $EMAT = 0.98 \times 120 + 0.02 \times 220 = 122 \text{ ns}$

Review Slides (3)


- memory frame? page? typical page size?
- page table? virtual -> physical translation?
- What is PTBR register? When to update it?
- Memory reads # for each reference?
- HW support for paging speed?
 - associative memory
 - TLB

Memory Protection

- Each page is associated with a set of **protection bits** in the page table
 - E.g., a bit to define read/write/execution permission
- Common use: **valid-invalid bit**
 - Valid: the page/frame is **in the process logical address space**, and is thus a legal page
 - Invalid: the page/frame is not in the process logical address space

Valid-Invalid Bit Example

- Potential issues:
 - Un-used page entry cause memory waste -> use page table length register (PTLR)
 - Process memory may NOT be on the boundary of a page -> memory limit register is still needed

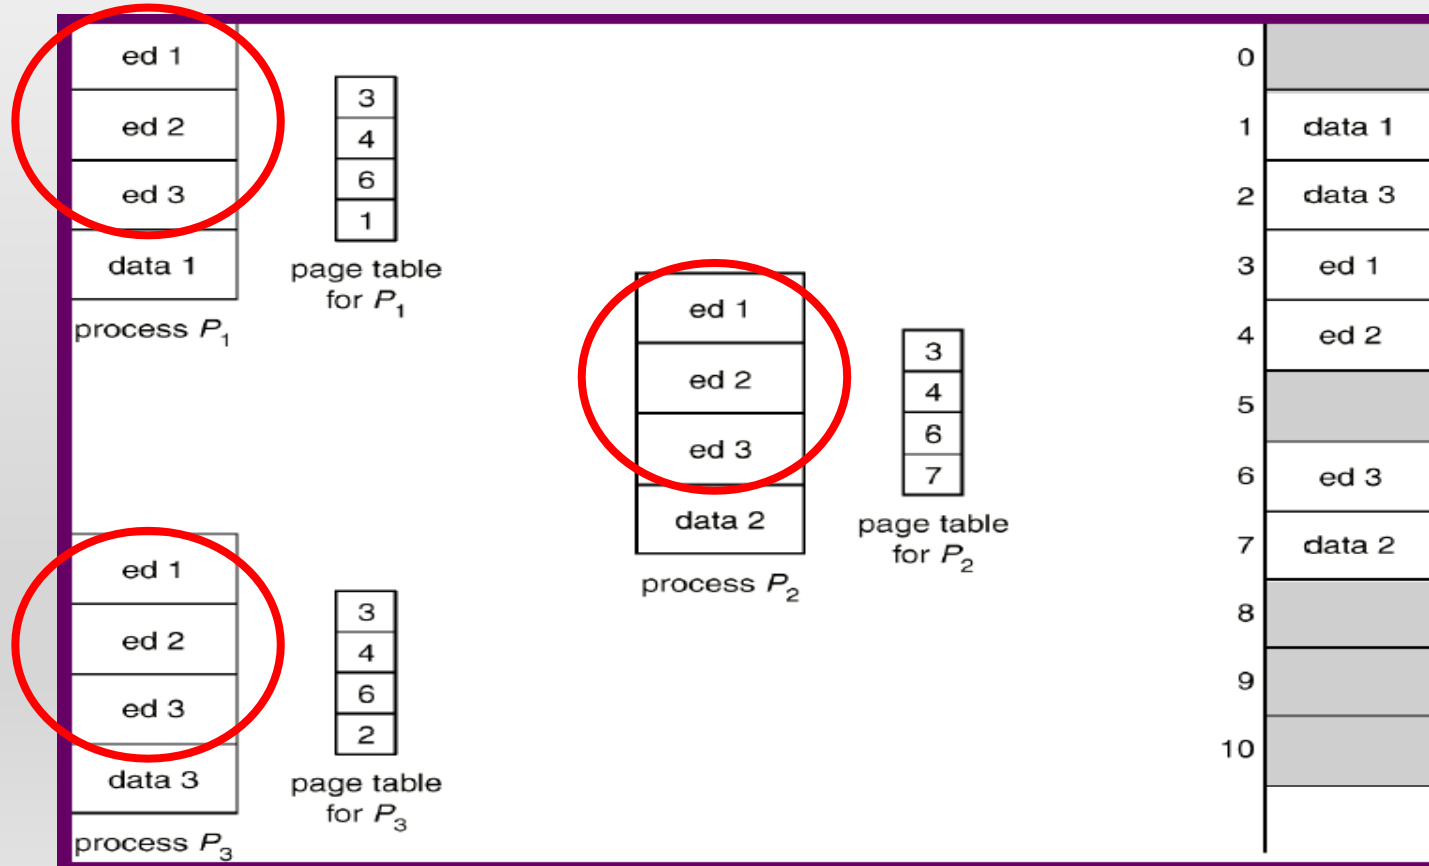


Shared Pages

- Paging allows processes to **share common code**, which must be **reentrant**
- **Reentrant code** (pure code)
 - It never changes during execution
 - text editors, compilers, web servers, etc
- **Only one copy** of the shared code needs to be kept in physical memory
- **Two (several) virtual addresses** are mapped to one physical address
- Process keeps a copy of its own private data and code

Shared Pages by Page Table

- Shared code must appear in the same location in the logical address space of all processes




Page Table Memory Structure

- Page table could be huge and difficult to load
 - 4GB (2^{32}) logical address space with 4KB (2^{12}) page
 - -> 1 million (2^{20}) page table entry
 - Assume each entry need 4 bytes (32bits)
 - -> Total size=4MB
 - Need to break it into several smaller page tables, better within a single page size (i.e. 4KB)
 - Or reduce the total size of page table
- Solutions:
 - Hierarchical Paging
 - Hash Page Tables
 - Inverted Page Table

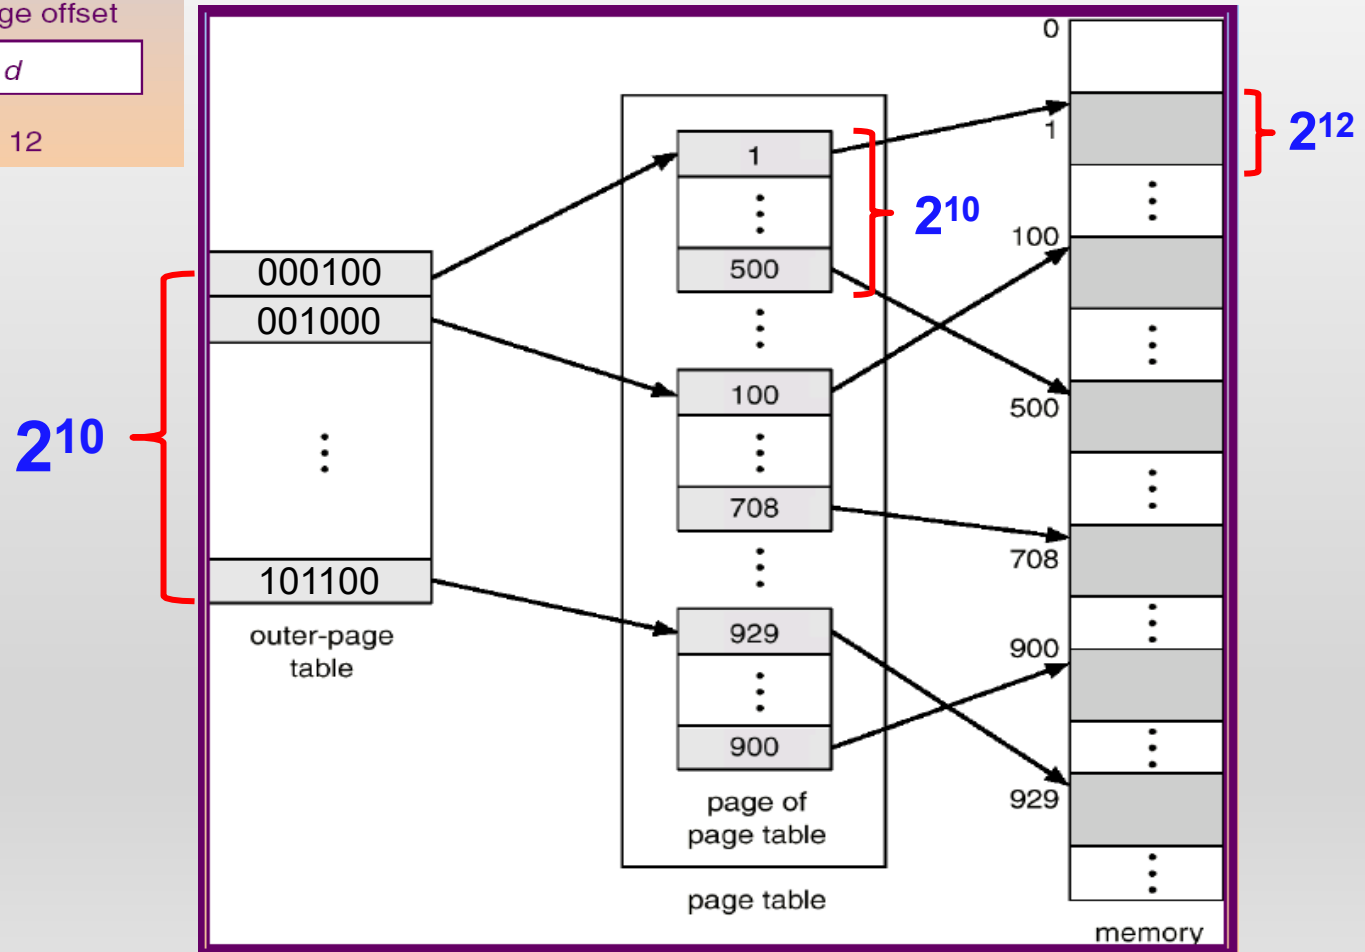
Hierarchical Paging

- Break up the **logical address space** into **multiple page tables**
 - Paged the page table
 - i.e. **n-level** page table
- Two-level paging (32-bit address with 4KB (2^{12}) page size)
- 12-bit offset (d) -> 4KB (2^{12}) page size
- 10-bit outer page number -> 1K (2^{10}) page table entries
- 10-bit inner page number -> 1K (2^{10}) page table entries
- **3 memory accesses**




Two-Level Page Table Example

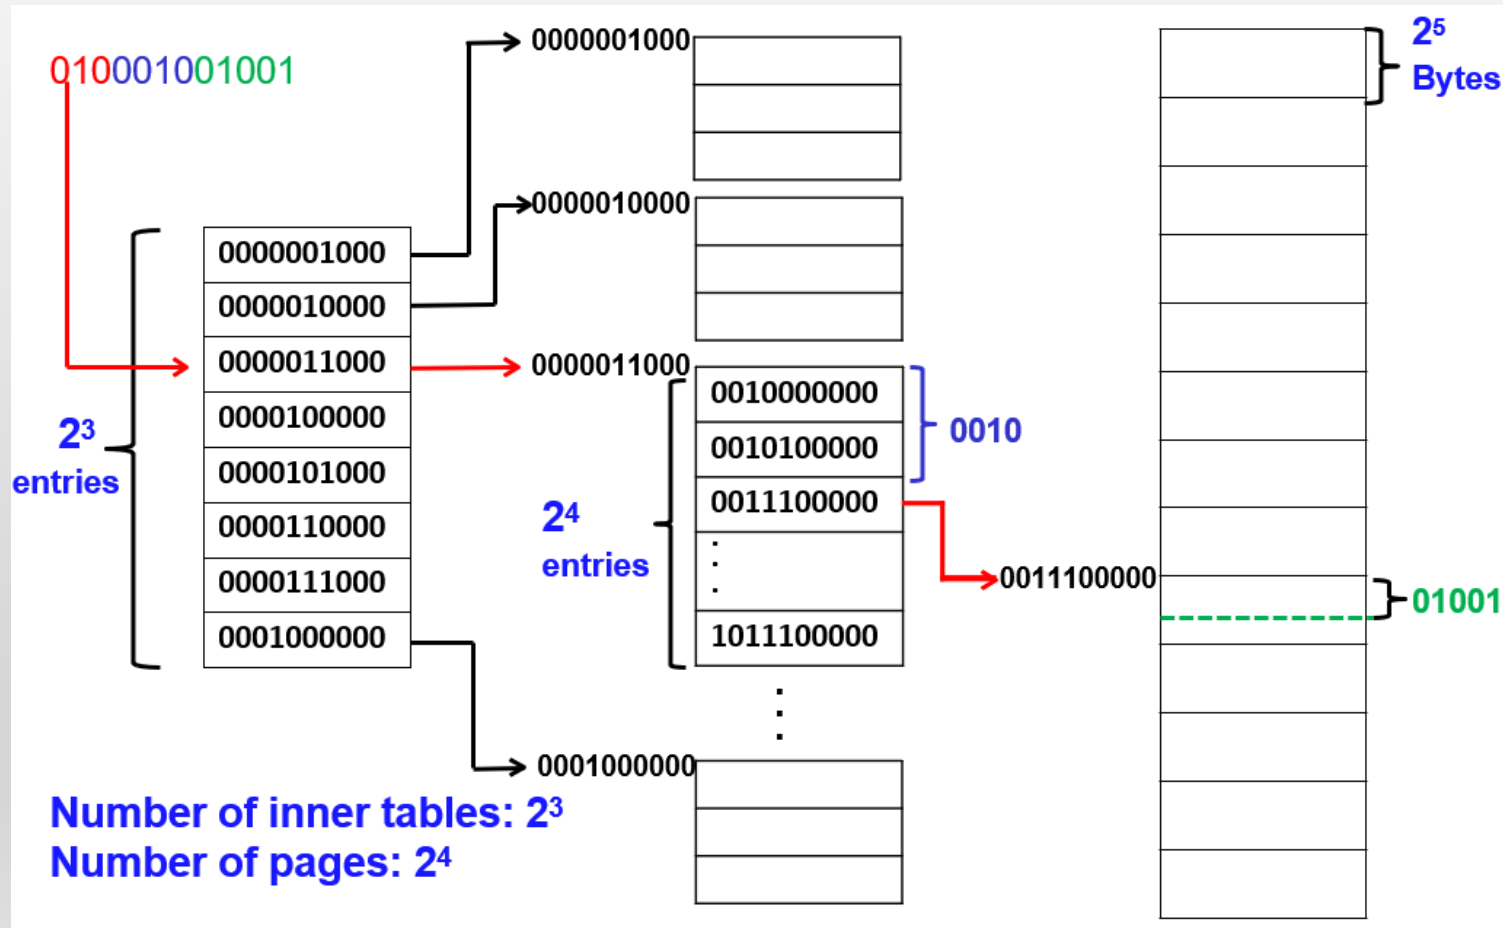
| page number | | page offset |
|-------------|-------|-------------|
| p_i | p_2 | d |
| 10 | 10 | 12 |



Two-Level Address Translation



Two-Level Page Table Translation Example




64-bit Address

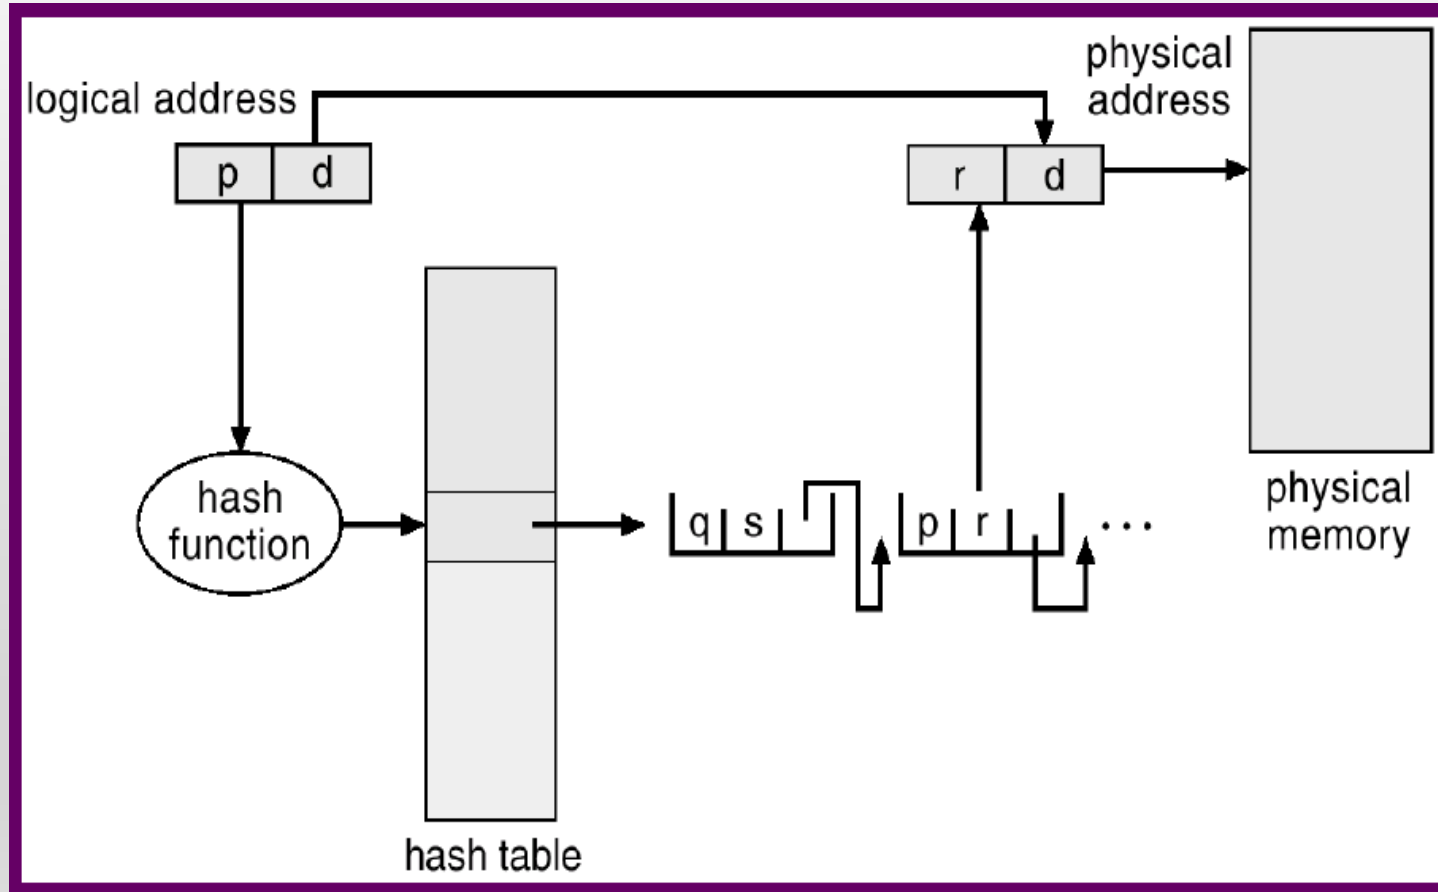
- How about 64-bit address? (assume each entry needs 4Bytes)
 - 42 (p1) + 10 (p2) + 12 (offset)
 - -> outer table requires $2^{42} \times 4B = 16TB$ contiguous memory!!!
 - 12 (p1)+10 (p2)+10 (p3)+10 (p4)+10 (p5)+12 (offset)
 - -> outer table requires $2^{12} \times 4B = 16KB$ contiguous memory
 - -> 6 memory accesses!!!
- Examples:
 - SPARC (32-bit) and Linux use 3-level paging
 - Motorola 68030 (32-bit) use 4-level paging

Hashed Page Table

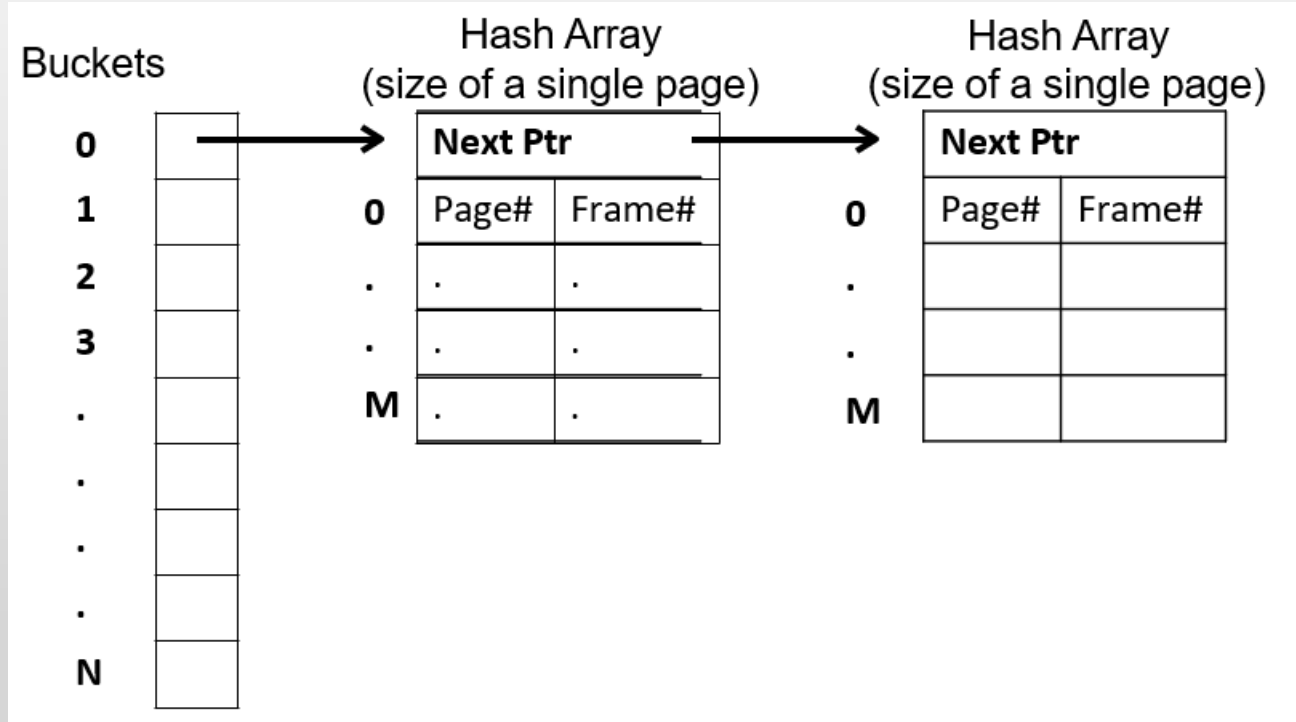
- Commonly-used for address > 32 bits
- Virtual page number is hashed into a **hash table**
- The size of the hash table varies
 - Larger hash table -> smaller chains in each entry
- Each entry in the hashed table contains
 - **(Virtual Page Number, Frame Number, Next Pointer)**
 - Pointers waste memory
 - Traverse linked list waste time & cause additional memory references



Hashed Page Table Address Translation




Improved Hashed Page Table Implementation



Inverted Page Table

- Maintains NO page table for each process
- Maintains a frame table for the whole memory
 - One entry for each real frame of memory
- Each entry in the frame table has
 - (PID, Page Number)
- Eliminate the memory needed for page tables but increase memory access time
 - Each access needs to search the whole frame table
 - Solution: use hashing for the frame table
- Hard to support shared page/memory

Inverted Page Table Addr Translation




Review Slides (4)

- memory protection by page table?
 - valid, invalid bits?
- page table memory structure?
 - hierarchical -> 2-level, 3-level, etc
 - hash table -> linked list
 - inverted page table
- How are pages shared by different processes?


Non-Contiguous Memory Allocation – Segmentation

Segmentation

- Memory-management scheme that supports user **view of memory**
- A program is a collection of segments. A segment is a logical unit such as:
 - main program
 - function, object
 - local/global variables,
 - stack, symbol table,
 - arrays, etc...



Logical View of Segmentation

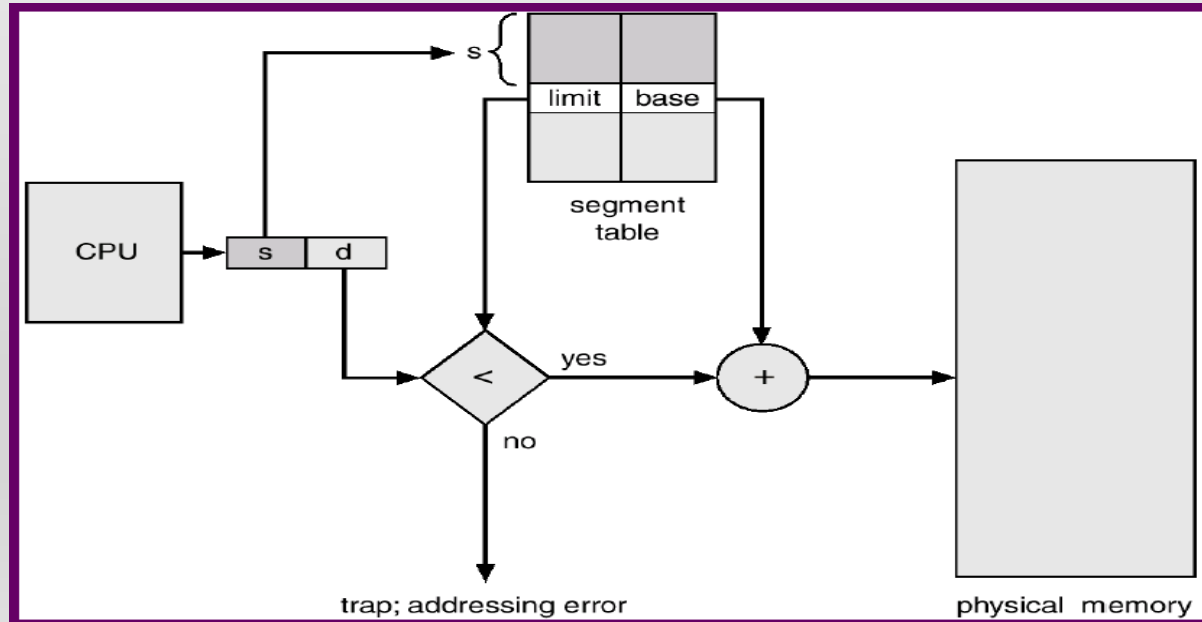


Segmentation Table

- Logical address: (seg#, offset)
 - Offset has the SAME length as physical addr.
- **Segmentation table** – maps two-dimensional physical addresses; each table entry has:
 - **Base (4 bytes)**: the start physical addr
 - **Limit (4 bytes)**: the length of the segment
- Segment-table base register (STBR):
 - the **physical addr** of the segmentation table
- Segment-table length register (STLR):
 - the # of segments

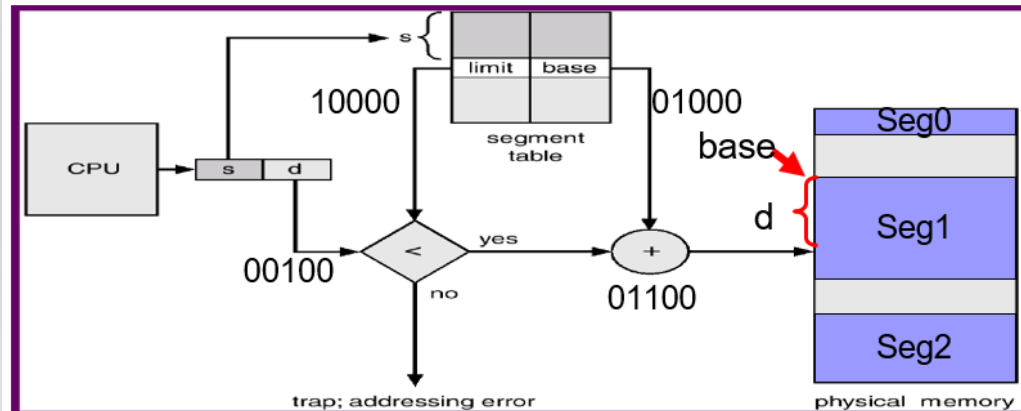
Segmentation Hardware

- Limit register is used to check offset length
- MMU allocate memory by assigning an appropriate **base address for each segment**
 - Physical address cannot overlap between segments

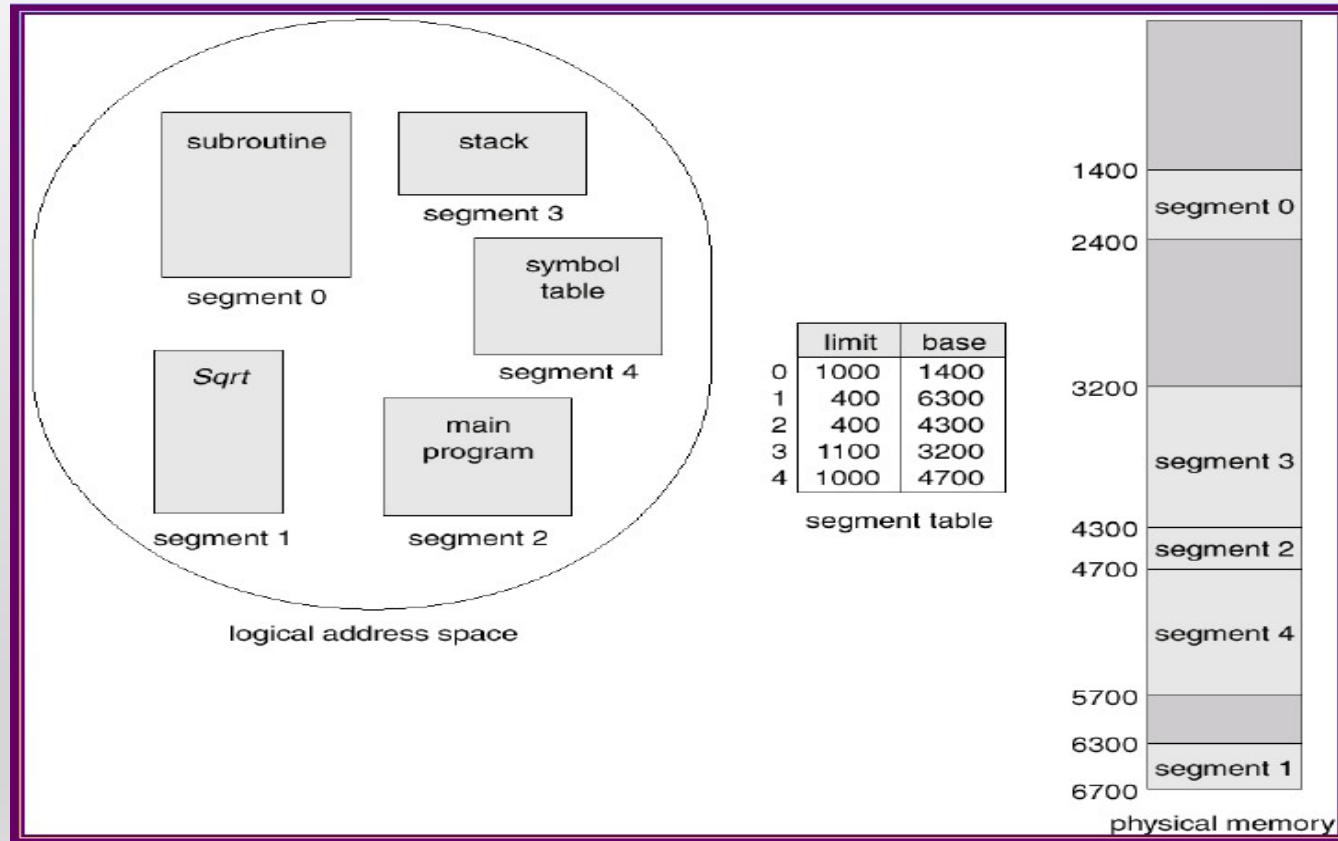


Address Translation Comparison

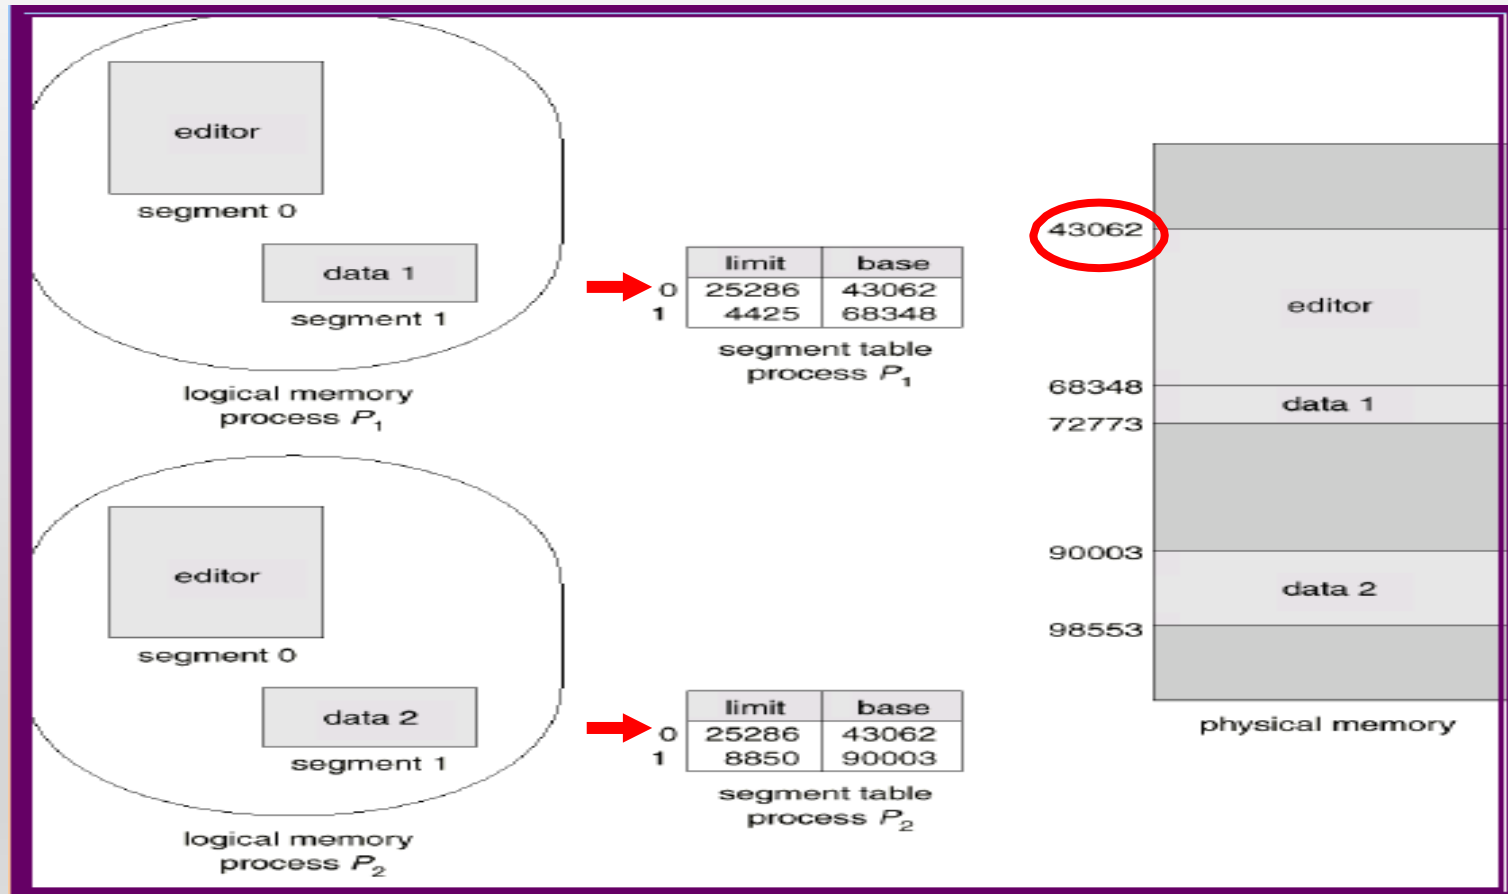
- Segment
 - Table entry: (segment base addr. , limit)
 - Segment base addr. can be arbitrary
 - The length of “offset” is the same as the physical memory size
- Page:
 - Table entry: (frame base addr.)
 - Frame base addr. = frame number * page size
 - The length of “offset” is the same as page size



Example of Segmentation



Sharing of Segments



Protection & Sharing

- Protection bits associated with segments
 - Read-only segment (code)
 - Read-write segments (data, heap, stack)
- Code sharing occurs at **segment level**
 - **Shared memory communication**
 - **Shared library**
- Share segment by having same base in two segment tables

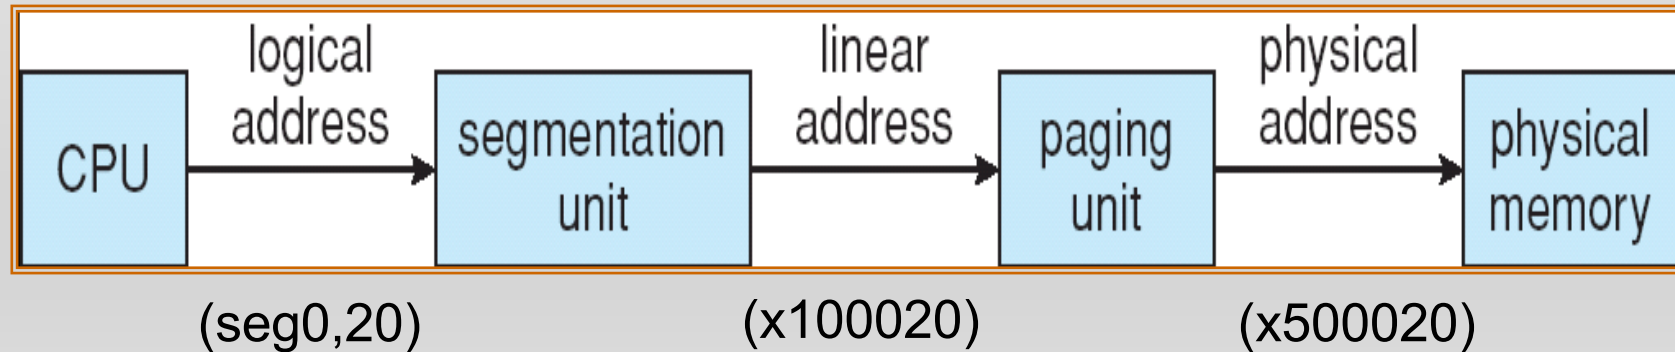
Segmentation with Paging

- Apply **segmentation in logical** address space
- Apply **paging in the physical** address space




Address Translation

- CPU generates logical address
 - Given to **segmentation unit**
 - -> produces **linear addresses**
 - Linear address given to **paging unit**
 - -> generates **physical address** in main memory
- Segmentation and paging units form the equivalent of MMU



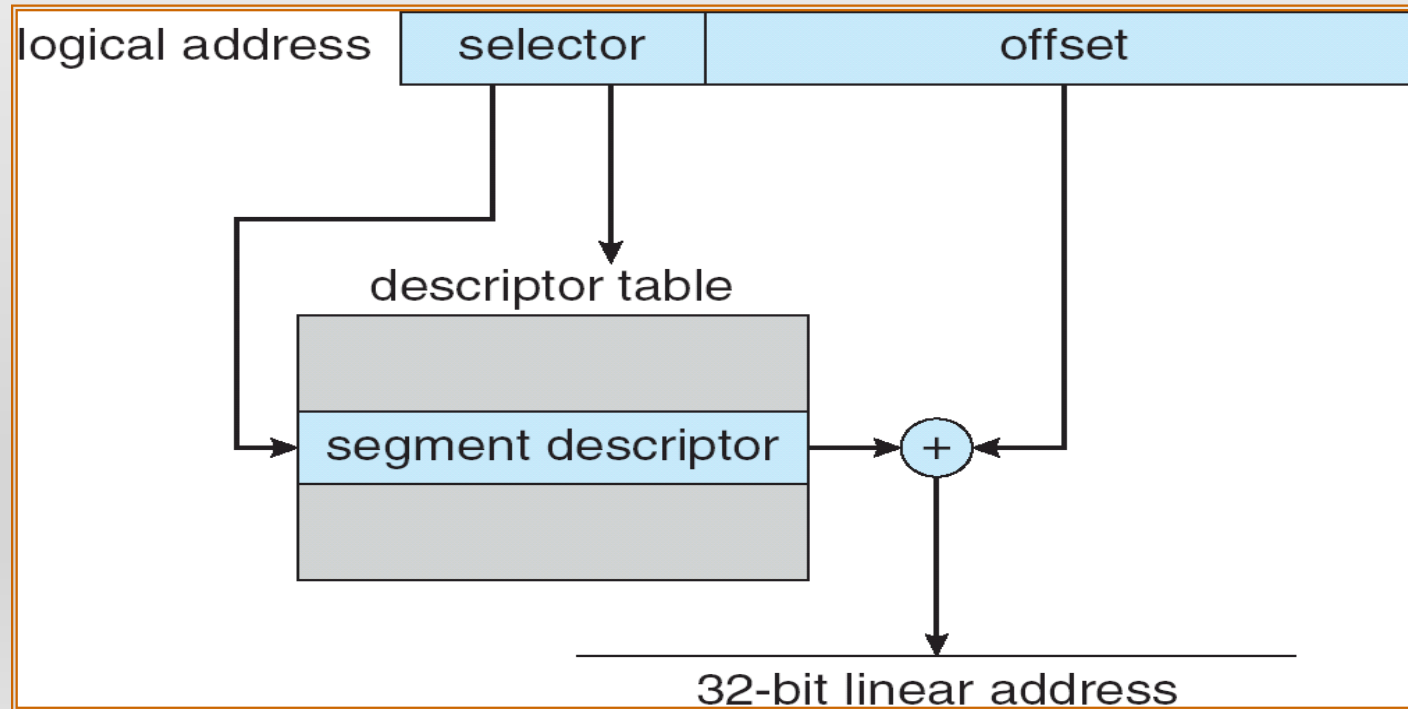
Example: The Intel Pentium

- **Logical address space** is divided into 2 partitions:
 - 1st: $8K(2^{13})$ segments (private), local descriptor table (LDT)
 - 2nd: $8K(2^{13})$ segments (shared), global descriptor table (GDT)
- Logical address:
 - max # of segments per process = $2^{14} = 16K$
 - size of a segment $\leq 2^{32} = 4GB$



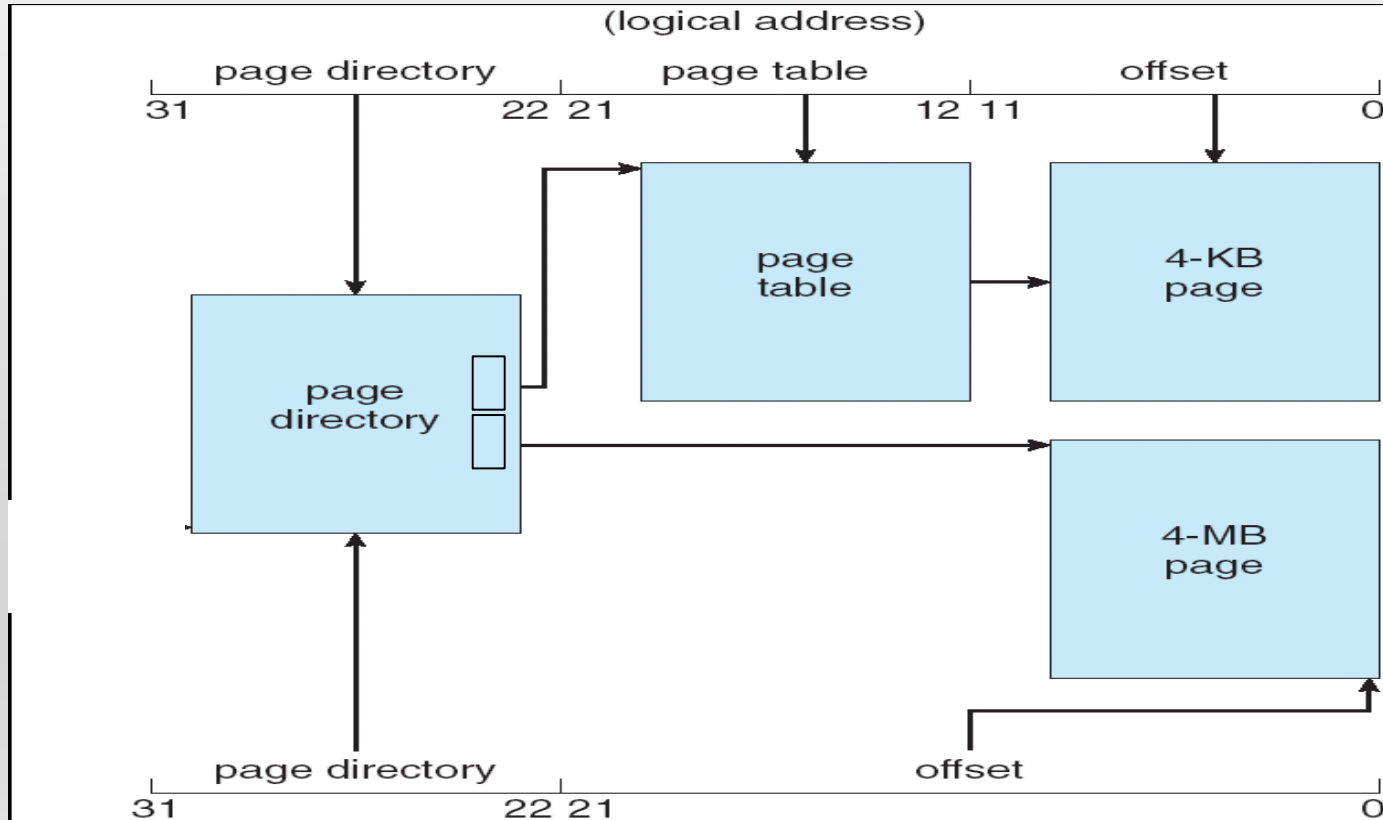
Intel Pentium Segmentation

- Segment descriptor
 - Segment base address and length
 - Access right and privileged level



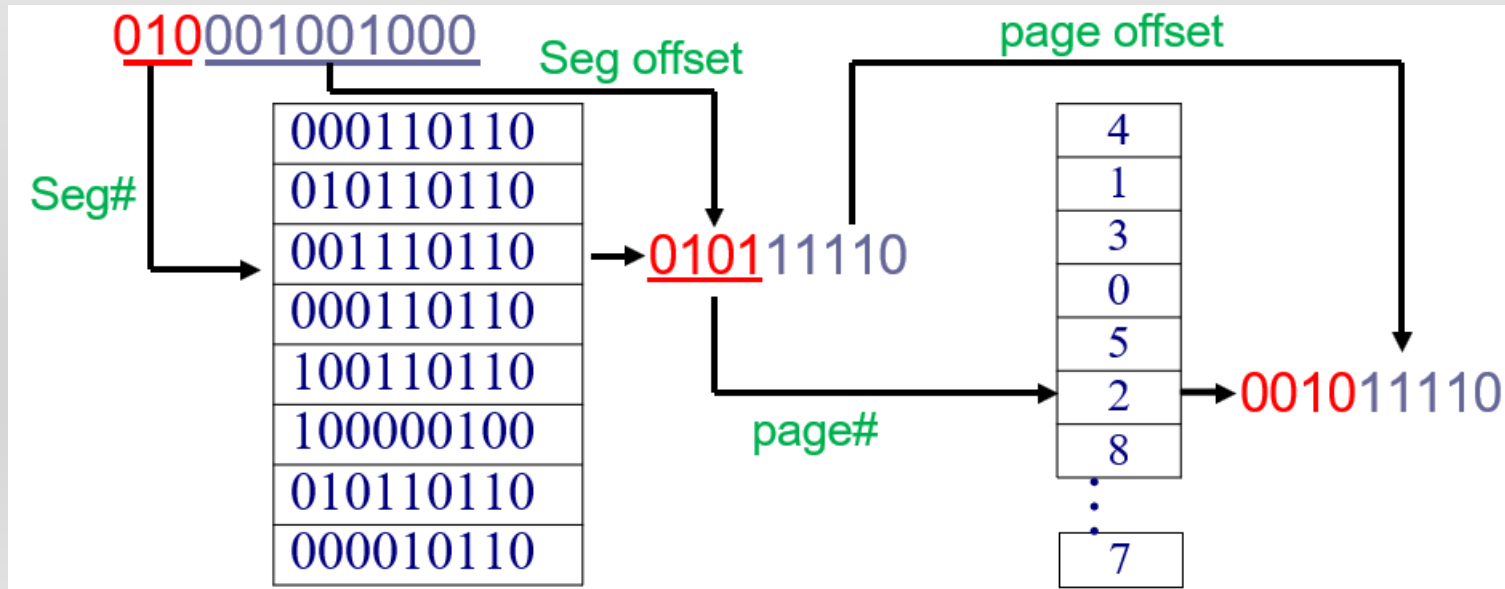
Intel Pentium Paging (Two-Level)

- Page size can be either 4KB or 4MB
 - Each page directory entry has a flag for indication



Example Question

- Let the **physical mem size is 512B**, the **page size is 32B** and the logical address of a program can have **8 segments**. Given a **12 bits hexadecimal logical address “448”**, translate the addr. With blow page and segment tables.
- linear addr:010111110, phy addr:001011110



Review Slides (5)

- Segmentation vs. Paging?

| | Paging | segmentation |
|---------------|----------------------------|------------------------------------|
| Length | Fixed | Varied |
| Fragmentation | Internal | External |
| Table entry | Page number → frame number | Seg ID → (base addr, limit length) |
| View | Physical memory | User program |

- Paged segmentation?

Reading Material & HW

- Chap 8
- Problem Set:
 - 8.1, 8.3, 8.4, 8.5, 8.12, 8.15, 8.16, 8.20, 8.23
- Interesting Reading:
 - M. Talluri, M. D. Hill, and Y. A. Khalidi. 1995. A new page table for 64-bit address spaces. SIGOPS Oper. Syst. Rev. 29, 5 (December 1995), 184-200.
 - http://pages.cs.wisc.edu/~markhill/papers/sosp95_pageta_bles.pdf

Q & A

Thank you for your attention