



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

*Operating System*

# Lecture 08: Virtual Memory Management

Shuo-Han Chen (陳碩漢),

[shch@nycu.edu.tw](mailto:shch@nycu.edu.tw)

Thur. 13:20 - 16:20 ED305

# Course Schedule

W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01 Due 10/5
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Typhoon – No class	V	
6	Oct. 9	Lec07: Memory Management	V	
7	Oct. 16	Lec08: Virtual Memory Management	V	HW02 Due 11/2
8	Oct. 23	Lec04: Process Scheduling	V	
9	Oct. 30	School Midterm Exam		
10	Nov. 6	Lec05: Process Synchronization	V	HW03
11	Nov. 13	Lec06: Deadlocks	V	
12	Nov. 20	School Event – No class		
13	Nov. 27	Lec09: File System Interface	V	HW04
14	Dec. 4	Lec10: File System Implementation	V	
15	Dec. 11	Lec11: Mass Storage System & Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

# Overview

- Background
- Demand Paging
- Process Creation
- Page Replacement
- Allocation of Frames
- Thrashing
- Operating System Examples


# Background

- Why we don't want to run a program that is entirely in memory....
  - Many code for handling unusual errors or conditions
  - Certain program routines or features are rarely used
  - The same library code used by many programs
  - Arrays, lists and tables allocated but not used
- -> We want better memory utilization

# Background

- **Virtual memory** - separation of user logical memory from physical memory
  - To run an extremely large process
    - **Logical address space** can be much larger than physical address space
  - To increase **CPU/resources utilization**
    - higher degree of multiprogramming degree
  - To **simplify programming** tasks
    - Free programmer from memory limitation
  - To run programs **faster**
    - less I/O would be needed to load or swap
- Virtual memory can be implemented via
  - **Demand paging**
  - Demand segmentation: more complicated due to variable sizes

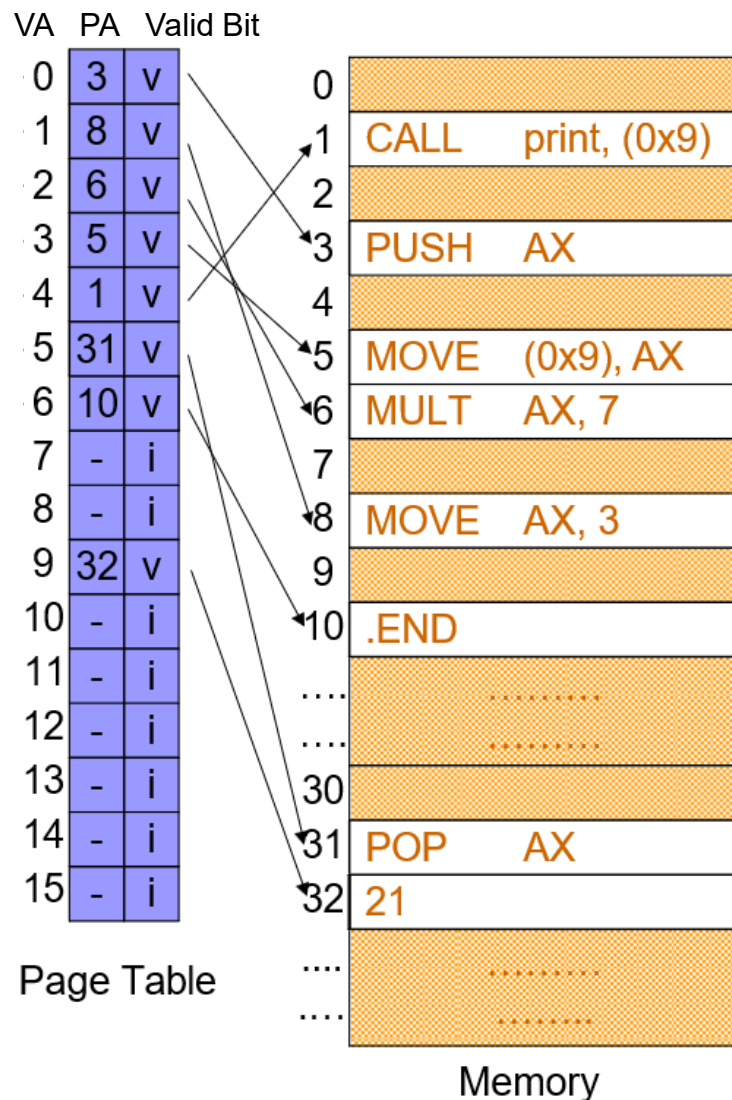
# Virtual Memory vs. Physical Memory



1. Initialize PCB, PC registers and Page Table.
2. Load Code into memory.
3. Running
4. Finish.

```
int data[8];
main( ) {
    data[3] = 3 * 7;
    print(data);
}
```

```
1 .START
2 PUSH    AX
3 MOVE    AX, 3
4 MULT    AX, 7
5 MOVE    (0x9), AX
6 CALL    print, (0x9)
7 POP     AX
8 .END
9 .SPACE (8)
```



# Demand Paging

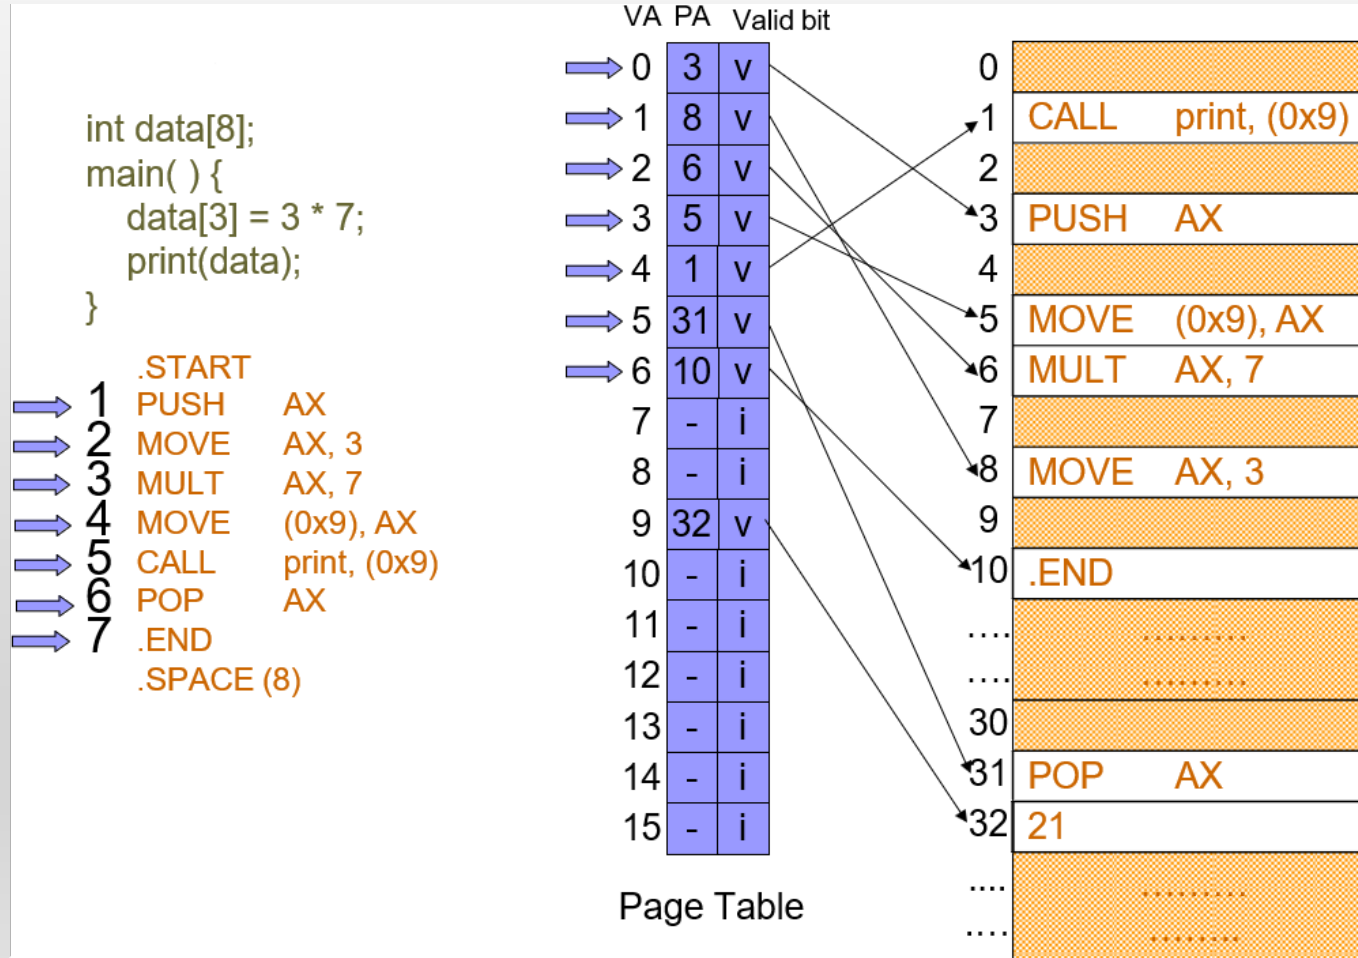
- A page rather than the whole process is brought into memory only when it is needed
  - Less I/O needed -> Faster response
  - Less memory needed -> More users
- Page is needed when there is a reference to the page
  - Invalid reference -> abort
  - Not-in-memory -> bring to memory via paging
- pure demand paging
  - Start a process with no page
  - Never bring a page into memory until it is required



# Demand Paging

- A **swapper** (midterm scheduler) manipulates the entire process, whereas a **pager** is concerned with the individual pages of a process
- Hardware support
  - **Page Table**: a **valid-invalid bit**
    - 1 -> page in memory
    - 0 -> page not in memory
    - Initially, all such bits are set to 0
  - Secondary memory (swap space, **backing store**): Usually, a high-speed disk (swap device) is use

# Demand Paging



0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical  
memory

	frame	valid-invalid bit
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	


physical memory



# Page Fault

- First reference to a page will trap to OS
- -> **page-fault trap**
- OS looks at the internal table (in PCB) to decide
  - Invalid reference -> abort
  - Just not in memory -> continue
- Get an empty frame
- Swap the page from disk (swap space) into the frame
- Reset page table, valid-invalid bit = 1
- **Restart instruction**

# Page Fault Handling Steps



# Page Replacement

- If there is no free frame when a page fault occurs
  - Swap a frame to backing store
  - Swap a page from backing store into the frame
  - Different page replacement algorithms pick different frames for replacement

# Demand Paging Performance

- Effective Access Time (EAT):  $(1 - p) \times ma + p \times pft$ 
  - P: **page fault rate**, ma : mem. access time, pft : page fault time
- Example:  $ma = 200ns$ ,  $pft = 8ms$ 
  - $EAT = (1 - p) \times 200ns + p \times 8ms$ 
    - $= 200ns + 7,999,800ns \times p$
- Access time is proportional to the page fault rate
  - If one access out of 1,000 causes a page fault, then
  - $EAT = 8.2$  microseconds. -> **slowdown by a factor of 40!**
  - For degradation less than 10%:
    - $220 > 200 + 7,999,800 \times p$ ,
    - $p < 0.0000025$  -> one access out of 399,990 to page fault

# Demand Paging Performance (Con't)

- Programs tend to have **locality** of reference
- Locality means program often accesses memory addresses that are close together
  - **A single page fault can bring in 4KB memory content**
  - Greatly reduce the occurrence of page fault
- major components of page fault time (about 8 ms)
  1. serve the page-fault interrupt
  2. **read in the page from disk (most expensive)**
  3. restart the process
  - The 1st and 3rd can be reduced to several hundred instructions
  - The page switch time is close to 8ms



# Process & Virtual Memory

- **Demand Paging**: only bring in the page containing the first instruction
- **Copy-on-Write**: the parent and the child process share the same frames initially, and frame-copy when a page is written
- **Memory-Mapped File**: map a file into the virtual address space to bypass file system calls (e.g., `read()`, `write()`)

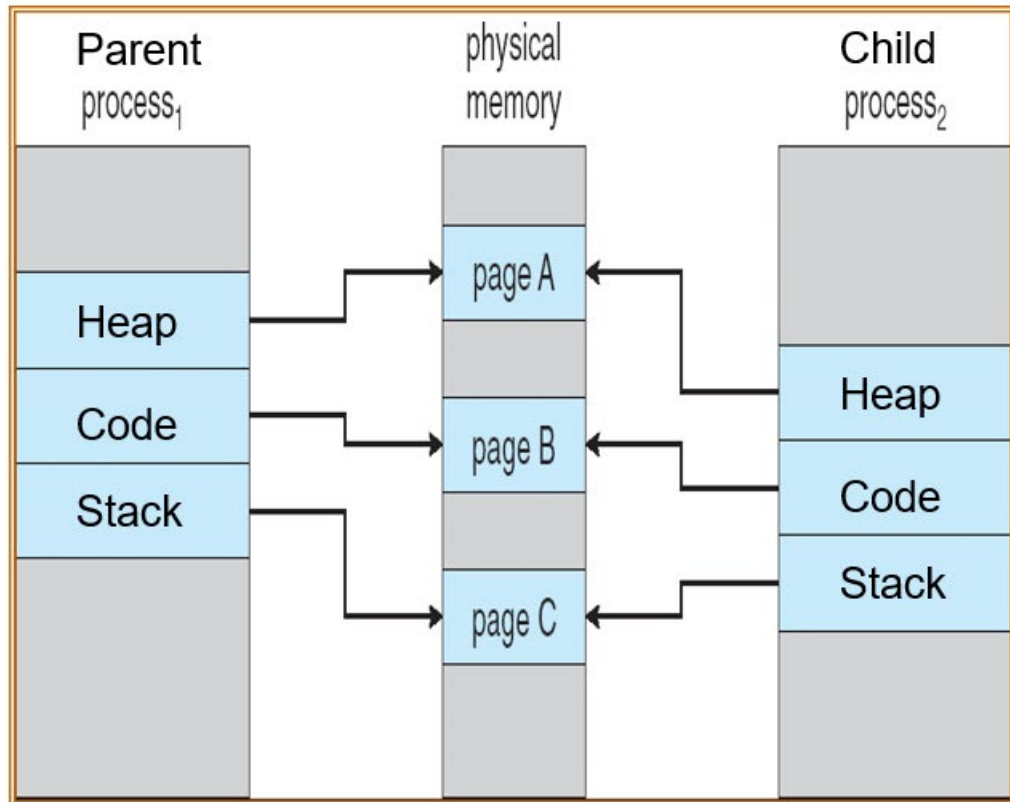
# Copy-on-Write

- Allow both the **parent** and the **child** process to **share the same frames** in memory
- If either process modifies a frame, only then a **frame is copied**
- COW allows efficient process creation (e.g., `fork()` )
- Free frames are allocated from a pool of **zeroed-out frames** (security reason)
  - The content of a frame is erased to 0

# When a child process is forked

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );

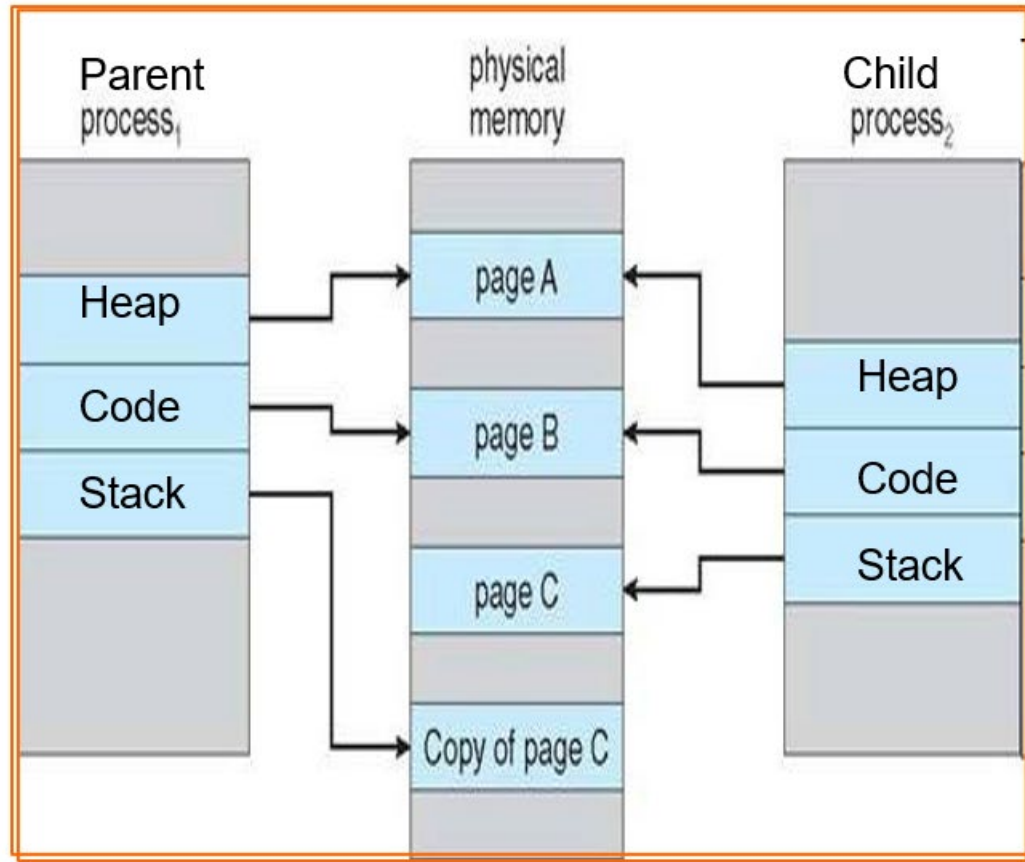
    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```



# After a page is modified

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );


    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```



# Memory-Mapped Files

- Approach:
  - MMF allows file I/O to be treated as routine memory access by mapping a disk block to a memory frame
  - A file is initially read using demand paging. Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- Benefit:
  - Faster file access by using memory access rather than `read()` and `write()` system calls
  - Allows several processes to map the SAME file allowing the pages in memory to be SHARED
- Concerns: Security, data lost, more programming efforts

# Memory-Mapped File Example



```

int buf;
int fd = open( filename, O_RDWR );
lseek( fd, 1024, SEEK_SET );
➔ read( fd, &buf, sizeof(int) );
➔ buf ++;
➔ lseek( fd, 1024, SEEK_SET );
➔ write( fd, &buf, sizeof(int) );
➔ close(fd);

```


**Memory  
mapped  
portion**

```


➔ int fd = open( filename, O_RDWR );
➔ int* area = mmap( 0, BUFSIZE,
    PROT_READ | PROT_WRITE,
    MAP_SHARED, fd, 1024 );
➔ area[0]++;
➔ close(fd);
➔ munmap( area, BUFSIZE );

```

## User space VM



## kernel file cache



# Review Slides ( I )

- Virtual memory? Physical Memory?
- Demand paging?
- Page table support for demand paging?
- OS handling steps for page fault?
- Page replacement?
- Copy-on-write? Usage?
- Memory-mapped file? Usage?




# Page Replacement Concept

- When a page fault occurs with no free frame
  - **swap out a process**, freeing all its frames, or
  - **page replacement**: find one not currently used and free it
    - Use **dirty bit** to reduce overhead of page transfers - **only modified pages are written to disk**
- Solve two major problems for demand paging
  - **frame-allocation algorithm**:
    - Determine **how many frames** to be allocated to a process
  - **page-replacement algorithm**:
    - select **which frame** to be replaced

# Page Replacement (Page Fault) Steps

- Find the location of the desired page on disk
- Find a free frame
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a victim frame
- Read the desired page into the (newly) free frame. Update the page & frame tables
- Restart the process instruction

# Page Replacement (Page Fault) Example



# Page Replacement Algorithms


- Goal: **lowest page-fault rate**
- Evaluation: running against a string of memory references (**reference string**) and computing the number of page faults
- Reference string:  
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Replacement Algorithms

- FIFO algorithm
- Optimal algorithm
- LRU algorithm
- Counting algorithm
  - LFU
  - MFU


# First-In-First-Out (FIFO) Algorithm

- The oldest page in a FIFO queue is replaced
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (available memory frames = 3)
- -> 9 page faults




# FIFO Illustrating Belady's Anomaly

- Does more allocated frames guarantee less page fault?
  - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
  - 4 frames (available memory frames = 4)
    - -> 10 page faults!
- Belady's anomaly
  - Greater allocated frames -> more page fault




# FIFO Illustrating Belady's Anomaly





# Optimal (Belady) Algorithm

- Replace the page that will not be used for the **longest period of time**
- -> need future knowledge
- 4 frames: **1, 2, 3, 4**, 1, 2, **5**, 1, 2, 3, **4**, 5 -> **6 page faults!**
- In practice, we don't have future knowledge
  - Only used for **reference & comparison**




# LRU Algorithm (Least Recently Used)

- An approximation of optimal algorithm:
  - looking backward, rather than forward
- It replaces the page that has not been used for the longest period of time
- It is often used, and is considered as quite good

# LRU Algorithm Implementations

- Counter implementation
  - page referenced: **time stamp** is copied into the counter
  - replacement: remove the one with oldest counter
    - **linear search is required...**
- Stack implementation
  - page referenced: move to top of the **double-linked list**
  - replacement: remove the page at the bottom
  - 4 frames: **1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5** -> **8 page faults!**



# Stack Algorithm

- A **property** of algorithms
- **Stack algorithm**: the set of pages in memory for  **$n$  frames** is always a **subset** of the set of pages that would be in memory with  **$n + 1$  frames**
- **Stack algorithms do not suffer from Belady's anomaly**
- Both **optimal** algorithm and **LRU** algorithm are stack algorithm

# LRU approximation algorithms



- **Few systems** provide sufficient hardware support for the LRU page-replacement
- additional-reference-bits algorithm
- second-chance algorithm
- enhanced second-chance algorithm

# Counting Algorithms

- **LFU Algorithm** (least frequently used)
  - keep a counter for each page
  - Idea: **An actively used page should have a large reference count**
- **MFU Algorithm** (most frequently used)
  - Idea: The page with the smallest count was probably just brought in and has yet to be used
- Both counting algorithm are not common
  - **implementation is expensive**
  - **do not approximate OPT algorithm very well**

# Introduction

- Each process needs minimum number of frames
- E.g.: IBM 370 - 6 pages to handle Storage to Storage MOVE instruction:
  - Both operands are in main storage, the first operand is  $B1(\text{Reg.ID})+D1$ , the second operand is  $B2(\text{Reg.ID})+D2$ ,  $L$  plus 1 is the length.
  - instruction is 6 bytes, may span 2 pages
  - Moving content could across 2 pages



# Frame Allocation

- Fixed allocation
  - **Equal allocation** - 100 frames, 5 processes -> 20 frames/process
  - **Proportional allocation** - Allocate according to the size of the process
- Priority allocation
  - using **proportional allocation based on priority**, instead of size
  - if process P generates a page fault
    - select for replacement one of its frames
    - select for replacement **from a process with lower priority**



# Frame Allocation

- Local allocation: each process select from its own set of allocated frames
- Global allocation: process selects a replacement frame from the set of all frames
  - one process can take away a frame of another process
  - e.g., allow a high-priority process to take frames from a low-priority process
  - good system performance and thus is common used
  - A minimum number of frames must be maintained for each process to prevent trashing

# Review Slides ( II )

- Page replacement steps?
- Page replacement algorithm goal?
- Dirty bit usage?
- Belady's anomaly?
- FIFO? Optimal? LRU?
- Fixed vs. priority frame allocation?
- Global vs. local frame allocation?

# Definition of Thrashing

- If a process does not have “enough” frames
  - the process does not have # frames it needs to support pages in active use
  - -> Very high paging activity
- A process is thrashing if it is spending more time paging than executing



# Thrashing

- Performance problem caused by thrashing (Assume global replacement is used)
  - processes **queued for I/O** to swap (page fault)
  - -> **low CPU utilization**
  - -> **OS increases the degree of multiprogramming**
  - -> new processes take frames from old processes
  - -> more page faults and thus more I/O
  - -> CPU utilization drops even further
- To prevent thrashing, must provide enough frames for each process:
  - **Working-set model, Page-fault frequency**

# Working-Set Model


- **Locality**: a set of pages that are **actively used together**
- Locality model: as a process executes, it **moves from locality to locality**
  - program structure (subroutine, loop, stack)
  - data structure (array, table)
- **Working-set model** (based on locality model)
  - working-set **window**: a parameter  $\Delta$  (delta)
  - working set: set of pages in most recent page  $\Delta$  references (**an approximation locality**)

# Working-Set Example


- If  $\Delta = 10$ :

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

# Working-Set Model


- Prevent thrashing using the working-set size
  - WSSi : working-set size for process i
  - $D = \sum WSS_i$  (total demand frames)
  - If  $D > m$  (available frames)  $\rightarrow$  thrashing
  - The OS monitors the WSSi of each process and allocates to the process enough frames
    - if  $D \ll m$ , increase degree of MP
    - if  $D > m$ , suspend a process
- Good :
  1. prevent thrashing while keeping the degree of multiprogramming as high as possible
  2. optimize CPU utilization
- Bad : too expensive for tracking

# Page Fault Frequency Scheme


- Page fault frequency directly measures and controls the page-fault rate to prevent thrashing
  - Establish upper and lower bounds on the desired page-fault rate of a process
  - If page fault rate exceeds the upper limit
    - allocate another frame to the process
  - If page fault rate falls below the lower limit
    - remove a frame from the process



# Page Fault Frequency Scheme



# Working Sets and Page Fault Rates



- Memory has locality property
- When the process moves to a new WS, the PF rate rises toward a peak

# Review Slides ( III )

- Thrashing definition?
- Process locality?
- When will thrashing happen? Solution?

# Reading Material & HW

- Chap 9
- Problems
  - 9.2, 9.4, 9.6, 9.8, 9.9, 9.12, 9.14, 9.17, 9.19, 9.21

# Backup

# Windows NT

- Uses demand paging with **clustering**. Clustering **brings in pages surrounding the faulting page**
- Processes are assigned **working-set minimums** and **working-set maximums**
- WS minimum: the minimum # of pages the process is guaranteed to be in memory
- A process can have pages up to its WS maximum
- When the amount of free memory in the system falls below a threshold, automatic working set trimming is performed
- **Working set trimming** removes pages from processes that have pages in excess of their WS minimum

Q & A

*Thank you for your attention*