



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

Operating System

Lecture 09: File System Interface

Shuo-Han Chen 陳碩漢

shch@nycu.edu.tw

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

Course Schedule


W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01 Due 10/5
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Typhoon – No class	V	
6	Oct. 9	Lec07: Memory Management	V	
7	Oct. 16	Lec08: Virtual Memory Management	V	HW02 Due 11/2
8	Oct. 23	Lec04: Multithreaded Programming	V	
9	Oct. 30	Midterm Exam		
10	Nov. 6	Lec05: Process Scheduling	V	Let's take a breath
11	Nov. 13	Lec06: Process Synchronization & Deadlocks	X	HW03
12	Nov. 20	School Event – No class	V	
13	Nov. 27	Lec09: File System Interface	V	
14	Dec. 4	Lec10: File System Implementation	V	HW04
15	Dec. 11	Lec11: Mass Storage System & Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

Overview

- File Concept
- Access Methods
- Directory Structure
- File System Mounting
- File Sharing
- Protection

File Concept

- **File**: a **logical storage unit** created by OS
 - v.s. **physical storage unit** in disk (sector, track)
- **File attributes**
 - Identifier: non-human-readable name
 - Name
 - Type
 - Location
 - Size
 - Protection
 - Last-access time, Last-updated time

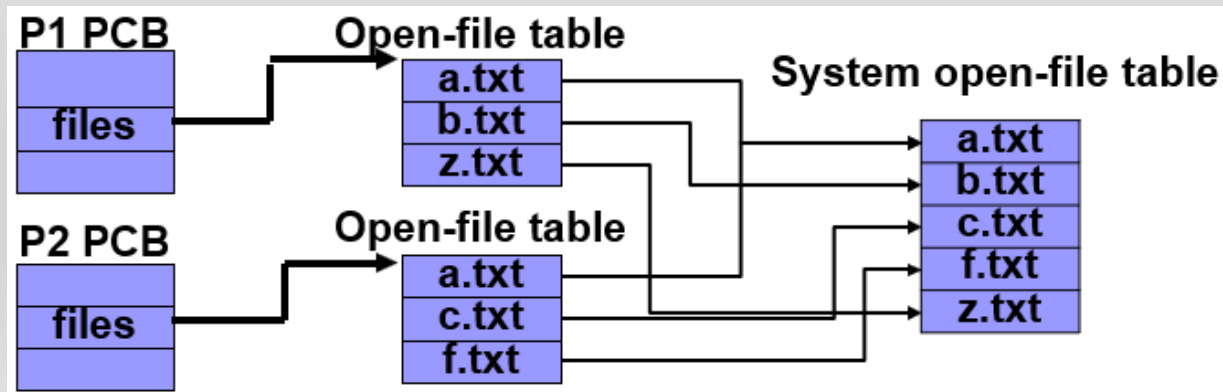


File Operations

- File operations include
 - Creating a file
 - Writing a file
 - Reading a file
 - Repositioning within a file (i.e. file seek)
 - Deleting a file
 - Truncating a file
- Process: open-file table
- OS: system-wide table

Open-File Tables

- **Per-process table**
 - Tracking all files opened by this process
 - Current **file pointer** for each opened file
 - **Access rights** and **accounting** information
- **System-wide table**
 - Each entry in the per-process table points to this table
 - **Process-independent information** such as **disk location**, **access dates**, **file size**
 - **Open count**



Open File Attributes


- Open-file attributes (**metadata**)
 - File pointer (per-process)
 - File open count (system table)
 - Disk location (system table)
 - Access rights (per-process)
- File types
 - .exe, .com, .obj, .cc, .mov, etc
 - **Hint** for OS to operate file in a **reasonable** way

file type	usual extension	function
executable	exe, com, bin or none	read to run machine-language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rrf, doc	various word-processor formats
library	lib, a, so, dll, mpeg, mov, rm	libraries of routines for programmers
print or view	arc, zip, tar	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes compressed, for archiving or storage
multimedia	mpeg, mov, rm	binary file containing audio or A/V information

Access Method

Access Methods

- Sequential access
 - Read/write next (block)
 - Reset: repositioning the file pointer to the beginning
 - Skip/rewind n records




Access Methods

- Direct (relative) access
 - Access an element at an **arbitrary position** in a sequence
 - File operations include the **block #** as parameter
 - Often use **random access** to refer the **access pattern** from direct access

sequential access	implementation for direct access
<i>reset</i>	<i>cp = 0;</i>
<i>read next</i>	<i>read cp;</i> <i>cp = cp+1;</i>
<i>write next</i>	<i>write cp;</i> <i>cp = cp+1;</i>

Index Access Methods

- Index: contains pointers to **blocks of a file**
- To find a record in a file:
 - search the index file -> find the pointer
 - use the pointer to directly access the record
- With a large file -> index could become too large



Review Slides (I)


- File vs. Sector, Track
- Open-file (in-memory) attributes
 - Per-process, system-wide?
- File-access methods?
 - Sequential access
 - Direct access
 - Index access

Directory Structure

Partition, Volume & Directory


- A **partition** (**formatted** or **raw**)
 - raw partition (no file system): UNIX **swap space**, **database**
 - Formatted partition with file system is called **volume**
 - a partition can be **a portion of a disk** or **group of multiple disks (distributed file system)**
 - Some storage devices (e.g.: floppy disk) does not and cannot have partition
- **Directories** are used by file system to store the information about the files in the partition

File-System Organization



Directory vs. File


- **Directory**: A collection of nodes containing information about all files
 - Both the directory structure and the files reside on disk



Directory Operations

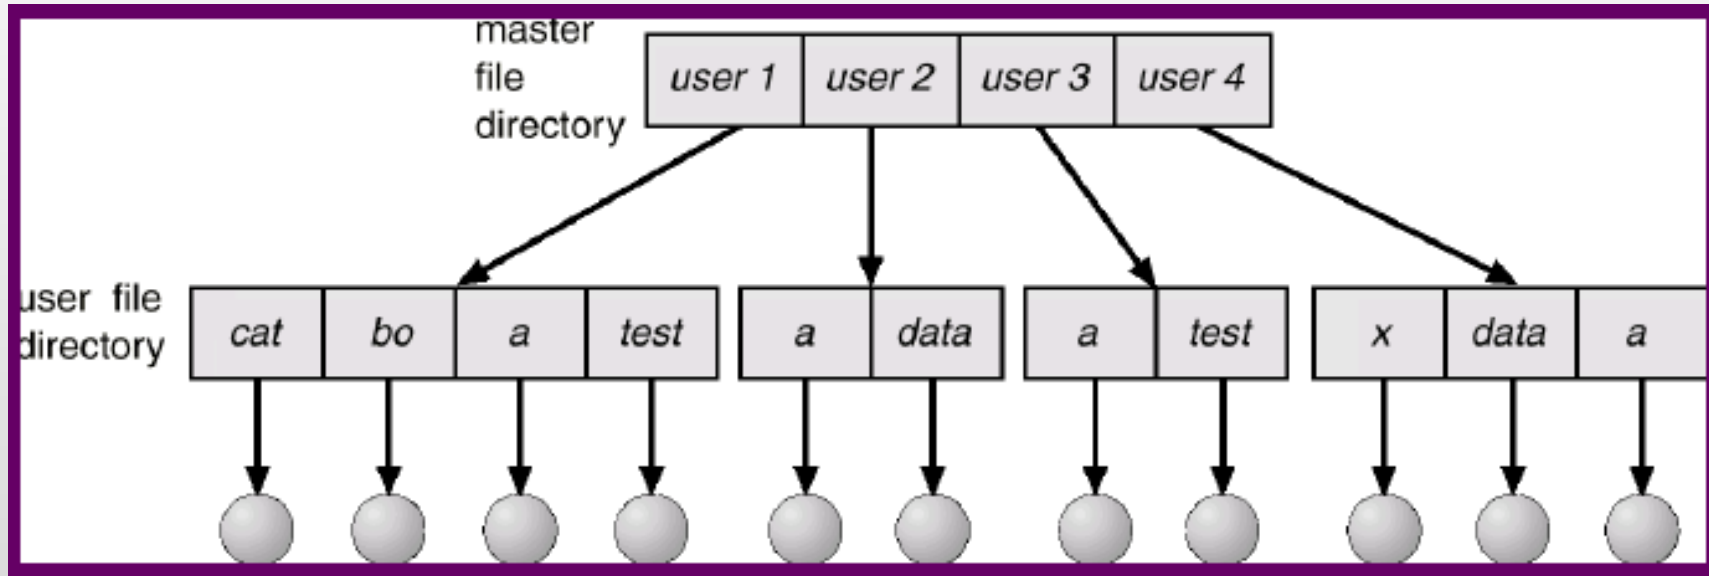
- Search for a file
- Create a file
- Delete a file
- List a directory
- Rename a file
- Traverse the file system

Single-Level Directory



- All files in one directory
 - Filename has to be unique
 - Poor efficiency in locating a file as number of files increases

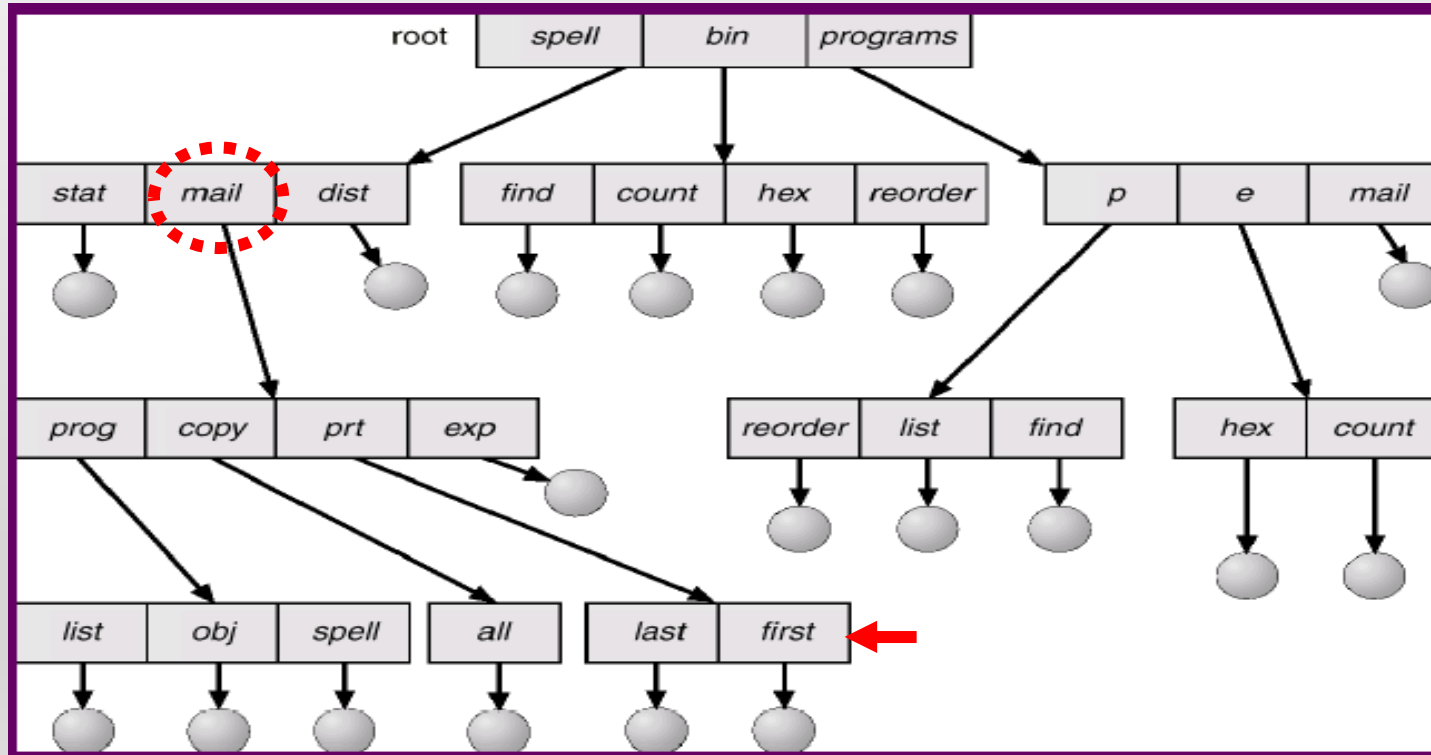
Two-Level Directory



- a separate dir for each user
- path = user name + file name
- single-level dir problems still exists per user

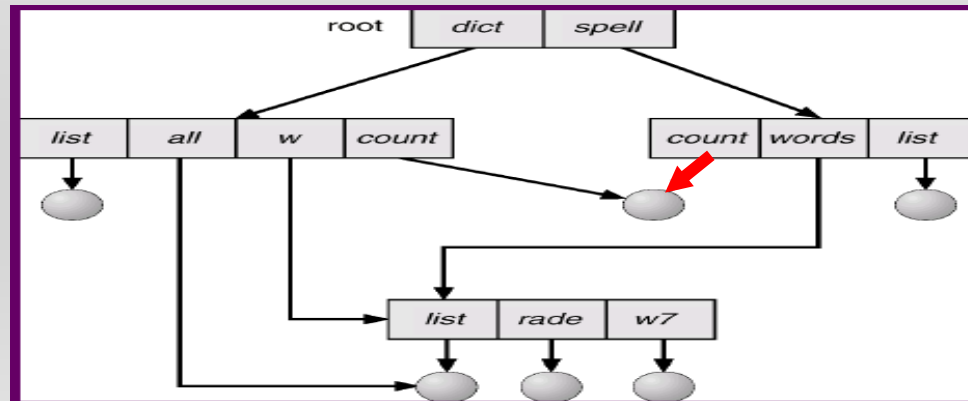
Tree-Structured Directory

- **Absolute path**: starting from the root
- **Relative path**: starting from a directory

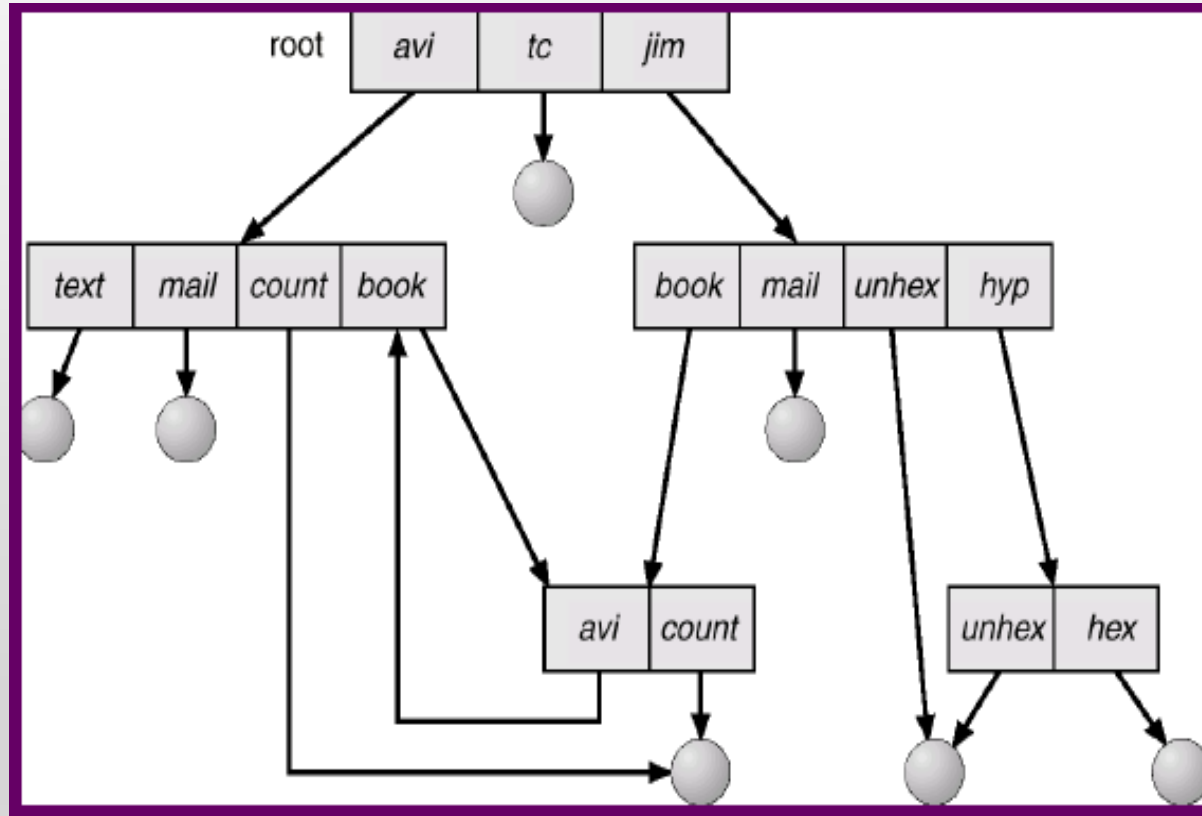


Acyclic-Graph Directory

- Use links to share files or directories
 - UNIX-like: **symbolic link** (ln -s /spell/count /dict/count)
- A file can have **multiple absolute paths**
- When does a file actually get deleted?
 - deleting the link but not the file
 - deleting the file but leaves the link -> **dangling pointer**
- deleting the file when **reference counters** is 0




General-Graph Directory



General-Graph Directory

- May contain cycles
 - Reference count does not work any more
 - E.g. self-referencing file
- How can we deal with cycles?
 - Garbage collection
 - First pass traverses the entire graph and marks accessible files or directories
 - Second pass collect and free everything that is un-marked
 - Poor performance on millions of files ...
 - Use cycle-detection algorithm when a link is created



Review Slides (II)

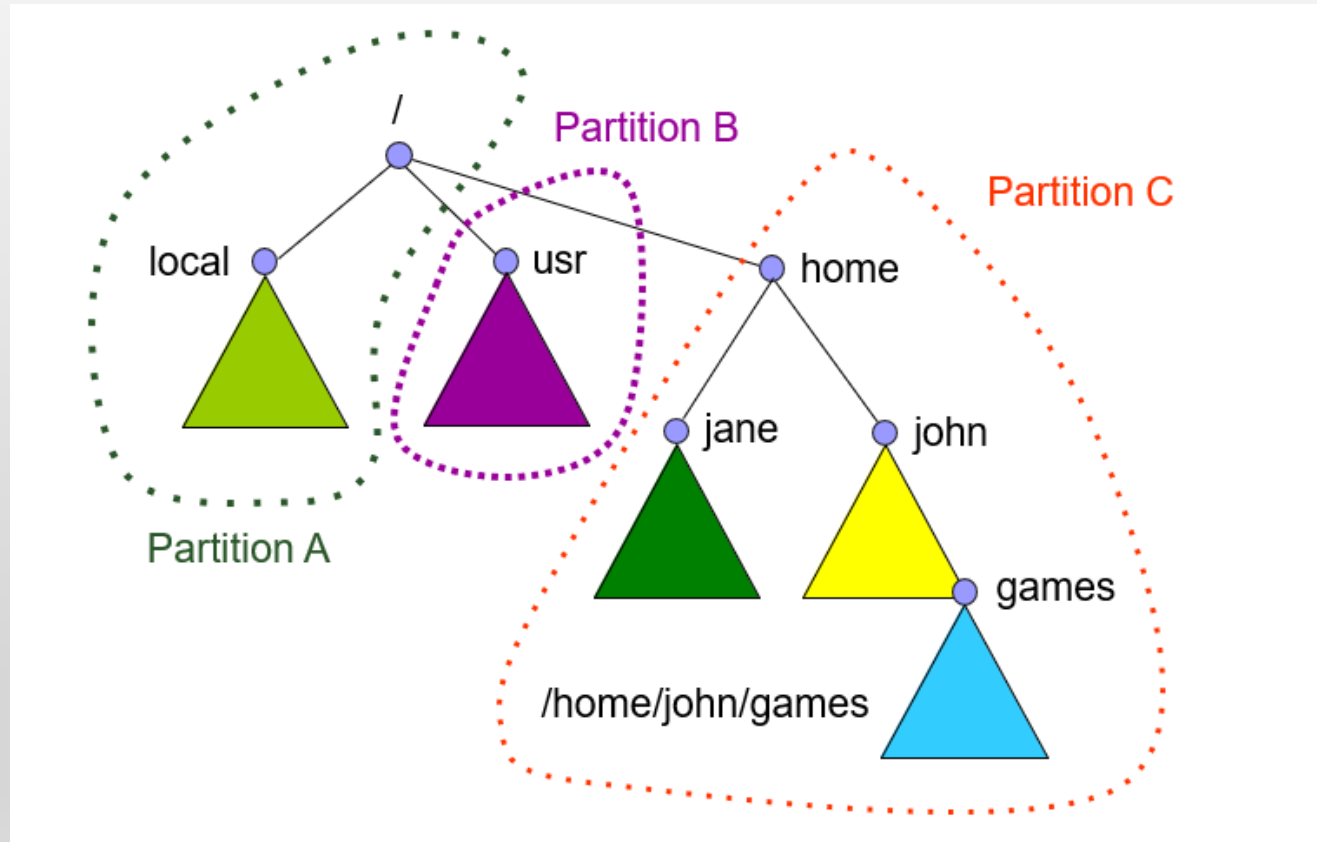
- Directory structure: pros & cons
 - One-level directory
 - Two-level directory
 - Tree-structured directory
 - Acyclic-graph directory
 - General-graph directory

File-System Mounting & File Sharing

File System Mounting


- A file system must be **mounted before** it can be accessed
- **Mount point:** **the root path** that a FS will be mounted to
- **Mount timing:**
 - **boot time**
 - **automatically at run-time**
 - **manually at run-time**

File System Mounting Example



File Sharing on Multiple Users


- Each user: (userID, groupID)
 - ID is associated with every ops/process/thread the user issues
- Each file has 3 sets of attributes
 - owner, group, others
- Owner attributes describe the **privileges** for the owner of the file
 - same for group/others attributes
 - group/others attributes are set by owner or root



Access-Control List

- We can create an **access-control list** (ACL) for each user
 - check requested file access against ACL
 - problem: unlimited # of users
- 3 classes of users -> 3 ACL (**RWX**) for **each file**
 - owner (e.g. 7 = RWX = 111)
 - group (e.g. 6 = RWX = 110)
 - public (others) (e.g. 4 = RWX = 100)

chmod 664 intro.ps



-rw-rw-r--	1	pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	5	pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	2	pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	2	pbg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	1	pbg	staff	9423	Feb 24 2003	program.c
-rwxr-xr-x	1	pbg	staff	20471	Feb 24 2003	program
drwx--x--x	4	pbg	faculty	512	Jul 31 10:31	lib/
drwx-----	3	pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	3	pbg	staff	512	Jul 8 09:35	test/

File Protection

- File owner/creator should be able to control
 - what can be done
 - by whom
 - -> Access control list (ACL)
- Files should be kept from
 - physical damage (reliability): i.e. RAID
 - improper access (protection): i.e. password

Review Slides (III)

- File system mounting point, timing?
- Access-control list? How does it function?

Reading Material & HW

- Chap 10
- Problems
 - 10.1: Consider a file system where a file can be deleted and its disk space reclaimed while links to that file still exist. What problems may occur if a new file is created in the same storage area or with the same absolute path name? How can these problems be avoided?
 - 10.4: Provide examples of applications that typically access files according to “sequential” and “random”.
 - 10.6: If the operating system knew that a certain application was going to access file data in a sequential manner, how could it exploit this information to improve performance?

Consistency Semantics

- When files are shared, ops from different users to the same file must be **synchronized**
- UNIX semantics
 - write is visible to all other users opening the same file
 - Open-file option: share the same file pointer
- Session semantics (AFS file system)
 - write is not visible to all other users
 - once a file is closed, changes are visible for sessions starting later -> current sessions do not see changes
- Immutable-Shared-Files semantics
 - once a file is declared shared, it cannot be modified

File Sharing on Remote File Systems


- Uses networking to allow file system access between systems
 - **Manually** via programs like **FTP**
 - **Semi automatically** via the **world wide web**
 - **Automatically, seamlessly** using **distributed file systems**

Client-server model

- Allows clients to mount remote file systems from servers
 - **Sever**: the machine that **owns the files** and **serves multiple clients**
 - **Client**: the machine that **accesses remote files**
 - Standard OS file calls are **translated into remote calls**
 - Client and user-on-client **identification is insecure or complicated**
- Example:
 - **NFS (network file sysytem)** for UNIX
 - **CIFS (common interface file system)** for Windows

Distributed Information Systems

- Distributed naming services
 - Provide **unified access** to the info for **remote computing**



Failure Modes

- Failures:
 - HW: disk, network cable, switch, server, etc.
 - SW: corruption or inconsistency of file, directory structure, etc.
- We need to recover:
 - Data: files, directory contents
 - Metadata: data and system management info.
- Stateful vs. Stateless communication protocol:
 - Stateless: treats each request as an independent transaction that is unrelated to any previous request (HTTP)
 - Stateful: info. maintained on both client and server is required

Q & A

Thank you for your attention