



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Institute of Artificial Intelligence Innovation

Department of Computer Science

*Operating System*

# Lecture 12: I/O Systems

Shuo-Han Chen 陳碩漢

[shch@nycu.edu.tw](mailto:shch@nycu.edu.tw)

Wed. 10:10 - 12:00 EC115 +

Fri. 11:10 – 12:00 Online

# Course Schedule

W	Date	Lecture	Online	Homework
1	Sept. 4	Lec00: Course Overview & Historical Prospective		
2	Sept. 11	Lec01: Introduction	V	
3	Sept. 18	Lec02: OS Structure	V	HW01 Due 10/5
4	Sept. 25	Lec03: Processes Concept	X	
5	Oct. 2	Typhoon – No class	V	
6	Oct. 9	Lec07: Memory Management	V	
7	Oct. 16	Lec08: Virtual Memory Management	V	HW02 Due 11/2
8	Oct. 23	Lec04: Multithreaded Programming	V	
9	Oct. 30	Midterm Exam		
10	Nov. 6	Lec05: Process Scheduling	V	Let's take a breath
11	Nov. 13	Lec06: Process Synchronization & Deadlocks	X	HW03
12	Nov. 20	School Event – No class	V	
13	Nov. 27	Lec09: File System Interface	V	
14	Dec. 4	Lec10: File System Implementation	V	HW04
15	Dec. 11	Lec11: Mass Storage System & Lec12: IO Systems	V	
16	Dec. 18	School Final Exam		

# Outline

- Overview
- I/O Hardware
- **I/O Methods**
- Kernel I/O Subsystem
- Performance
- Application Interface

# Overview

- The two main jobs of a computer
  - I/O and Computation
- I/O devices: tape, HD, mouse, joystick, network card, screen, flash disks, etc
- I/O subsystem: the methods to control all I/O devices
- Two conflicting trends
  - Standardization of HW/SW interfaces
  - Board variety of I/O devices

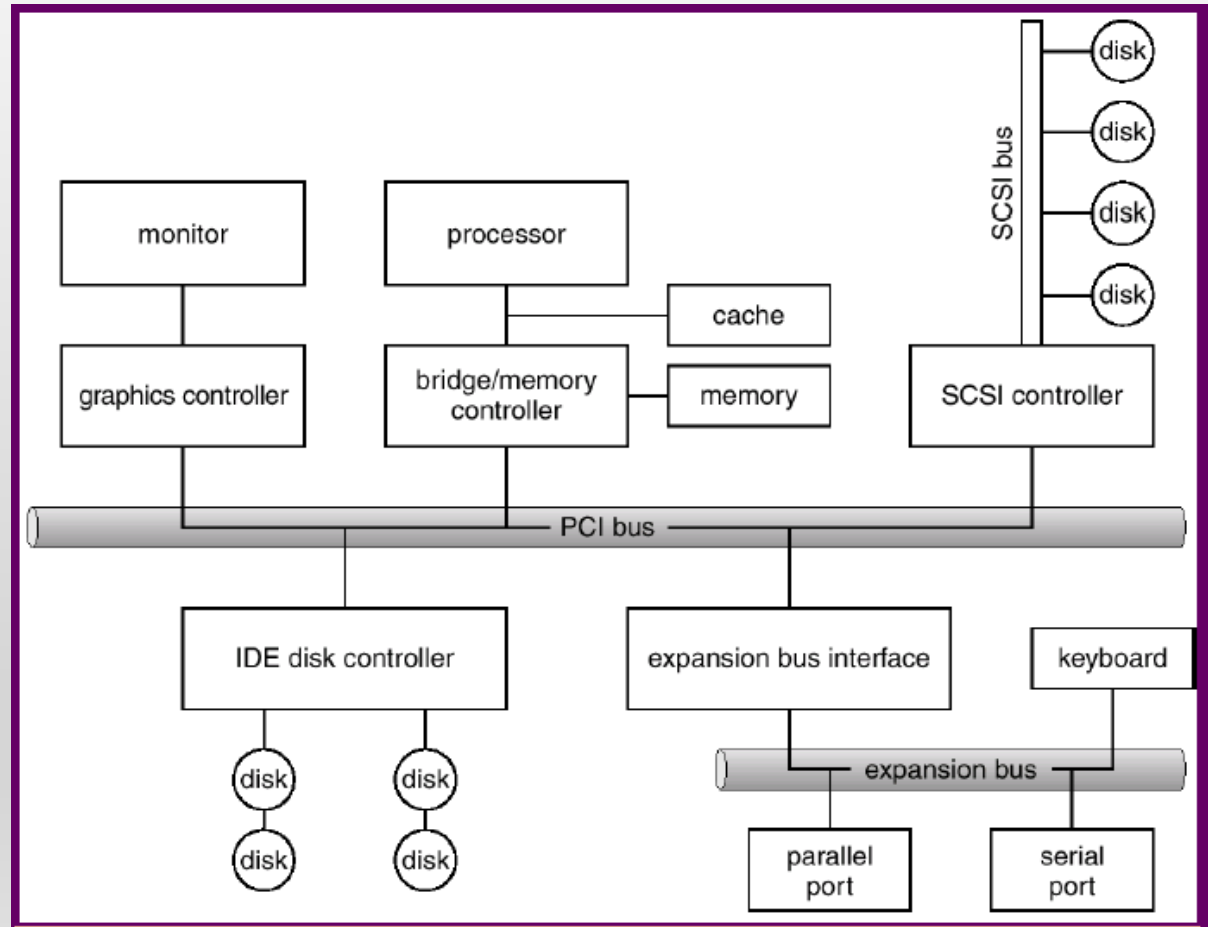
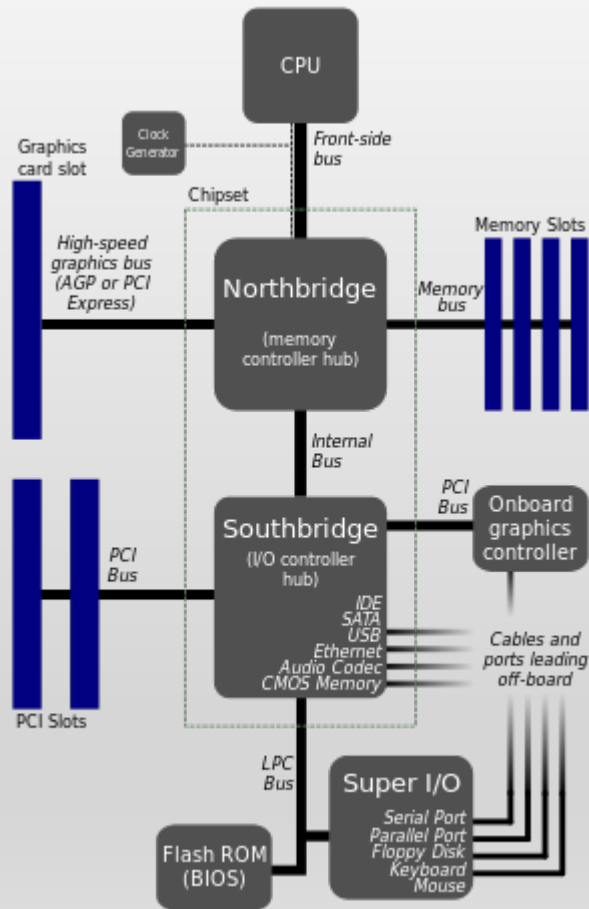
# Overview

- **Device drivers:** a uniform device-access **interface to the I/O subsystem**
  - Similar to system calls between apps and OS
- Device categories
  - Storage devices: disks, tapes
  - Transmission devices: network cards, modems
  - Human-interface devices: keyboard, screen, mouse
  - Specialized devices: joystick, touchpad

# I/O Hardware

- **Port:** A **connection point** between I/O devices and the host
  - E.g.: USB ports
- **Bus:** A set of **wires** and a well-defined **protocol** that specifies messages sent over the wires
  - E.g.: PCI bus
- **Controller:** A collection of electronics that can **operate** a port, a bus, or a device
  - A controller could have its own processor, memory, etc. (E.g.: SCSI controller)

# Typical PC Bus Structure



# Basic I/O Method (Port-mapped I/O)

- Each I/O port (device) is identified by a unique **port address**
- Each I/O port consists of **four registers** (1~4Bytes)
  - **Data-in register**: read by the host to get input
  - **Data-out register**: written by the host to send output
  - **Status register**: read by the host to check I/O status
  - **Control register**: written by the host to control the device
- Program interact with an I/O port through **special I/O instructions (different from mem. access)**
  - X86: IN, OUT



# Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

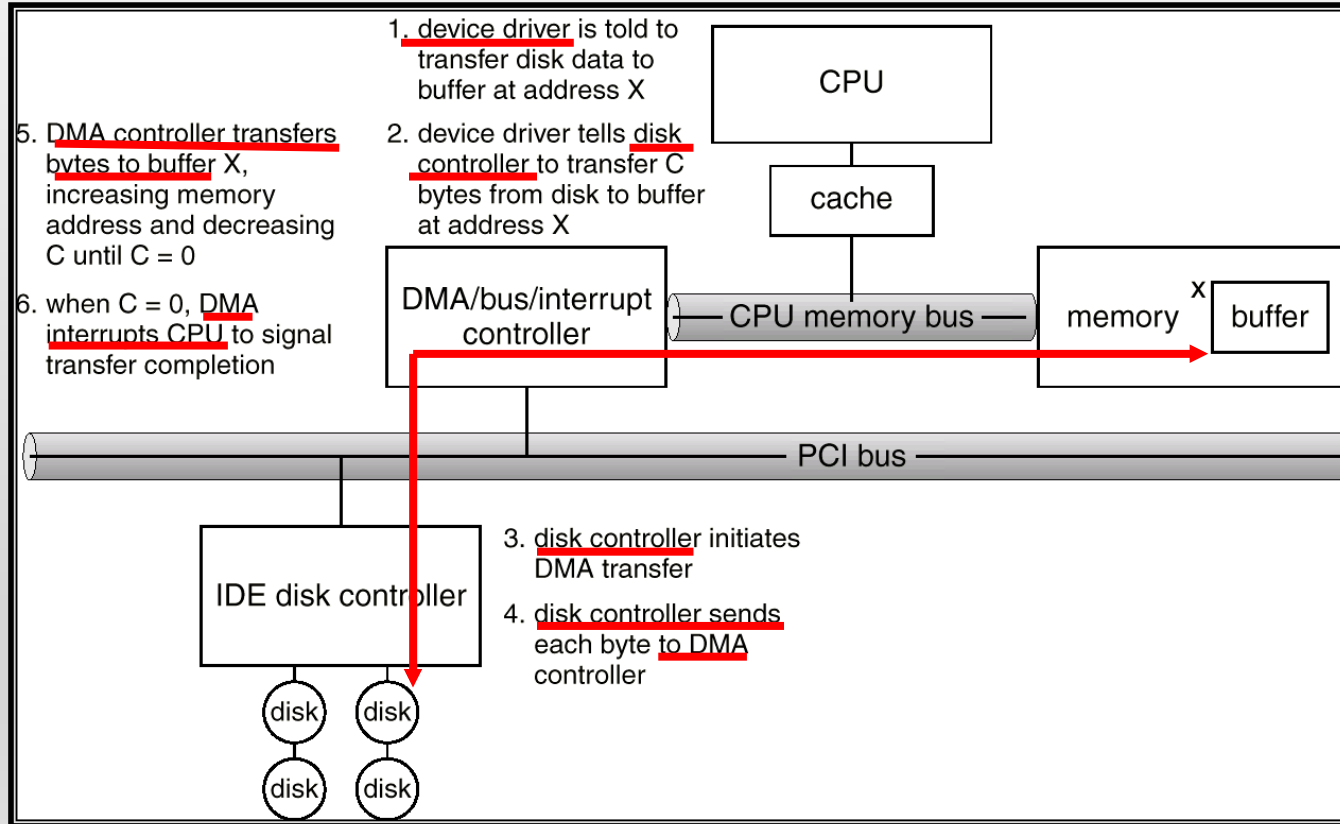
# I/O Methods Categorization

- Depending on how to address a device:
  - Port-mapped I/O
    - Use different address space from memory
    - Access by special I/O instruction (e.g. **IN**, **OUT**)
  - Memory-mapped I/O
    - Reserve specific memory space for device
    - Access by standard data-transfer instruction (e.g. **MOV**)
    - Good: More efficient for large memory I/O (e.g. graphic card)
    - Bad: Vulnerable to accidental modification, error

# I/O Methods Categorization

- Depending on how to interact with a device:
  - **Poll** (busy-waiting): processor periodically check status register of a device
  - **Interrupt**: device notify processor of its completion
- Depending on who to control the transfer:
  - **Programmed I/O**: transfer controlled by **CPU**
  - **Direct memory access** (DMA) I/O: controlled by **DMA controller** (a special purpose controller)
    - Design for **large data transfer**
    - Commonly used with **memory-mapped I/O** and **interrupt**

# Six-Step Process to Perform DMA (Direct Memory Access)



# Review Slides ( I )

- Definition of I/O port? Bus? Controller?
- I/O device and CPU communication?
  - Port-mapped vs. Memory-mapped
  - Poll vs. Interrupt
  - Programmed I/O vs. DMA
- Steps to handle an interrupt I/O and DMA request?

# Kernel I/O Subsystem

# I/O Subsystem

- **I/O Scheduling** - improve system performance by ordering the jobs in I/O queue
  - e.g. disk I/O order scheduling
- **Buffering** - store data in memory while transferring between I/O devices
  - **Speed mismatch** between devices
  - Devices with **different data-transfer sizes**
  - Support copy semantics

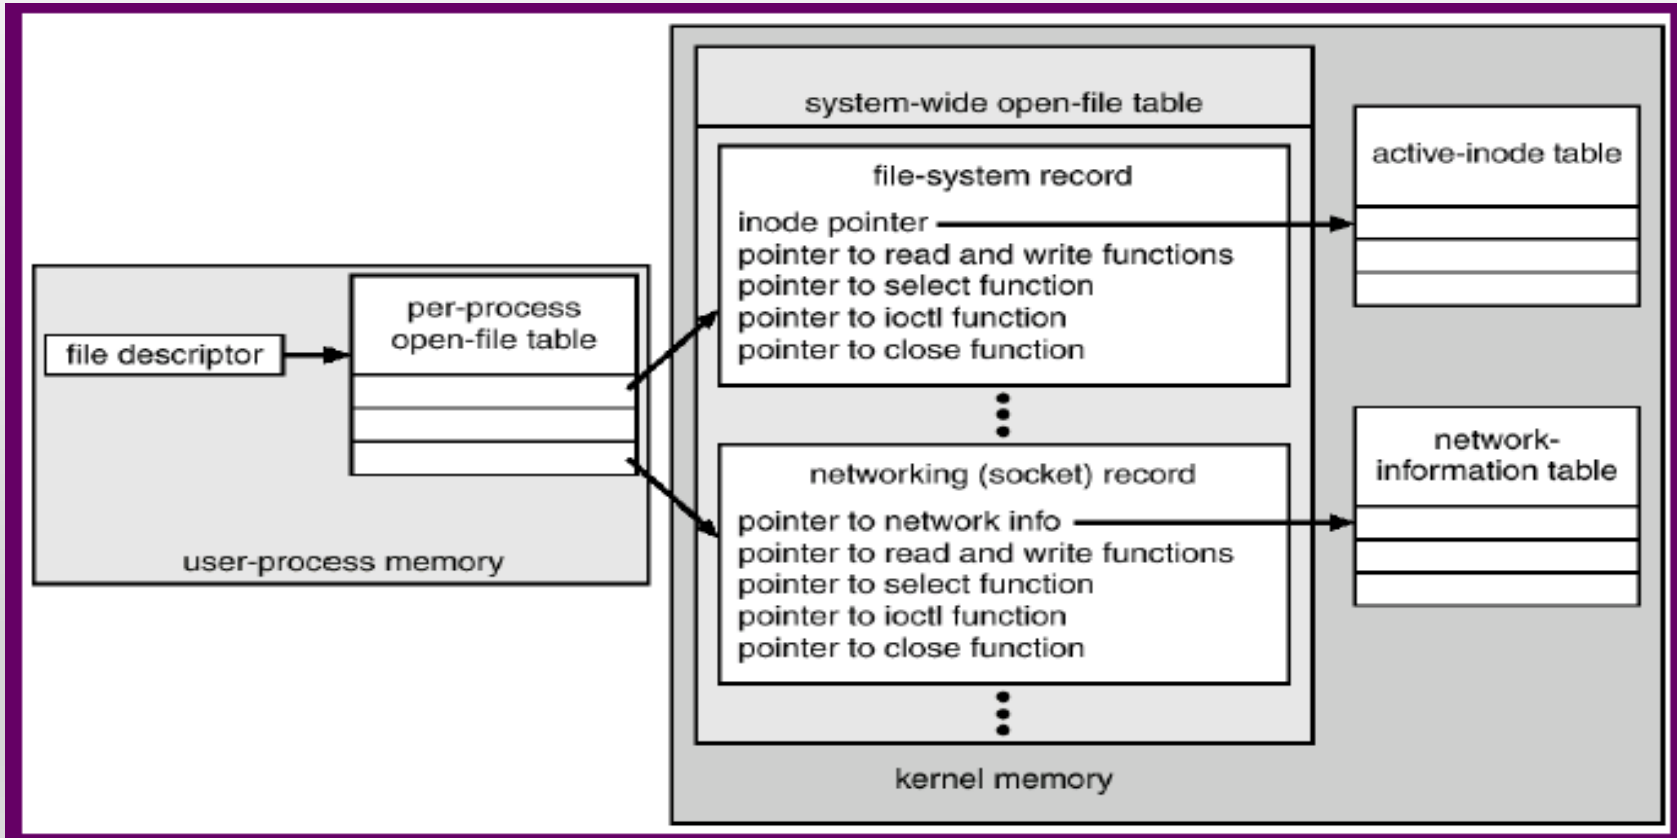
# I/O Subsystem

- **Caching** - fast memory that holds copies of data
  - Always just a copy
  - Key to **performance**
- **Spooling** - holds output for a **device**
  - e.g. printing (cannot accept **interleaved files**)
- **Error handling** - when I/O error happens
  - e.g. SCSI devices returns error information
- **I/O protection**
  - Privileged instructions

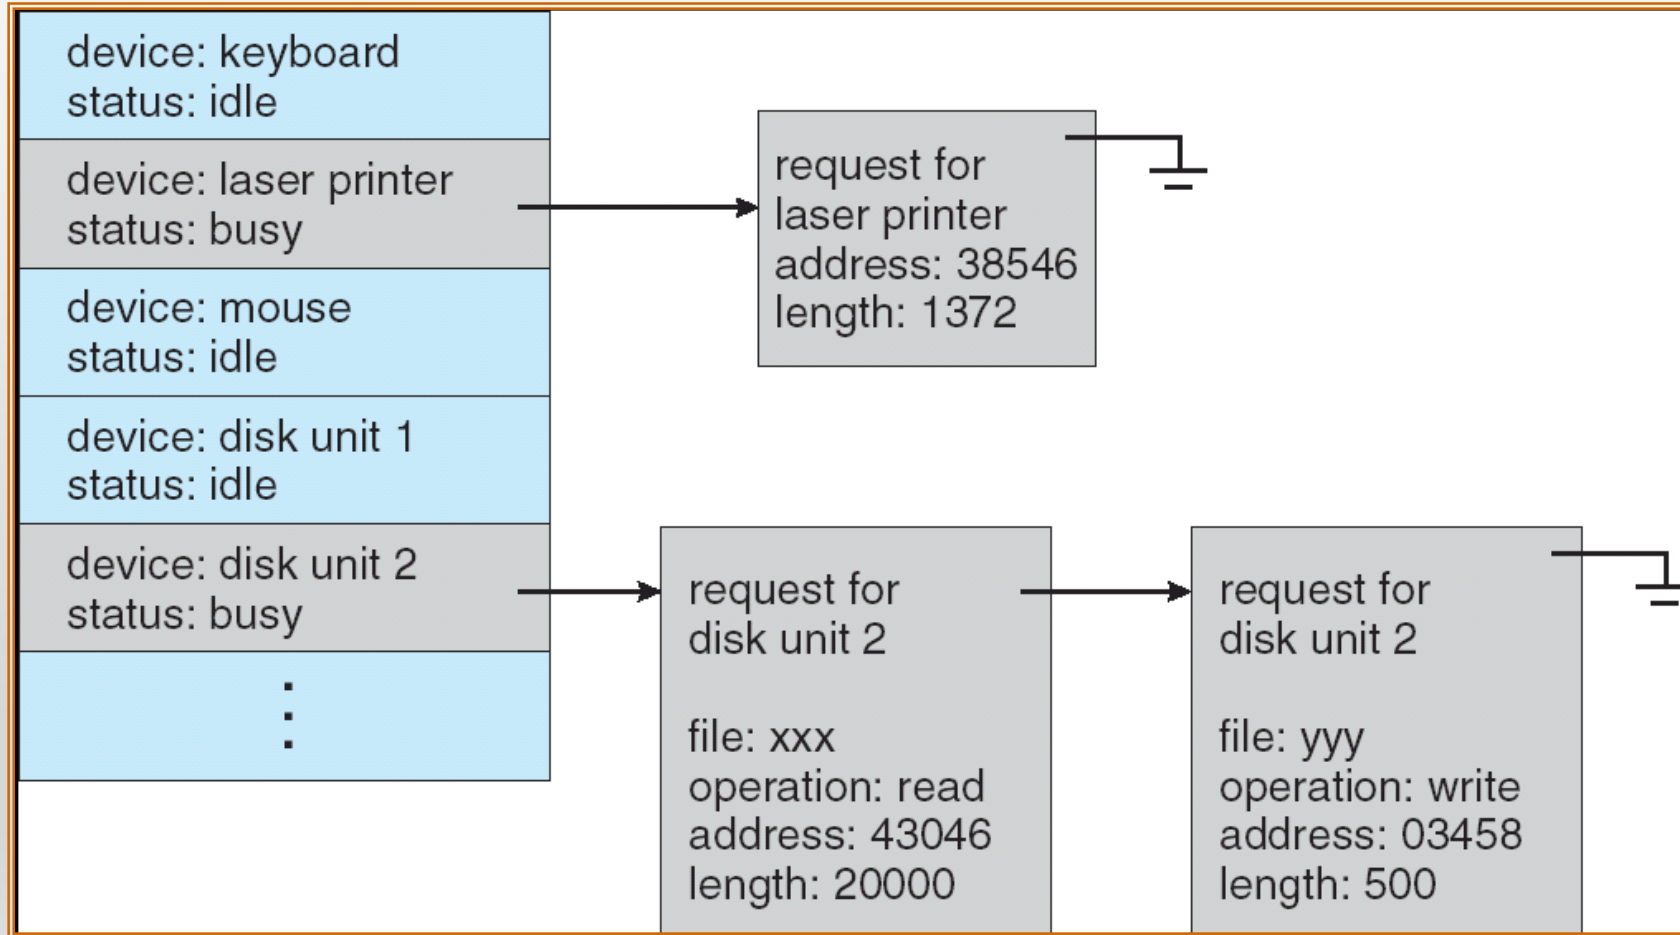


# UNIX I/O Kernel Data Structure

- Linux treats all I/O devices like a file



# Device-status Table



# Blocking and Nonblocking I/O

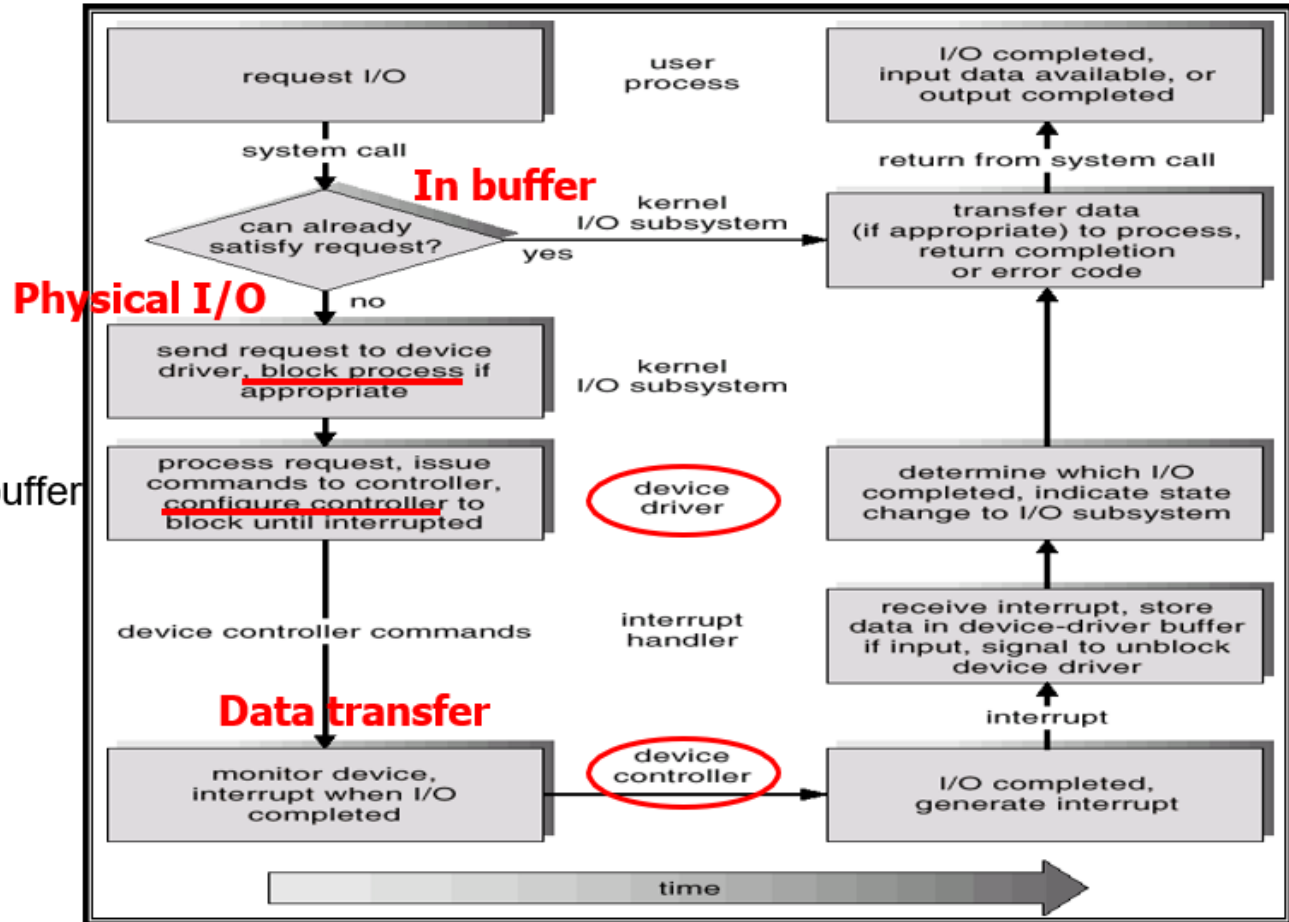
- Blocking - process suspended until I/O completed
  - Easy to use and understand
  - Insufficient for some needs
  - Use for **synchronous** communication & I/O
- Nonblocking
  - Implemented via **multi-threading**
  - Returns quickly with count of bytes read or written
  - Use for **asynchronous** communication & I/O

# Life Cycle of An I/O Request

Check buffer  
cache

Move process  
from run queue  
to wait queue

Allocate kernel buffer  
Schedule I/O



# Performance

- I/O is a major factor in **system** performance
  - It places heavy demands on the CPU to **execute device driver code**
  - The resulting **context switches** stress the CPU and its hardware caches
  - I/O loads down the memory bus during **data copy** between controllers and physical memory, ...
  - **Interrupt handling** is a relatively expensive task
    - Busy-waiting could be more efficient than interrupt- driven **if I/O time is small**

# Improving Performance

- Reduce **number of context switches**
- Reduce **data copying**
- Reduce **interrupts** by using large transfers, smart controllers, polling
- Use **DMA**
- Balance CPU, memory, bus, and I/O performance for highest throughput

# Review Slides ( II )

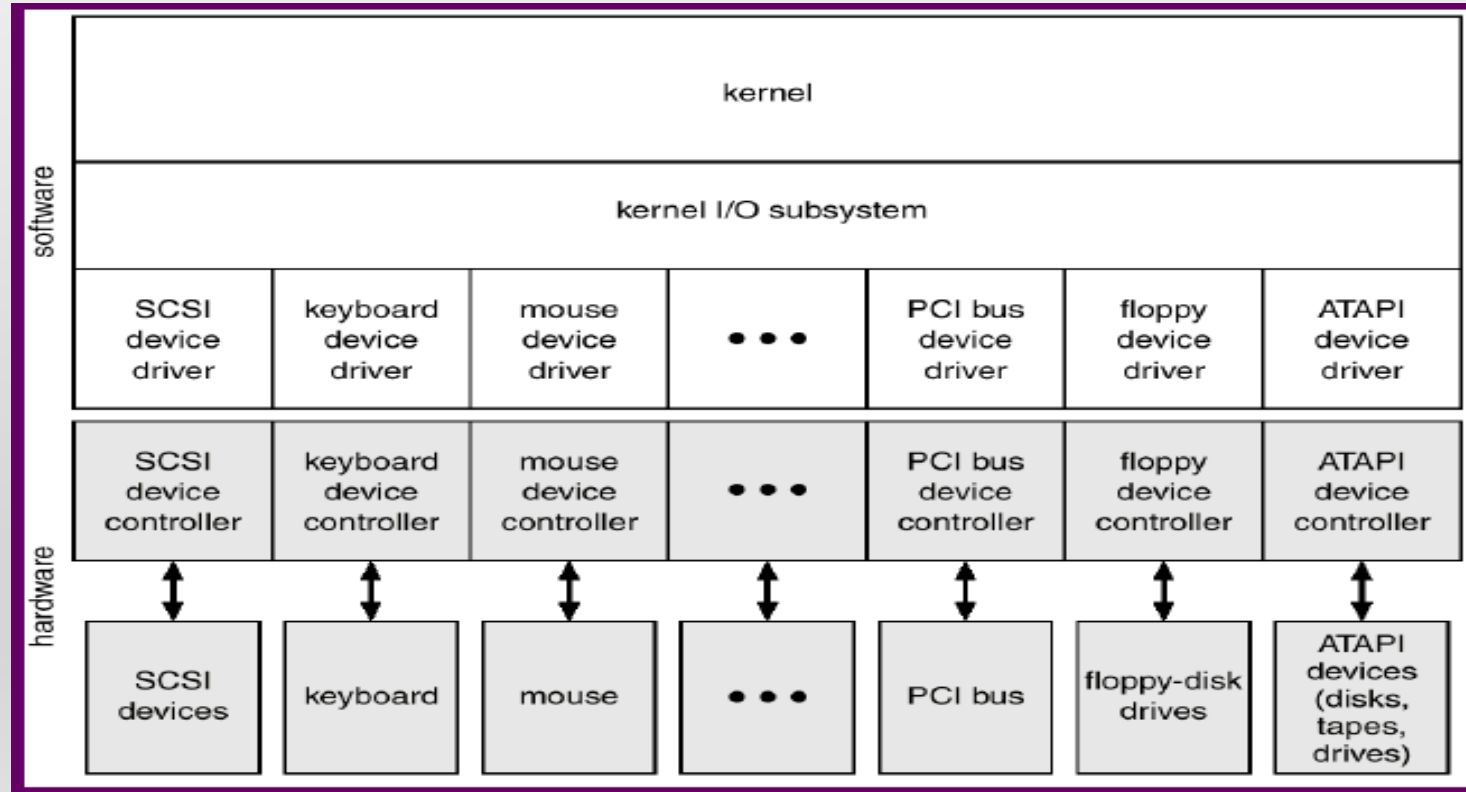
- What are the key I/O services
  - Scheduling
  - Cache
  - Buffering
  - Spooling
  - Error handling
  - I/ protection
- How to improve system performance?

# Application I/O Interface



# A Kernel I/O Structure

- **Device drivers**: a uniform device-access **interface to the I/O subsystem**; hide the differences among device controllers from the I/O sub-system of OS



# Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read&write	CD-ROM graphics controller disk

# I/O Device Class

- Device class is fairly standard across different OS
  - Block I/O
  - Char-stream I/O
  - Memory-mapped file access
  - Network sockets
  - Clock & timer interfaces
- Back-door interfaces (e.g. `ioctl()` )
  - Enable an application to access any functionality implemented by a device driver **without the need to invent a new system call**

# Block & Char Devices

- **Block devices:** disk drives
  - system calls: `read( )`, `write( )`, `seek( )`
  - Memory-mapped file can be layered on top
- **Char-stream** devices: mouse, keyboard, serial ports
  - system calls: `get( )`, `put( )`
  - Libraries layered on top allow line editing

# Network Devices

- Varying enough from block and character to have own interface
  - System call: `send()`, `recv()`, `select()`
  - `select()` returns which socket is waiting to send or receive, eliminates the need of busy waiting
- Many other approaches
  - pipes, FIFOs, STREAMS, message queues

# Reading Material & HW

- 13.1 - 13.6
- Problem Set
  - 13.2
  - 13.5
  - 13.6
  - 13.8

# Interrupt Vector Table

- Intel Pentium Processor:

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19-31	(Intel reserved, do not use)
32-255	maskable interrupts

# CPU and device Interrupt handshake

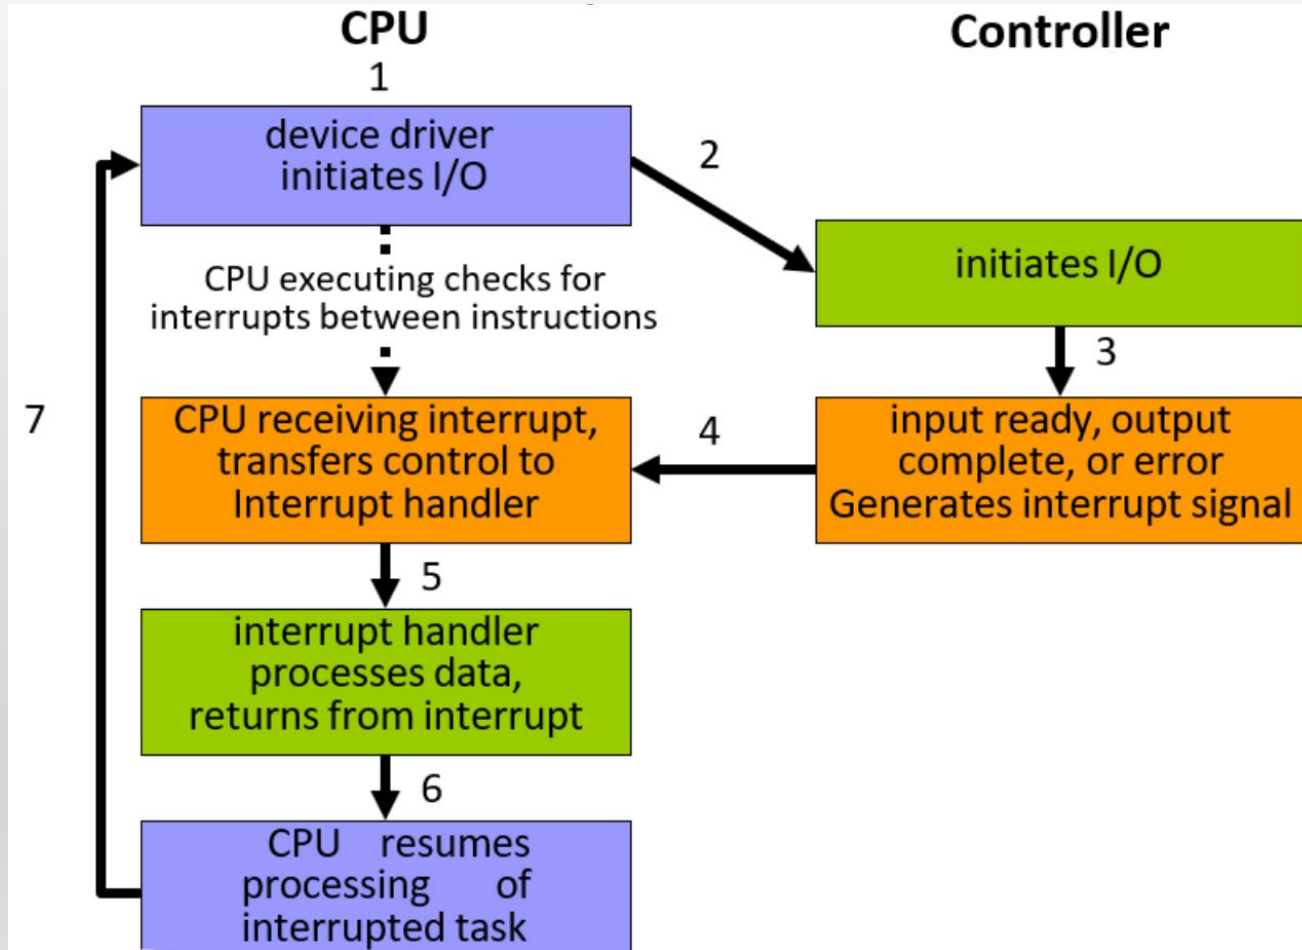
1. Device asserts **interrupt request (IRQ)**
2. CPU checks the **interrupt request line** at the **beginning of each instruction cycle**
3. **Save** the status and address of **interrupted process**
4. CPU acknowledges the interrupt and search the **interrupt vector** table for interrupt handler routines
5. CPU fetches the next instruction from the **interrupt handler routine**
6. **Restore interrupted process** after executing interrupt handler routine



# Interrupt Prioritization

- **Maskable interrupt:** interrupt with priority lower than current priority is not recognized until pending interrupt is complete
- **Non-maskable interrupt (NMI):** highest- priority, never masked
  - Often used for power-down, memory error

# Interrupt-Driven I/O



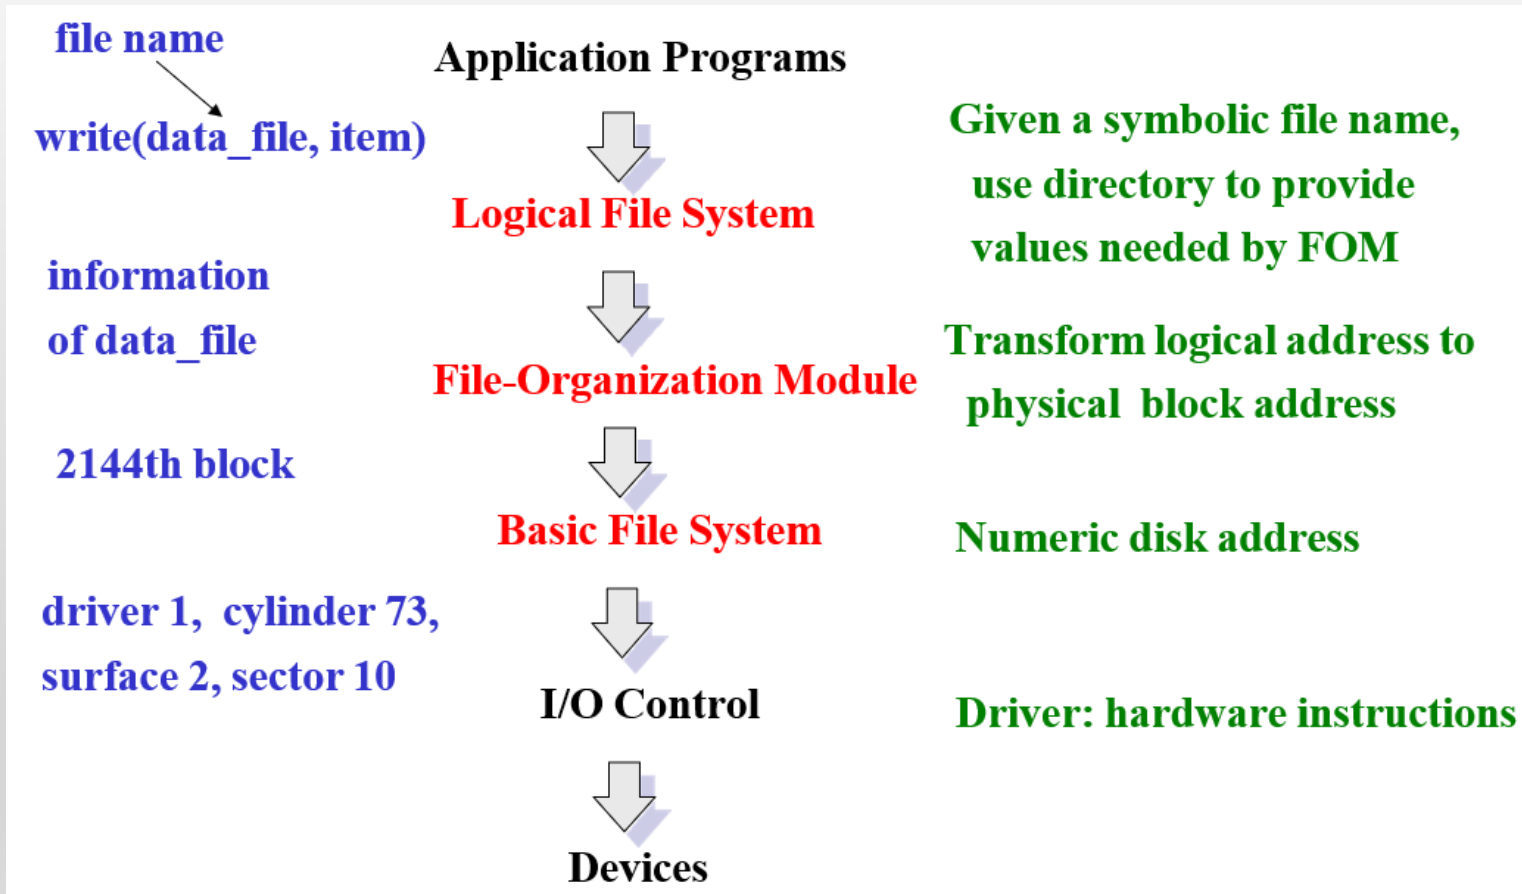
# Summary of Services in I/O Subsystem

- The management of the name space for files and devices
- Access control to files and devices
- Operation control
- File system space allocation
- Disk allocation
- Buffering, caching, and spooling
- I/O scheduling
- Device status monitoring, error handling, and failure recovery
- Device driver configuration and initialization

# I/O Requests to Hardware Operations

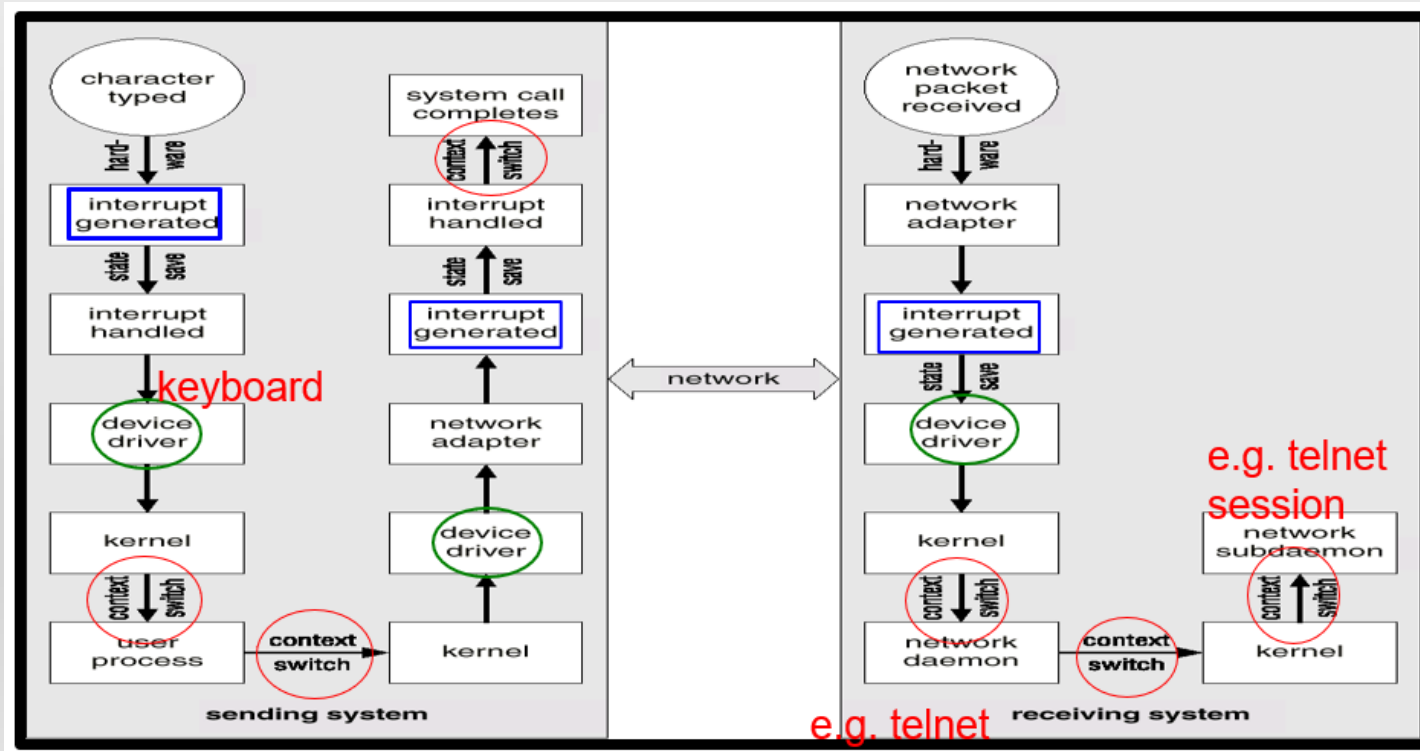
- Consider reading a file from disk for a process
  - Determine device holding file
  - Translate name to device representation
  - Physically read data from disk into buffer
  - Make data available to requesting process
  - Return control to process

# Layered File System revisited



# Intercomputer Communications

- Network traffic could cause high context switch rate
- Interrupt generated during keyboard & network I/O
- Context switch occurs between prog. & kernel (drivers)

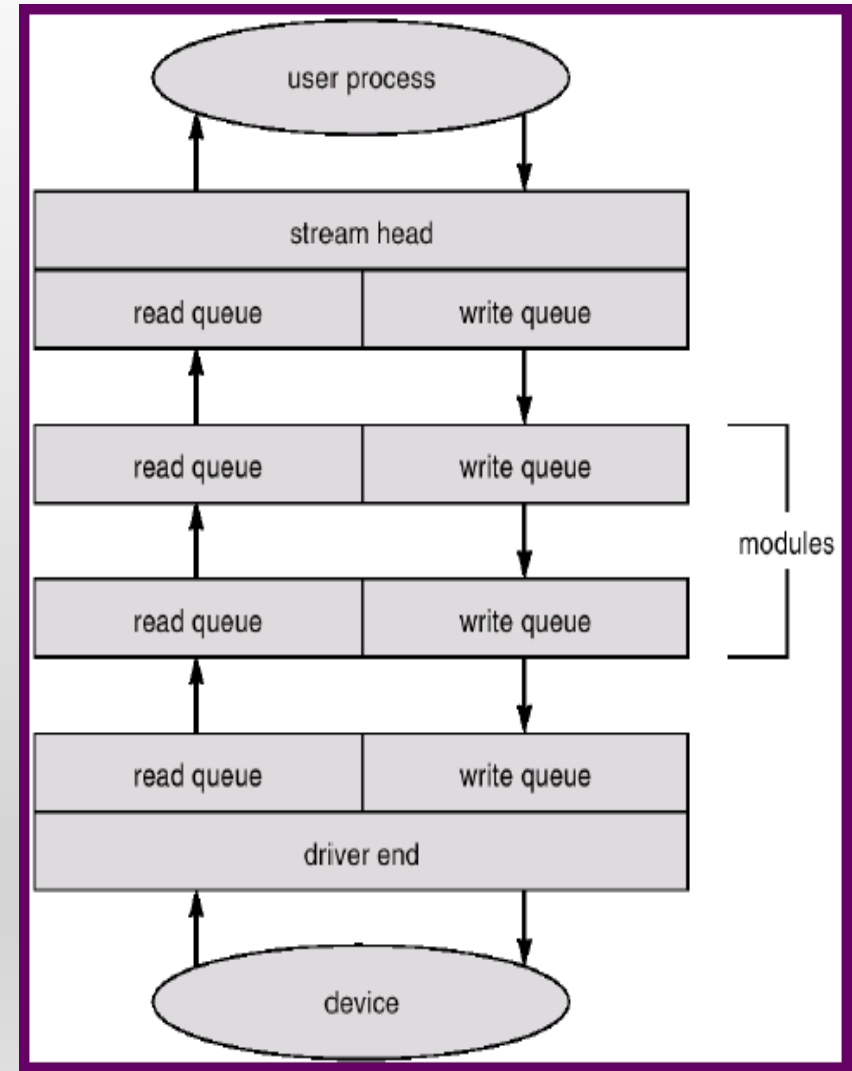


# STREAMS

- A **full-duplex communication** channel between a user-level process and a device
- STREAM provides a framework for a modular and incremental approach to writing device drivers and network protocols

# The STREAM Structure

- A STREAM consists of
  - STREAM head interfaces with user process
  - Driver end interfaces with the device
  - zero or more STREAM modules between them
- Each module contains a read and a write queue
- Message passing is used to communicate between queues





Q & A

*Thank you for your attention*